

MASSIMO DI PIERRO

ANNOTATED ALGORITHMS IN PYTHON

WITH APPLICATIONS IN PHYSICS, BIOLOGY, AND FINANCE (2ND ED)

EXPERTS4SOLUTIONS

Copyright 2013 by Massimo Di Pierro. All rights reserved.

THE CONTENT OF THIS BOOK IS PROVIDED UNDER THE TERMS OF THE CREATIVE COMMONS PUBLIC LICENSE BY-NC-ND 3.0.

<http://creativecommons.org/licenses/by-nc-nd/3.0/legalcode>

THE WORK IS PROTECTED BY COPYRIGHT AND/OR OTHER APPLICABLE LAW. ANY USE OF THE WORK OTHER THAN AS AUTHORIZED UNDER THIS LICENSE OR COPYRIGHT LAW IS PROHIBITED.

BY EXERCISING ANY RIGHTS TO THE WORK PROVIDED HERE, YOU ACCEPT AND AGREE TO BE BOUND BY THE TERMS OF THIS LICENSE. TO THE EXTENT THIS LICENSE MAY BE CONSIDERED TO BE A CONTRACT, THE LICENSOR GRANTS YOU THE RIGHTS CONTAINED HERE IN CONSIDERATION OF YOUR ACCEPTANCE OF SUCH TERMS AND CONDITIONS.

Limit of Liability/Disclaimer of Warranty: While the publisher and author have used their best efforts in preparing this book, they make no representations or warranties with respect to the accuracy or completeness of the contents of this book and specifically disclaim any implied warranties of merchantability or fitness for a particular purpose. No warranty may be created or extended by sales representatives or written sales materials. The advice and strategies contained herein may not be suitable for your situation. You should consult with a professional where appropriate. Neither the publisher nor the author shall be liable for any loss of profit or any other commercial damages, including but not limited to special, incidental, consequential, or other damages.

For more information about appropriate use of this material, contact:

Massimo Di Pierro
School of Computing
DePaul University
243 S Wabash Ave
Chicago, IL 60604 (USA)
Email: massimo.dipierro@gmail.com

Library of Congress Cataloging-in-Publication Data:

ISBN: 978-0-9911604-0-2

Build Date: June 6, 2017

to my parents

Contents

1	Introduction	15
1.1	Main Ideas	16
1.2	About Python	19
1.3	Book Structure	19
1.4	Book Software	21
2	Overview of the Python Language	23
2.1	About Python	23
2.1.1	Python versus Java and C++ syntax	24
2.1.2	help, dir	24
2.2	Types of variables	25
2.2.1	int and long	26
2.2.2	float and decimal	27
2.2.3	complex	30
2.2.4	str	30
2.2.5	list and array	31
2.2.6	tuple	33
2.2.7	dict	35
2.2.8	set	36
2.3	Python control flow statements	38
2.3.1	for...in	38
2.3.2	while	40
2.3.3	if...elif...else	41
2.3.4	try...except...else...finally	41
2.3.5	def...return	43

2.3.6	lambda	46
2.4	Classes	47
2.4.1	Special methods and operator overloading	49
2.4.2	class Financial Transaction	50
2.5	File input/output	51
2.6	How to import modules	52
2.6.1	math and cmath	53
2.6.2	os	53
2.6.3	sys	54
2.6.4	datetime	54
2.6.5	time	55
2.6.6	urllib and json	55
2.6.7	pickle	58
2.6.8	sqlite	59
2.6.9	numpy	64
2.6.10	matplotlib	66
2.6.11	ocl	72
3	Theory of Algorithms	75
3.1	Order of growth of algorithms	76
3.1.1	Best and worst running times	79
3.2	Recurrence relations	83
3.2.1	Reducible recurrence relations	85
3.3	Types of algorithms	88
3.3.1	Memoization	90
3.4	Timing algorithms	93
3.5	Data structures	94
3.5.1	Arrays	94
3.5.2	List	94
3.5.3	Stack	95
3.5.4	Queue	95
3.5.5	Sorting	96
3.6	Tree algorithms	98
3.6.1	Heapsort and priority queues	98
3.6.2	Binary search trees	102

3.6.3	Other types of trees	104
3.7	Graph algorithms	104
3.7.1	Breadth-first search	106
3.7.2	Depth-first search	108
3.7.3	Disjoint sets	109
3.7.4	Minimum spanning tree: Kruskal	110
3.7.5	Minimum spanning tree: Prim	112
3.7.6	Single-source shortest paths: Dijkstra	114
3.8	Greedy algorithms	116
3.8.1	Huffman encoding	116
3.8.2	Longest common subsequence	119
3.8.3	Needleman–Wunsch	121
3.8.4	Continuous Knapsack	123
3.8.5	Discrete Knapsack	124
3.9	Artificial intelligence and machine learning	128
3.9.1	Clustering algorithms	128
3.9.2	Neural network	132
3.9.3	Genetic algorithms	138
3.10	Long and infinite loops	140
3.10.1	P, NP, and NPC	140
3.10.2	Cantor’s argument	141
3.10.3	Gödel’s theorem	142
4	Numerical Algorithms	145
4.1	Well-posed and stable problems	145
4.2	Approximations and error analysis	146
4.2.1	Error propagation	148
4.2.2	buckingham	149
4.3	Standard strategies	150
4.3.1	Approximate continuous with discrete	150
4.3.2	Replace derivatives with finite differences	151
4.3.3	Replace nonlinear with linear	153
4.3.4	Transform a problem into a different one	154
4.3.5	Approximate the true result via iteration	155
4.3.6	Taylor series	155

4.3.7	Stopping Conditions	162
4.4	Linear algebra	164
4.4.1	Linear systems	164
4.4.2	Examples of linear transformations	171
4.4.3	Matrix inversion and the Gauss–Jordan algorithm	173
4.4.4	Transposing a matrix	175
4.4.5	Solving systems of linear equations	176
4.4.6	Norm and condition number again	177
4.4.7	Cholesky factorization	180
4.4.8	Modern portfolio theory	182
4.4.9	Linear least squares, χ^2	185
4.4.10	Trading and technical analysis	189
4.4.11	Eigenvalues and the Jacobi algorithm	191
4.4.12	Principal component analysis	194
4.5	Sparse matrix inversion	196
4.5.1	Minimum residual	197
4.5.2	Stabilized biconjugate gradient	198
4.6	Solvers for nonlinear equations	200
4.6.1	Fixed-point method	201
4.6.2	Bisection method	202
4.6.3	Newton’s method	203
4.6.4	Secant method	204
4.7	Optimization in one dimension	204
4.7.1	Bisection method	205
4.7.2	Newton’s method	205
4.7.3	Secant method	205
4.7.4	Golden section search	206
4.8	Functions of many variables	208
4.8.1	Jacobian, gradient, and Hessian	209
4.8.2	Newton’s method (solver)	211
4.8.3	Newton’s method (optimize)	212
4.8.4	Improved Newton’s method (optimize)	212
4.9	Nonlinear fitting	213
4.10	Integration	217
4.10.1	Quadrature	219

4.11	Fourier transforms	221
4.12	Differential equations	225
5	Probability and Statistics	229
5.1	Probability	229
5.1.1	Conditional probability and independence	231
5.1.2	Discrete random variables	232
5.1.3	Continuous random variables	234
5.1.4	Covariance and correlations	236
5.1.5	Strong law of large numbers	238
5.1.6	Central limit theorem	238
5.1.7	Error in the mean	240
5.2	Combinatorics and discrete random variables	240
5.2.1	Different plugs in different sockets	240
5.2.2	Equivalent plugs in different sockets	241
5.2.3	Colored cards	242
5.2.4	Gambler's fallacy	243
6	Random Numbers and Distributions	245
6.1	Randomness, determinism, chaos and order	245
6.2	Real randomness	246
6.2.1	Memoryless to Bernoulli distribution	247
6.2.2	Bernoulli to uniform distribution	248
6.3	Entropy generators	249
6.4	Pseudo-randomness	249
6.4.1	Linear congruential generator	250
6.4.2	Defects of PRNGs	252
6.4.3	Multiplicative recursive generator	252
6.4.4	Lagged Fibonacci generator	253
6.4.5	Marsaglia's add-with-carry generator	253
6.4.6	Marsaglia's subtract-and-borrow generator	254
6.4.7	Lüscher's generator	254
6.4.8	Knuth's polynomial congruential generator . . .	254
6.4.9	PRNGs in cryptography	255
6.4.10	Inverse congruential generator	256

6.4.11	Mersenne twister	256
6.5	Parallel generators and independent sequences	257
6.5.1	Non-overlapping blocks	258
6.5.2	Leapfrogging	259
6.5.3	Lehmer trees	260
6.6	Generating random numbers from a given distribution .	260
6.6.1	Uniform distribution	261
6.6.2	Bernoulli distribution	262
6.6.3	Biased dice and table lookup	263
6.6.4	Fishman–Yarberry method	264
6.6.5	Binomial distribution	266
6.6.6	Negative binomial distribution	268
6.6.7	Poisson distribution	270
6.7	Probability distributions for continuous random variables	272
6.7.1	Uniform in range	272
6.7.2	Exponential distribution	273
6.7.3	Normal/Gaussian distribution	275
6.7.4	Pareto distribution	278
6.7.5	In and on a circle	279
6.7.6	In and on a sphere	279
6.8	Resampling	280
6.9	Binning	281
7	Monte Carlo Simulations	283
7.1	Introduction	283
7.1.1	Computing π	283
7.1.2	Simulating an online merchant	286
7.2	Error analysis and the bootstrap method	289
7.3	A general purpose Monte Carlo engine	291
7.3.1	Value at risk	292
7.3.2	Network reliability	294
7.3.3	Critical mass	296
7.4	Monte Carlo integration	299
7.4.1	One-dimensional Monte Carlo integration	299
7.4.2	Two-dimensional Monte Carlo integration	301

7.4.3	n -dimensional Monte Carlo integration	302
7.5	Stochastic, Markov, Wiener, and processes	303
7.5.1	Discrete random walk (Bernoulli process)	304
7.5.2	Random walk: Ito process	305
7.6	Option pricing	306
7.6.1	Pricing European options: Binomial tree	308
7.6.2	Pricing European options: Monte Carlo	310
7.6.3	Pricing any option with Monte Carlo	312
7.7	Markov chain Monte Carlo (MCMC) and Metropolis	314
7.7.1	The Ising model	317
7.8	Simulated annealing	321
7.8.1	Protein folding	321
8	Parallel Algorithms	325
8.1	Parallel architectures	326
8.1.1	Flynn taxonomy	327
8.1.2	Network topologies	328
8.1.3	Network characteristics	331
8.2	Parallel metrics	332
8.2.1	Latency and bandwidth	332
8.2.2	Speedup	334
8.2.3	Efficiency	335
8.2.4	Isoefficiency	336
8.2.5	Cost	337
8.2.6	Cost optimality	338
8.2.7	Amdahl's law	338
8.3	Message passing	339
8.3.1	Broadcast	344
8.3.2	Scatter and collect	345
8.3.3	Reduce	347
8.3.4	Barrier	349
8.3.5	Global running times	350
8.4	mpi4py	351
8.5	Master-Worker and Map-Reduce	351
8.6	pyOpenCL	355

8.6.1	A first example with PyOpenCL	356
8.6.2	Laplace solver	358
8.6.3	Portfolio optimization (in parallel)	363
9	Appendices	367
9.1	Appendix A: Math Review and Notation	367
9.1.1	Symbols	367
9.1.2	Set theory	368
9.1.3	Logarithms	371
9.1.4	Finite sums	372
9.1.5	Limits ($n \rightarrow \infty$)	373
	Index	379
	Bibliography	383

1

Introduction

This book is assembled from lectures given by the author over a period of 10 years at the School of Computing of DePaul University. The lectures cover multiple classes, including Analysis and Design of Algorithms, Scientific Computing, Monte Carlo Simulations, and Parallel Algorithms. These lectures teach the core knowledge required by any scientist interested in numerical algorithms and by students interested in computational finance.

The notes are not comprehensive, yet they try to identify and describe the most important concepts taught in those courses using a few common tools and unified notation.

In particular, these notes do not include proofs; instead, they provide definitions and annotated code. The code is built in a modular way and is reused as much as possible throughout the book so that no step of the computations is left to the imagination. Each function defined in the code is accompanied by one or more examples of practical applications.

We take an interdisciplinary approach by providing examples in finance, physics, biology, and computer science. This is to emphasize that, although we often compartmentalize knowledge, there are very few ideas and methodologies that constitute the foundations of them all. Ultimately, this book is about problem solving using computers. The algorithms you

will learn can be applied to different disciplines. Throughout history, it is not uncommon that an algorithm invented by a physicist would find application in, for example, biology or finance.

Almost all of the algorithms written in this book can be found in the `nlib` library:

<https://github.com/mdipierro/nlib>

1.1 Main Ideas

Even if we cover many different algorithms and examples, there are a few central ideas in this book that we try to emphasize over and over.

The first idea is that we can simplify the solution of a problem by using an approximation and then systematically improve our approximation by iterating and computing corrections.

The divide-and-conquer methodology can be seen as an example of this approach. We do this with the insertion sort when we sort the first two numbers, then we sort the first three, then we sort the first four, and so on. We do it with merge sort when we sort each set of two numbers, then each set of four, then each set of eight, and so on. We do it with the Prim, Kruskal, and Dijkstra algorithms when we iterate over the nodes of a graph, and as we acquire knowledge about them, we use it to update the information about the shortest paths.

We use this approach in almost all our numerical algorithms because any differentiable function can be approximated with a linear function:

$$f(x + \delta x) \simeq f(x) + f'(x)\delta x \quad (1.1)$$

We use this formula in the Newton's method to solve nonlinear equations and optimization problems, in one or more dimensions.

We use the same approximation in the fix point method, which we use to solve equations like $f(x) = 0$; in the minimum residual and conjugate gradient methods; and to solve the Laplace equation in the last chapter of

the book. In all these algorithms, we start with a random guess for the solution, and we iteratively find a better one until convergence.

The second idea of the book is that certain quantities are random, but even random numbers have patterns that we can capture using instruments like distributions and correlations. The presence of these patterns helps us model those systems that may have a random output (e.g., nuclear reactions, financial systems) and also helps us in computations. In fact, we can use random numbers to compute quantities that are not random (Monte Carlo methods). The most common approximation that we make in different parts of the book is that when a random variable x is localized at a point with a given uncertainty, δx , then its distribution is Gaussian. Thanks to the properties of Gaussian random numbers, we conclude the following:

- Using the linear approximation (our first big idea), if $z = f(x)$, the uncertainty in the output is

$$\delta z = f'(x)\delta x \quad (1.2)$$

- If we add two independent Gaussian random variables $z = x + y$, the uncertainty in the output is

$$\delta z = \sqrt{\delta x^2 + \delta y^2} \quad (1.3)$$

- If we add N independent and identically distributed Gaussian variables $z = \sum x_i$, the uncertainty in the output is

$$\delta z = \sqrt{N}\delta x \quad (1.4)$$

We use this over and over, for example, when relating the volatility over different time intervals (daily, yearly).

- If we compute an average of N independent and identically distributed Gaussian random variables, $z = 1/N \sum x_i$, the uncertainty in the average is

$$\delta z = \delta x / \sqrt{N} \quad (1.5)$$

We use this to estimate the error on the average in a Monte Carlo computation. In that case, we write it as $d\mu = \sigma / \sqrt{N}$, and σ is the standard deviation of $\{x_i\}$.

The third idea is that the time it takes to run an iterative algorithm is proportional to the number of iterations. It is therefore our goal to minimize the number of iterations required to reach a target precision. We develop a language to compare algorithms based on their running time and classify algorithms into categories. This is useful to choose the best algorithm based on the problem at hand.

In the chapter on parallel algorithms, we learn how to distribute those iterations over multiple parallel processes and how to break individual iterations into independent steps that can be executed concurrently on parallel processes, to reduce the total time required to obtain a solution within a given target precision. In the parallel case, the running time acquires an overhead that depends on the communication patterns between the parallel processes, the communication latency, and bandwidth.

In the ultimate analysis, we can even try to understand ourselves as a parallel machine that models the input from the world by approximations. The brain is a graph that can be modeled by a neural network. The learning process is an ongoing optimization process in which the brain adjusts its synapses to produce better and better responses. The decision process mimics a search tree. We solve problems by searching for the most similar problems that we have encountered before, then we refine the solution. Our DNA is a code that evolved to efficiently compress the information necessary to grow us from a single cell into a complex being. We evolved according to evolutionary mechanisms that can be modeled using genetic algorithms. We can find our similarities with other organisms using the longest common subsequence algorithm. We can reconstruct our evolutionary tree using shortest-path algorithms and find out how we came to be.

1.2 About Python

The programming language used in this book is Python [1] version 2.7. This is because Python algorithms are very similar to the corresponding pseudo-code, and therefore this language is easy to read and understand compared to other languages such as C++ or Java. Moreover, Python is a popular language in many Universities and Companies (including Google).

The goal of the book is to explain the algorithms by building them from scratch. It is not our goal to teach the user about existing libraries that may be (and often are) faster than our implementation. Two notable examples are NumPy [2] and SciPy [3]. These libraries provide a Python interface to the BLAS and LAPACK libraries for linear algebra and applications. Although we wholeheartedly recommend using them when developing production code, we believe they are not appropriate for teaching the algorithms themselves because those algorithms are written in C, FORTRAN, and assembly languages and are not easy to read.

1.3 Book Structure

This book is divided into the following chapters:

- This introduction.
- An introduction to the Python programming language. The introduction assumes the reader is not new to basic programming concepts, such as conditionals, loops, and function calls, and teaches the basic syntax of the Python language, with particular focus on those built-in modules that are important for scientific applications (math, cmath, decimal, random) and a few others.
- Chapter 3 is a short review of the general theory of algorithms with applications. There we review how to determine the running time of an algorithm from simple loops to more complex recursive algorithms. We review basic data structures used to store information such as lists,

arrays, stacks, queues, trees, and graphs. We also review the classification of basic algorithms such as divide-and-conquer, dynamic programming, and greedy algorithms. In the examples, we peek into complex algorithms such as Shannon–Fano compression, a maze solver, a clustering algorithm, and a neural network.

- In chapter 4, we talk about traditional numerical algorithms, in particular, linear algebra, solvers, optimizers, integrators, and Fourier–Laplace transformations. We start by reviewing the concept of Taylor series and their convergence to understand approximations, sources of error, and convergence. We then use those concepts to build more complex algorithms by systematically improving their first-order (linear) approximation. Linear algebra serves us as a tool to approximate and implement functions of many variables.
- In chapter 5, we provide a review of probability and statistics and implement basic Python functions to perform statistical analysis of random variables.
- In chapter 6, we discuss algorithms to generate random numbers from many distributions. Python already has a built-in module to generate random numbers, and in subsequent chapters, we utilize it, yet in this chapter, we discuss in detail how pseudo random number generators work and their pitfalls.
- In chapter 7, we write about Monte Carlo simulations. This is a numerical technique that utilizes random numbers to solve otherwise deterministic problems. For example, in chapter 4, we talk about numerical integration in one dimension. Those algorithms can be extended to perform numerical integration in a few (two, three, sometimes four) dimensions, but they fail for very large numbers of dimensions. That is where Monte Carlo integration comes to our rescue, as it increasingly becomes the integration method of choice as the number of variables increases. We present applications of Monte Carlo simulations.
- In chapter 8, we discuss parallel algorithms. There are many paradigms for parallel programming these days, and the tendency is toward inhomogeneous architectures. Although we review many different

types of architectures, we focus on three programming paradigms that have been very successful: message-passing, map-reduce, and multi-threaded GPU programming. In the message-passing case, we create a simple “parallel simulator” (psim) in Python that allows us to understand the basic ideas behind message passing and issues with different network topologies. In the GPU case, we use pyOpenCL [4] and ocl [5], a Python-to-OpenCL compiler that allows us to write Python code and convert it in real time to OpenCL for running on the GPU.

- Finally, in the appendix, we provide a compendium of useful formulas and definitions.

1.4 Book Software

We utilize the following software libraries developed by the author and available under an Open Source BSD License:

- <http://github.com/mdipierro/nlib>
- <http://github.com/mdipierro/buckingham>
- <http://github.com/mdipierro/psim>
- <http://github.com/mdipierro/ocl>

We also utilize the following third party libraries:

- <http://www.numpy.org/>
- <http://matplotlib.org/>
- <https://github.com/michaelfairley/mincemeatpy>
- <http://mpi4py.scipy.org/>
- <http://mathematician.de/software/pyopencl>

All the code included in these notes is released by the author under the three-clause BSD License.

Acknowledgments

Many thanks to Alan Etkins, Brian Fox, Dan Bowker, Ethan Sudman, Holly Monteith, Konstantinos Moutselos, Michael Gheith, Paula Mikrut, Sean Neilan, and John Plamondon for reviewing different editions of this book. We also thank all the students of our classes for their useful comments and suggestions. Finally, we thank Wikipedia, from which we borrowed a few ideas and examples.

Overview of the Python Language

2.1 About Python

Python is a general-purpose high-level programming language. Its design philosophy emphasizes programmer productivity and code readability. It has a minimalist core syntax with very few basic commands and simple semantics. It also has a large and comprehensive standard library, including an Application Programming Interface (API) to many of the underlying operating system (OS) functions. Python provides built-in objects such as linked lists (`list`), tuples (`tuple`), hash tables (`dict`), arbitrarily long integers (`long`), complex numbers, and arbitrary precision decimal numbers.

Python supports multiple programming paradigms, including object-oriented (`class`), imperative (`def`), and functional (`lambda`) programming. Python has a dynamic type system and automatic memory management using reference counting (similar to Perl, Ruby, and Scheme).

Python was first released by Guido van Rossum in 1991 [6]. The language has an open, community-based development model managed by the nonprofit Python Software Foundation. There are many interpreters and compilers that implement the Python language, including one in Java (Jython), one built on .Net (IronPython), and one built in Python itself

(PyPy). In this brief review, we refer to the reference C implementation created by Guido.

You can find many tutorials, the official documentation, and library references of the language on the official Python website. [1]

For additional Python references, we can recommend the books in ref. [6] and ref. [7].

You may skip this chapter if you are already familiar with the Python language.

2.1.1 Python versus Java and C++ syntax

	Java/C++	Python
assignment	<i>a = b;</i>	<i>a = b</i>
comparison	<i>if (a == b)</i>	<i>if a == b:</i>
loops	<i>for(a = 0; a < n; a++)</i>	<i>for a in range(0, n):</i>
block	Braces {...}	indentation
function	<i>float f(float a) {</i>	<i>def f(a):</i>
function call	<i>f(a)</i>	<i>f(a)</i>
arrays/lists	<i>a[i]</i>	<i>a[i]</i>
member	<i>a.member</i>	<i>a.member</i>
nothing	<i>null / void*</i>	<i>None</i>

As in Java, variables that are primitive types (bool, int, float) are passed by copy, but more complex types, unlike C++, are passed by reference. This means when we pass an object to a function, in Python, we do not make a copy of the object, we simply define an alternate name for referencing the object in the function.

2.1.2 help, dir

The Python language provides two commands to obtain documentation about objects defined in the current scope, whether the object is built in or user defined.

We can ask for help about an object, for example, “1”:

```

1 >>> help(1)
2 Help on int object:
3
4 class int(object)
5 |   int(x[, base]) -> integer
6 |
7 |   Convert a string or number to an integer, if possible. A floating point
8 |   argument will be truncated towards zero (this does not include a string
9 |   representation of a floating point number!) When converting a string, use
10 |   the optional base. It is an error to supply a base when converting a
11 |   non-string. If the argument is outside the integer range a long object
12 |   will be returned instead.
13 |
14 |   Methods defined here:
15 |
16 |   __abs__(...)
17 |       x.__abs__() <==> abs(x)
18 ...

```

and because “1” is an integer, we get a description about the `int` class and all its methods. Here the output has been truncated because it is very long and detailed.

Similarly, we can obtain a list of object attributes (including methods) for any object using the command `dir`. For example:

```

1 >>> dir(1)
2 ['__abs__', '__add__', '__and__', '__class__', '__cmp__', '__coerce__',
3  '__delattr__', '__div__', '__divmod__', '__doc__', '__float__',
4  '__floordiv__', '__getattr__', '__getnewargs__', '__hash__', '__hex__',
5  '__index__', '__init__', '__int__', '__invert__', '__long__', '__lshift__',
6  '__mod__', '__mul__', '__neg__', '__new__', '__nonzero__', '__oct__',
7  '__or__', '__pos__', '__pow__', '__radd__', '__rand__', '__rdiv__',
8  '__rdivmod__', '__reduce__', '__reduce_ex__', '__repr__', '__rfloordiv__',
9  '__rlshift__', '__rmod__', '__rmul__', '__ror__', '__rpow__', '__rrshift__',
10 '__rshift__', '__rsub__', '__rtruediv__', '__rxor__', '__setattr__',
11 '__str__', '__sub__', '__truediv__', '__xor__']

```

2.2 Types of variables

Python is a dynamically typed language, meaning that variables do not have a type and therefore do not have to be declared. Variables may also change the type of value they hold through their lives. Values, on the

other hand, do have a type. You can query a variable for the type of value it contains:

```

1 >>> a = 3
2 >>> print type(a)
3 <type 'int'>
4 >>> a = 3.14
5 >>> print type(a)
6 <type 'float'>
7 >>> a = 'hello python'
8 >>> print type(a)
9 <type 'str'>

```

Python also includes, natively, data structures such as lists and dictionaries.

2.2.1 int and long

There are two types representing integer numbers: `int` and `long`. The difference is that `int` corresponds to the microprocessor's native bit length. Typically, this is 32 bits and can hold signed integers in range $[-2^{31}, +2^{31})$, whereas the `long` type can hold almost any arbitrary integer. It is important that Python automatically converts one into the other as necessary, and you can mix and match the two types in computations. Here is an example:

```

1 >>> a = 1024
2 >>> type(a)
3 <type 'int'>
4 >>> b = a**128
5 >>> print b
6 20815864389328798163850480654728171077230524494533409610638224700807216119346720
7 59602447888346464836968484322790856201558276713249664692981627981321135464152584
8 82590187784406915463666993231671009459188410953796224233873542950969577339250027
9 68876520583464697770622321657076833170056511209332449663781837603694136444406281
10 042053396870977465916057756101739472373801429441421111406337458176
11 >>> print type(b)
12 <type 'long'>

```

Computers represent 32-bit integer numbers by converting them to base 2. The conversion works in the following way:

```

1 def int2binary(n, nbits=32):
2     if n<0:
3         return [1 if bit==0 else 0 for bit in int2binary(-n-1,nbits)]

```

```

4 bits = [0]*nbits
5 for i in range(nbits):
6     n, bits[i] = divmod(n,2)
7 if n: raise OverflowError
8 return bits

```

The case $n < 0$ is called *two's complement* and is defined as the value obtained by subtracting the number from the largest power of 2 (2^{32} for 32 bits). Just by looking at the most significant bit, one can determine the sign of the binary number (1 for negative and 0 for zero or positive).

2.2.2 float and decimal

There are two ways to represent decimal numbers in Python: using the native double precision (64 bits) representation, `float`, or using the `decimal` module.

Most numerical problems are dealt with simply using `float`:

```

1 >>> pi = 3.141592653589793
2 >>> two_pi = 2.0 * pi

```

Floating point numbers are internally represented as follows:

$$x = \pm m 2^e \quad (2.1)$$

where x is the number, m is called the *mantissa* and is zero or a number in the range $[1,2)$, and e is called the *exponent*. The sign, m , and e can be computed using the following algorithm, which also writes their representation in binary:

```

1 def float2binary(x,nm=4,ne=4):
2     if x==0:
3         return 0, [0]*nm, [0]*ne
4     sign,mantissa, exponent = (1 if x<0 else 0),abs(x),0
5     while abs(mantissa)>=2:
6         mantissa,exponent = 0.5*mantissa,exponent+1
7     while 0<abs(mantissa)<1:
8         mantissa,exponent = 2.0*mantissa,exponent-1
9     mantissa = int2binary(int(2**(nm-1)*mantissa),nm)
10    exponent = int2binary(exponent,ne)
11    return sign, mantissa, exponent

```

Because the exponent is stored in a fixed number of bits (11 for a 64-bit floating point number), exponents smaller than -1022 and larger than 1023 cannot be represented. An arithmetic operation that returns a number smaller than $2^{-1022} \simeq 10^{-308}$ cannot be represented and results in an underflow error. An operation that returns a number larger than $2^{1023} \simeq 10^{308}$ also cannot be represented and results in an overflow error.

Here is an example of overflow:

```
1 >>> a = 10.0**200
2 >>> a*a
3 inf
```

And here is an example of underflow:

```
1 >>> a = 10.0**-200
2 >>> a*a
3 0.0
```

Another problem with finite precision arithmetic is the loss of precision in computation. Consider the case of the difference between two numbers with very different orders of magnitude. To compute the difference, the CPU reduces them to the same exponent (the largest of the two) and then computes the difference in the two mantissas. If two numbers differ for a factor 2^k , then the mantissa of the smallest number, in binary, needs to be shifted by k positions, thus resulting in a loss of information because the k least significant bits in the mantissa are ignored. If the difference between the two numbers is greater than a factor 2^{52} , all bits in the mantissa of the smallest number are ignored, and the smallest number becomes completely invisible.

Following is a practical example that produces an incorrect result:

```
1 >>> a = 1.0
2 >>> b = 2.0**53
3 >>> a+b-b
4 0.0
```

a simple example of what occurs internally in a processor to add two floating point numbers together. The IEEE 754 standard states that for 32-bit floating point numbers, the exponent has a range of -126 to $+127$:

```
1 262 in IEEE 754: 0 10000111 000001100000000000000000 (+ e:8 m:1.0234375)
```

```

2 3 in IEEE 754: 0 10000000 100000000000000000000000 (+ e:1 m:1.5)
3 265 in IEEE 754: 0 10000111 000010010000000000000000

```

To add 262.0 to 3.0, the exponents must be the same. The exponent of the lesser number is increased to the exponent of the greater number. In this case, 3's exponent must be increased by 7. Increasing the exponent by 7 means the mantissa must be shifted seven binary digits to the right:

```

1 0 10000111 000001100000000000000000
2 0 10000111 000000110000000000000000 (The implied '1' is also pushed seven
   places to the right)
3 -----
4 0 10000111 000010010000000000000000 which is the IEEE 754 format for 265.0

```

In the case of two numbers in which the exponent is greater than the number of digits in the mantissa, the smaller number is shifted right off the end. The effect is a zero added to the larger number.

In some cases, only some of the bits of the smaller number's mantissa are lost if a partial addition occurs.

This precision issue is always present but not always obvious. It may consist of a small discrepancy between the true value and the computed value. This difference may increase during the computation, in particular, in iterative algorithms, and may be sizable in the result of a complex algorithm.

Python also has a module for decimal floating point arithmetic that allows decimal numbers to be represented exactly. The class `Decimal` incorporates a notion of significant places (unlike the hardware-based binary floating point, the decimal module has a user-alterable precision):

```

1 >>> from decimal import Decimal, getcontext
2 >>> getcontext().prec = 28 # set precision
3 >>> Decimal(1) / Decimal(7)
4 Decimal('0.1428571428571428571428571428571429')

```

Decimal numbers can be used almost everywhere in place of floating point number arithmetic but are slower and should be used only where arbitrary precision arithmetic is required. It does not suffer from the overflow, underflow, and precision issues described earlier:

```

1 >>> from decimal import Decimal
2 >>> a = Decimal(10.0)**300

```

3
4

2.2.3

Python has native support for complex numbers. The imaginary unit is represented by the character `j`:

1
2
3
4
5
6
7
8
9

The real and imaginary parts of a complex number are stored as 64-bit floating point numbers.

Normal arithmetic operations are supported. The `cmath` module contains trigonometric and other functions for complex numbers. For example,

1
2
3
4

2.2.4

Python supports the use of two different types of strings: ASCII strings and Unicode strings. ASCII strings are delimited by `'...'`, `"..."`, `'''...'''`, or `"""..."""`. Triple quotes delimit multiline strings. Unicode strings start with a `u`, followed by the string containing Unicode characters. A Unicode string can be converted into an ASCII string by choosing an encoding (e.g., UTF8):

1
2
3

After executing these three commands, the resulting `a` is an ASCII string

storing UTF8 encoded characters.

It is also possible to write variables into strings in various ways:

```

1 >>> print 'number is ' + str(3)
2 number is 3
3 >>> print 'number is %s' % (3)
4 number is 3
5 >>> print 'number is %(number)s' % dict(number=3)
6 number is 3

```

The final notation is more explicit and less error prone and is to be preferred.

Many Python objects, for example, numbers, can be serialized into strings using `str` or `repr`. These two commands are very similar but produce slightly different output. For example,

```

1 >>> for i in [3, 'hello']:
2 ...     print str(i), repr(i)
3 3 3
4 hello 'hello'

```

For user-defined classes, `str` and `repr` can be defined and redefined using the special operators `__str__` and `__repr__`. These are briefly described later in this chapter. For more information on the topic, refer to the official Python documentation [8].

Another important characteristic of a Python string is that it is an iterable object, similar to a list:

```

1 >>> for i in 'hello':
2 ...     print i
3 h
4 e
5 l
6 l
7 o

```

2.2.5 list and array

The distinction between lists and arrays is usually in their implementation and in the relative difference in speed of the operations they can perform. Python defines a type called `list` that internally is implemented more like an array.

32 ANNOTATED ALGORITHMS IN PYTHON

The main methods of Python lists are `append`, `insert`, and `delete`. Other useful methods include `count`, `index`, `reverse`, and `sort`:

```
1 >>> b = [1, 2, 3]
2 >>> print type(b)
3 <type 'list'>
4 >>> b.append(8)
5 >>> b.insert(2, 7) # insert 7 at index 2 (3rd element)
6 >>> del b[0]
7 >>> print b
8 [2, 7, 3, 8]
9 >>> print len(b)
10 4
11 >>> b.append(3)
12 >>> b.reverse()
13 >>> print b, " 3 appears ", b.count(3), " times. The number 7 appears at index "
14 [3, 8, 3, 7, 2] 3 appears 2 times. The number 7 appears at index 3
```

Lists can be sliced:

```
1 >>> a = [2, 7, 3, 8]
2 >>> print a[:3]
3 [2, 7, 3]
4 >>> print a[1:]
5 [7, 3, 8]
6 >>> print a[-2:]
7 [3, 8]
```

and concatenated/joined:

```
1 >>> a = [2, 7, 3, 8]
2 >>> a = [2, 3]
3 >>> b = [5, 6]
4 >>> print a + b
5 [2, 3, 5, 6]
```

A list is iterable; you can loop over it:

```
1 >>> a = [1, 2, 3]
2 >>> for i in a:
3 ...     print i
4 1
5 2
6 3
```

A list can also be sorted in place with the `sort` method:

```
1 >>> a.sort()
```

There is a very common situation for which a *list comprehension* can be

used. Consider the following code:

```

1 >>> a = [1,2,3,4,5]
2 >>> b = []
3 >>> for x in a:
4 ...     if x % 2 == 0:
5 ...         b.append(x * 3)
6 >>> print b
7 [6, 12]
```

This code clearly processes a list of items, selects and modifies a subset of the input list, and creates a new result list. This code can be entirely replaced with the following list comprehension:

```

1 >>> a = [1,2,3,4,5]
2 >>> b = [x * 3 for x in a if x % 2 == 0]
3 >>> print b
4 [6, 12]
```

Python has a module called `array`. It provides an efficient array implementation. Unlike lists, array elements must all be of the same type, and the type must be either a char, short, int, long, float, or double. A type of char, short, int, or long may be either signed or unsigned. Notice these are C-types, not Python types.

```

1 >>> from array import array
2 >>> a = array('d', [1,2,3,4,5])
3 array('d', [1.0, 2.0, 3.0, 4.0, 5.0])
```

An array object can be used in the same way as a list, but its elements must all be of the same type, specified by the first argument of the constructor (“d” for double, “l” for signed long, “f” for float, and “c” for character). For a complete list of available options, refer to the official Python documentation.

Using “array” over “list” can be faster, but more important, the “array” storage is more compact for large arrays.

2.2.6 tuple

A tuple is similar to a list, but its size and elements are immutable. If a tuple element is an object, the object itself is mutable, but the reference to the object is fixed. A tuple is defined by elements separated by a comma

and optionally delimited by round parentheses:

```
1 >>> a = 1, 2, 3
2 >>> a = (1, 2, 3)
```

The round brackets are required for a tuple of zero elements such as

```
1 >>> a = () # this is an empty tuple
```

A trailing comma is required for a one-element tuple but not for two or more elements:

```
1 >>> a = (1) # not a tuple
2 >>> a = (1,) # this is a tuple of one element
3 >>> b = (1,2) # this is a tuple of two elements
```

Since lists are mutable; this works:

```
1 >>> a = [1, 2, 3]
2 >>> a[1] = 5
3 >>> print a
4 [1, 5, 3]
```

the element assignment does not work for a tuple:

```
1 >>> a = (1, 2, 3)
2 >>> print a[1]
3 2
4 >>> a[1] = 5
5 Traceback (most recent call last):
6   File "<stdin>", line 1, in <module>
7 TypeError: 'tuple' object does not support item assignment
```

A tuple, like a list, is an iterable object. Notice that a tuple consisting of a single element must include a trailing comma:

```
1 >>> a = (1)
2 >>> print type(a)
3 <type 'int'>
4 >>> a = (1,)
5 >>> print type(a)
6 <type 'tuple'>
```

Tuples are very useful for efficient packing of objects because of their immutability. The brackets are often optional. You may easily get each element of a tuple by assigning multiple variables to a tuple at one time:

```
1 >>> a = (2, 3, 'hello')
2 >>> (x, y, z) = a
3 >>> print x
4 2
5 >>> print z
```

```

6 hello
7 >>> a = 'alpha', 35, 'sigma' # notice the rounded brackets are optional
8 >>> p, r, q = a
9 print r
10 35

```

2.2.7 dict

A Python dict-ionary is a hash table that maps a key object to a value object:

```

1 >>> a = {'k': 'v', 'k2': 3}
2 >>> print a['k']
3 v
4 >>> print a['k2']
5 3
6 >>> 'k' in a
7 True
8 >>> 'v' in a
9 False

```

You will notice that the format to define a dictionary is the same as the JavaScript Object Notation [JSON]. Dictionaries may be nested:

```

1 >>> a = {'x': 3, 'y': 54, 'z': {'a': 1, 'b': 2}}
2 >>> print a['z']
3 {'a': 1, 'b': 2}
4 >>> print a['z']['a']
5 1

```

Keys can be of any hashable type (int, string, or any object whose class implements the `__hash__` method). Values can be of any type. Different keys and values in the same dictionary do not have to be of the same type. If the keys are alphanumeric characters, a dictionary can also be declared with the alternative syntax:

```

1 >>> a = dict(k='v', h2=3)
2 >>> print a['k']
3 v
4 >>> print a
5 {'h2': 3, 'k': 'v'}

```

Useful methods are `has_key`, `keys`, `values`, `items`, and `update`:

```

1 >>> a = dict(k='v', k2=3)
2 >>> print a.keys()
3 ['k2', 'k']

```

```

4 >>> print a.values()
5 [3, 'v']
6 >>> a.update({'n1': 'new item'})      # adding a new item
7 >>> a.update(dict(n2='newer item'))   # alternate method to add a new item
8 >>> a['n3'] = 'newest item'           # another method to add a new item
9 >>> print a.items()
10 [('k2', 3), ('k', 'v'), ('n3', 'newest item'), ('n2', 'newer item'), ('n1', 'new
    item')]

```

The items method produces a list of tuples, each containing a key and its associated value.

Dictionary elements and list elements can be deleted with the command `del`:

```

1 >>> a = [1, 2, 3]
2 >>> del a[1]
3 >>> print a
4 [1, 3]
5 >>> a = dict(k='v', h2=3)
6 >>> del a['h2']
7 >>> print a
8 {'k': 'v'}

```

Internally, Python uses the hash operator to convert objects into integers and uses that integer to determine where to store the value. Using a key that is not hashable will cause an `unhashable type` error:

```

1 >>> hash("hello world")
2 -1500746465
3 >>> k = [1,2,3]
4 >>> a = {k: '4'}
5 Traceback (most recent call last):
6   File "<stdin>", line 1, in <module>
7 TypeError: unhashable type: 'list'

```

2.2.8 set

A set is something between a list and a dictionary. It represents a non-ordered list of unique elements. Elements in a set cannot be repeated. Internally, it is implemented as a hash table, similar to a set of keys in a dictionary. A set is created using the set constructor. Its argument can be a list, a tuple, or an iterator:

```

1 >>> s = set([1,2,3,4,5,5,5]) # notice duplicate elements are removed

```

```

2 >>> print s
3 set([1,2,3,4,5])
4 >>> s = set((1,2,3,4,5))
5 >>> print s
6 set([1,2,3,4,5])
7 >>> s = set(i for i in range(1,6))
8 >>> print s
9 set([1, 2, 3, 4, 5])

```

Sets are not ordered lists therefore appending to the end is not applicable. Instead of append, add elements to a set using the add method:

```

1 >>> s = set()
2 >>> s.add(2)
3 >>> s.add(3)
4 >>> s.add(2)
5 >>> print s
6 set([2, 3])

```

Notice that the same element cannot be added twice (2 in the example). There is no exception or error thrown when trying to add the same element more than once.

Because sets are not ordered, the order in which you add items is not necessarily the order in which they will be returned:

```

1 >>> s = set([6, 'b', 'beta', -3.4, 'a', 3, 5.3])
2 >>> print (s)
3 set(['a', 3, 6, 5.3, 'beta', 'b', -3.4])

```

The set object supports normal set operations like union, intersection, and difference:

```

1 >>> a = set([1,2,3])
2 >>> b = set([2,3,4])
3 >>> c = set([2,3])
4 >>> print a.union(b)
5 set([1, 2, 3, 4])
6 >>> print a.intersection(b)
7 set([2, 3])
8 >>> print a.difference(b)
9 set([1])
10 >>> if len(c) == len(a.intersection(c)):
11 ...     print "c is a subset of a"
12 ... else:
13 ...     print "c is not a subset of a"
14 ...
15 c is a subset of a

```

To check for membership,

```
1 >>> 2 in a
2 True
```

2.3 Python control flow statements

Python uses indentation to delimit blocks of code. A block starts with a line ending with colon and continues for all lines that have a similar or higher indentation as the next line:

```
1 >>> i = 0
2 >>> while i < 3:
3 ...     print i
4 ...     i = i + 1
5 0
6 1
7 2
```

It is common to use four spaces for each level of indentation. It is a good policy not to mix tabs with spaces, which can result in (invisible) confusion.

2.3.1 for...in

In Python, you can loop over iterable objects:

```
1 >>> a = [0, 1, 'hello', 'python']
2 >>> for i in a:
3 ...     print i
4 0
5 1
6 hello
7 python
```

In the preceding example, you will notice that the loop index “i” takes on the values of each element in the list [0, 1, ‘hello’, ‘python’] sequentially. The Python range keyword creates a list of integers automatically that may be used in a “for” loop without manually creating a long list of numbers.

```
1 >>> a = range(0,5)
2 >>> print a
3 [0, 1, 2, 3, 4]
```

```

4 >>> for i in a:
5 ...     print i
6 0
7 1
8 2
9 3
10 4

```

The parameters for `range(a,b,c)` are as follows: the first parameter is the starting value of the list. The second parameter is the next value if the list contains one more element. The third parameter is the increment value.

The keyword `range` can also be called with one parameter. It is matched to “b” with the first parameter defaulting to 0 and the third to 1:

```

1 >>> print range(5)
2 [0, 1, 2, 3, 4]
3 >>> print range(53,57)
4 [53,54,55,56]
5 >>> print range(102,200,10)
6 [102, 112, 122, 132, 142, 152, 162, 172, 182, 192]
7 >>> print range(0,-10,-1)
8 [0, -1, -2, -3, -4, -5, -6, -7, -8, -9]

```

The keyword `range` is very convenient for creating a list of numbers; however, as the list grows in length, the memory required to store the list also grows. A more efficient option is to use the keyword `xrange`, which generates an iterable range instead of the entire list of elements.

This is equivalent to the C/C++/C#/Java syntax:

```

1 for(int i=0; i<4; i=i+1) { ... }

```

Another useful command is `enumerate`, which counts while looping and returns a tuple consisting of (index, value):

```

1 >>> a = [0, 1, 'hello', 'python']
2 >>> for (i, j) in enumerate(a): # the ( ) around i, j are optional
3 ...     print i, j
4 0 0
5 1 1
6 2 hello
7 3 python

```

You can jump out of a loop using `break`:

```

1 >>> for i in [1, 2, 3]:
2 ...     print i
3 ...     break

```

1

You can jump to the next loop iteration without executing the entire code block with `continue`:

```
1 >>> for i in [1, 2, 3]:
2 ...     print i
3 ...     continue
4 ...     print 'test'
5 1
6 2
7 3
```

Python also supports list comprehensions, and you can build lists using the following syntax:

```
1 >>> a = [i*i for i in [0, 1, 2, 3]:
2 >>> print a
3 [0, 1, 4, 9]
```

Sometimes you may need a counter to “count” the elements of a list while looping:

```
1 >>> a = [e*(i+1) for (i,e) in enumerate(['a','b','c','d'])]
2 >>> print a
3 ['a', 'bb', 'ccc', 'dddd']
```

2.3.2 while

Comparison operators in Python follow the C/C++/Java operators of `==`, `!=`, `<`, `>`, and so on. However, Python also accepts the `<>` operator as not equal to and is equivalent to `!=`. Logical operators are `and`, `or`, and `not`.

The `while` loop in Python works much as it does in many other programming languages, by looping an indefinite number of times and testing a condition before each iteration. If the condition is `False`, the loop ends:

```
1 >>> i = 0
2 >>> while i < 10:
3 ...     i = i + 1
4 >>> print i
5 10
```

The `for` loop was introduced earlier in this chapter.

There is no `loop...until` or `do...while` construct in Python.

2.3.3 if...elif...else

The use of conditionals in Python is intuitive:

```

1 >>> for i in range(3):
2 ...     if i == 0:
3 ...         print 'zero'
4 ...     elif i == 1:
5 ...         print 'one'
6 ...     else:
7 ...         print 'other'
8 zero
9 one
10 other

```

The `elif` means “else if.” Both `elif` and `else` clauses are optional. There can be more than one `elif` but only one `else` statement. Complex conditions can be created using the `not`, `and`, and `or` logical operators:

```

1 >>> for i in range(3):
2 ...     if i == 0 or (i == 1 and i + 1 == 2):
3 ...         print '0 or 1'

```

2.3.4 try...except...else...finally

Python can throw - pardon, raise - exceptions:

```

1 >>> try:
2 ...     a = 1 / 0
3 ... except Exception, e:
4 ...     print 'oops: %s' % e
5 ... else:
6 ...     print 'no problem here'
7 ... finally:
8 ...     print 'done'
9 oops: integer division or modulo by zero
10 done

```

If an exception is raised, it is caught by the `except` clause, and the `else` clause is not executed. The `finally` clause is always executed.

There can be multiple `except` clauses for different possible exceptions:

```

1 >>> try:
2 ...     raise SyntaxError
3 ... except ValueError:
4 ...     print 'value error'

```

42 ANNOTATED ALGORITHMS IN PYTHON

```
5 ... except SyntaxError:
6 ...     print 'syntax error'
7 syntax error
```

The finally clause is guaranteed to be executed while the except and else are not. In the following example, the function returns within a try block. This is bad practice, but it shows that the finally will execute regardless of the reason the try block is exited:

```
1 >>> def f(x):
2 ...     try:
3 ...         r = x*x
4 ...         return r # bad practice
5 ...     except:
6 ...         print "exception occurred %s" % e
7 ...     else:
8 ...         print "nothing else to do"
9 ...     finally:
10 ...        print "Finally we get here"
11 ...
12 >>> y = f(3)
13 Finally we get here
14 >>> print "result is ", y
15 result is 9
```

For every try, you must have either an except or a finally, while the else is optional.

Here is a list of built-in Python exceptions:

```
1 BaseException
2 +-- SystemExit
3 +-- KeyboardInterrupt
4 +-- Exception
5     +-- GeneratorExit
6     +-- StopIteration
7     +-- StandardError
8         | +-- ArithmeticError
9         | | +-- FloatingPointError
10        | | +-- OverflowError
11        | | +-- ZeroDivisionError
12        | +-- AssertionError
13        | +-- AttributeError
14        | +-- EnvironmentError
15        | | +-- IOError
16        | | +-- OSError
17        | | +-- WindowsError (Windows)
18        | | +-- VMSError (VMS)
```

```

19 |     +-- EOFError
20 |     +-- ImportError
21 |     +-- LookupError
22 |     |     +-- IndexError
23 |     |     +-- KeyError
24 |     +-- MemoryError
25 |     +-- NameError
26 |     |     +-- UnboundLocalError
27 |     +-- ReferenceError
28 |     +-- RuntimeError
29 |     |     +-- NotImplementedError
30 |     +-- SyntaxError
31 |     |     +-- IndentationError
32 |     |         +-- TabError
33 |     +-- SystemError
34 |     +-- TypeError
35 |     +-- ValueError
36 |     |     +-- UnicodeError
37 |     |         +-- UnicodeDecodeError
38 |     |         +-- UnicodeEncodeError
39 |     |         +-- UnicodeTranslateError
40 | +-- Warning
41 |     +-- DeprecationWarning
42 |     +-- PendingDeprecationWarning
43 |     +-- RuntimeWarning
44 |     +-- SyntaxWarning
45 |     +-- UserWarning
46 |     +-- FutureWarning
47 |     +-- ImportWarning
48 |     +-- UnicodeWarning

```

For a detailed description of each of these, refer to the official Python documentation.

Any object can be raised as an exception, but it is good practice to raise objects that extend one of the built-in exception classes.

2.3.5 `def...return`

Functions are declared using `def`. Here is a typical Python function:

```

1 >>> def f(a, b):
2 ...     return a + b
3 >>> print f(4, 2)
4 6

```

There is no need (or way) to specify the type of an argument(s) or the

44 ANNOTATED ALGORITHMS IN PYTHON

return value(s). In this example, a function `f` is defined that can take two arguments.

Functions are the first code syntax feature described in this chapter to introduce the concept of *scope*, or *namespace*. In the preceding example, the identifiers `a` and `b` are undefined outside of the scope of function `f`:

```
1 >>> def f(a):
2 ...     return a + 1
3 >>> print f(1)
4 2
5 >>> print a
6 Traceback (most recent call last):
7   File "<pyshell#22>", line 1, in <module>
8     print a
9 NameError: name 'a' is not defined
```

Identifiers defined outside of the function scope are accessible within the function; observe how the identifier `a` is handled in the following code:

```
1 >>> a = 1
2 >>> def f(b):
3 ...     return a + b
4 >>> print f(1)
5 2
6 >>> a = 2
7 >>> print f(1) # new value of a is used
8 3
9 >>> a = 1 # reset a
10 >>> def g(b):
11 ...     a = 2 # creates a new local a
12 ...     return a + b
13 >>> print g(2)
14 4
15 >>> print a # global a is unchanged
16 1
```

If `a` is modified, subsequent function calls will use the new value of the global `a` because the function definition binds the storage location of the identifier `a`, not the value of `a` itself at the time of function declaration; however, if `a` is assigned-to inside function `g`, the global `a` is unaffected because the new local `a` hides the global value. The external-scope reference can be used in the creation of *closures*:

```
1 >>> def f(x):
2 ...     def g(y):
3 ...         return x * y
```

```

4     ...     return g
5 >>> doubler = f(2) # doubler is a new function
6 >>> tripler = f(3) # tripler is a new function
7 >>> quadrupler = f(4) # quadrupler is a new function
8 >>> print doubler(5)
9 10
10 >>> print tripler(5)
11 15
12 >>> print quadrupler(5)
13 20

```

Function `f` creates new functions; note that the scope of the name `g` is entirely internal to `f`. Closures are extremely powerful.

Function arguments can have default values and can return multiple results as a tuple (notice the parentheses are optional and are omitted in the example):

```

1 >>> def f(a, b=2):
2     ...     return a + b, a - b
3 >>> x, y = f(5)
4 >>> print x
5 7
6 >>> print y
7 3

```

Function arguments can be passed explicitly by name; therefore the order of arguments specified in the caller can be different than the order of arguments with which the function was defined:

```

1 >>> def f(a, b=2):
2     ...     return a + b, a - b
3 >>> x, y = f(b=5, a=2)
4 >>> print x
5 7
6 >>> print y
7 -3

```

Functions can also take a runtime-variable number of arguments. Parameters that start with `*` and `**` must be the last two parameters. If the `**` parameter is used, it must be last in the list. Extra values passed in will be placed in the `*identifier` parameter, whereas named values will be placed into the `**identifier`. Notice that when passing values into the function, the unnamed values must be before any and all named values:

```

1 >>> def f(a, b, *extra, **extraNamed):
2     ...     print "a = ", a

```

```

3 ...     print "b = ", b
4 ...     print "extra = ", extra
5 ...     print "extranamed = ", extraNamed
6 >>> f(1, 2, 5, 6, x=3, y=2, z=6)
7 a = 1
8 b = 2
9 extra = (5, 6)
10 extranamed = {'y': 2, 'x': 3, 'z': 6}

```

Here the first two parameters (1 and 2) are matched with the parameters `a` and `b`, while the tuple `5, 6` is placed into `extra` and the remaining items (which are in a dictionary format) are placed into `extraNamed`.

In the opposite case, a list or tuple can be passed to a function that requires individual positional arguments by unpacking them:

```

1 >>> def f(a, b):
2 ...     return a + b
3 >>> c = (1, 2)
4 >>> print f(*c)
5 3

```

and a dictionary can be unpacked to deliver keyword arguments:

```

1 >>> def f(a, b):
2 ...     return a + b
3 >>> c = {'a':1, 'b':2}
4 >>> print f(**c)
5 3

```

2.3.6 lambda

The keyword `lambda` provides a way to define a short unnamed function:

```

1 >>> a = lambda b: b + 2
2 >>> print a(3)
3 5

```

The expression “`lambda [a]:[b]`” literally reads as “a function with arguments `[a]` that returns `[b]`.” The `lambda` expression is itself unnamed, but the function acquires a name by being assigned to identifier `a`. The scoping rules for `def` apply to `lambda` equally, and in fact, the preceding code, with respect to `a`, is identical to the function declaration using `def`:

```

1 >>> def a(b):
2 ...     return b + 2

```

```

3 >>> print a(3)
4 5

```

The only benefit of `lambda` is brevity; however, brevity can be very convenient in certain situations. Consider a function called `map` that applies a function to all items in a list, creating a new list:

```

1 >>> a = [1, 7, 2, 5, 4, 8]
2 >>> map(lambda x: x + 2, a)
3 [3, 9, 4, 7, 6, 10]

```

This code would have doubled in size had `def` been used instead of `lambda`. The main drawback of `lambda` is that (in the Python implementation) the syntax allows only for a single expression; however, for longer functions, `def` can be used, and the extra cost of providing a function name decreases as the length of the function grows.

Just like `def`, `lambda` can be used to *curry* functions: new functions can be created by wrapping existing functions such that the new function carries a different set of arguments:

```

1 >>> def f(a, b): return a + b
2 >>> g = lambda a: f(a, 3)
3 >>> g(2)
4 5

```

Python functions created with either `def` or `lambda` allow refactoring of existing functions in terms of a different set of arguments.

2.4 Classes

Because Python is dynamically typed, Python classes and objects may seem odd. In fact, member variables (attributes) do not need to be specifically defined when declaring a class, and different instances of the same class can have different attributes. Attributes are generally associated with the instance, not the class (except when declared as “class attributes,” which is the same as “static member variables” in C++/Java).

Here is an example:

```

1 >>> class MyClass(object): pass
2 >>> myinstance = MyClass()

```

```

3 >>> myinstance.myvariable = 3
4 >>> print myinstance.myvariable
5 3

```

Notice that `pass` is a do-nothing command. In this case, it is used to define a class `MyClass` that contains nothing. `MyClass()` calls the constructor of the class (in this case, the default constructor) and returns an object, an instance of the class. The `(object)` in the class definition indicates that our class extends the built-in object class. This is not required, but it is good practice.

Here is a more involved class with multiple methods:

```

1 >>> class Complex(object):
2 ...     z = 2
3 ...     def __init__(self, real=0.0, imag=0.0):
4 ...         self.real, self.imag = real, imag
5 ...     def magnitude(self):
6 ...         return (self.real**2 + self.imag**2)**0.5
7 ...     def __add__(self, other):
8 ...         return Complex(self.real+other.real, self.imag+other.imag)
9 >>> a = Complex(1,3)
10 >>> b = Complex(2,1)
11 >>> c = a + b
12 >>> print c.magnitude()
13 5

```

Functions declared inside the class are methods. Some methods have special reserved names. For example, `__init__` is the constructor. In the example, we created a class to store the `real` and the `imag` part of a complex number. The constructor takes these two variables and stores them into `self` (not a keyword but a variable that plays the same role as `this` in Java and `(*this)` in C++; this syntax is necessary to avoid ambiguity when declaring nested classes, such as a class that is local to a method inside another class, something Python allows but Java and C++ do not).

The `self` variable is defined by the first argument of each method. They all must have it, but they can use another variable name. Even if we use another name, the first argument of a method always refers to the object calling the method. It plays the same role as the `this` keyword in Java and C++.

Method `__add__` is also a special method (all special methods start and

end in double underscore) and it overloads the + operator between `self` and `other`. In the example, `a+b` is equivalent to a call to `a.__add__(b)`, and the `__add__` method receives `self=a` and `other=b`.

All variables are local variables of the method, except variables declared outside methods, which are called *class variables*, equivalent to C++ *static member variables*, which hold the same value for all instances of the class.

2.4.1 Special methods and operator overloading

Class attributes, methods, and operators starting with a double underscore are usually intended to be private (e.g., to be used internally but not exposed outside the class), although this is a convention that is not enforced by the interpreter.

Some of them are reserved keywords and have a special meaning:

- `__len__`
- `__getitem__`
- `__setitem__`

They can be used, for example, to create a container object that acts like a list:

```

1 >>> class MyList(object):
2 >>>     def __init__(self, *a): self.a = list(a)
3 >>>     def __len__(self): return len(self.a)
4 >>>     def __getitem__(self, key): return self.a[key]
5 >>>     def __setitem__(self, key, value): self.a[key] = value
6 >>> b = MyList(3, 4, 5)
7 >>> print b[1]
8 4
9 >>> b.a[1] = 7
10 >>> print b.a
11 [3, 7, 5]
```

Other special operators include `__getattr__` and `__setattr__`, which define the get and set methods (getters and setters) for the class, and `__add__`, `__sub__`, `__mul__`, and `__div__`, which overload arithmetic operators. For the use of these operators, we refer the reader to the chapter on linear

algebra, where they will be used to implement algebra for matrices.

2.4.2 class Financial Transaction

As one more example of a class, we implement a class that represents a financial transaction. We can think of a simple transaction as a single money transfer of quantity a that occurs at a given time t . We adopt the convention that a positive amount represents money flowing in and a negative value represents money flowing out.

The present value (computed at time t_0) for a transaction occurring at time t days from now of amount A is defined as

$$PV(t, A) = Ae^{-tr} \quad (2.2)$$

where r is the daily risk-free interest rate. If t is measured in days, r has to be the daily risk-free return. Here we will assume it defaults to $r = 0.05/365$ (5% annually).

Here is a possible implementation of the transaction:

```

1 from datetime import date
2 from math import exp
3 today = date.today()
4 r_free = 0.05/365.0
5
6 class FinancialTransaction(object):
7     def __init__(self, t, a, description=''):
8         self.t = t
9         self.a = a
10        self.description = description
11    def pv(self, t0=today, r=r_free):
12        return self.a*exp(r*(t0-self.t).days)
13    def __str__(self):
14        return '%.2f dollars in %i days (%s)' % \
15            (self.a, self.t, self.description)

```

Here we assume t and t_0 are `datetime.date` objects that store a date. The date constructor takes the year, the month, and the day separated by a comma. The expression `(t0-t).days` computes the distance in days between t_0 and t .

Similarly, we can implement a *Cash Flow* class to store a list of transactions,

with the `add` method to add a new transaction to the list. The present value of a cash flow is the sum of the present values of each transaction:

```

1 class CashFlow(object):
2     def __init__(self):
3         self.transactions = []
4     def add(self, transaction):
5         self.transactions.append(transaction)
6     def pv(self, t0, r=r_free):
7         return sum(x.pv(t0, r) for x in self.transactions)
8     def __str__(self):
9         return '\n'.join(str(x) for x in self.transactions)

```

What is the net present value at the beginning of 2012 for a bond that pays \$1000 the 20th of each month for the following 24 months (assuming a fixed interest rate of 5% per year)?

```

1 >>> bond = CashFlow()
2 >>> today = date(2012, 1, 1)
3 >>> for year in range(2012, 2014):
4 ...     for month in range(1, 13):
5 ...         coupon = FinancialTransaction(date(year, month, 20), 1000)
6 ...         bond.add(coupon)
7 >>> print round(bond.pv(today, r=0.05/365), 0)
8 22826

```

This means the cost for this bond should be \$22,826.

2.5 File input/output

In Python, you can open and write in a file with

```

1 >>> file = open('myfile.txt', 'w')
2 >>> file.write('hello world')
3 >>> file.close()

```

Similarly, you can read back from the file with

```

1 >>> file = open('myfile.txt', 'r')
2 >>> print file.read()
3 hello world

```

Alternatively, you can read in binary mode with “rb,” write in binary mode with “wb,” and open the file in append mode “a” using standard C notation.

The `read` command takes an optional argument, which is the number of

bytes. You can also jump to any location in a file using `seek` :

You can read back from the file with `read`:

```
1 >>> print file.seek(6)
2 >>> print file.read()
3 world
```

and you can close the file with:

```
1 >>> file.close()
```

2.6 How to import modules

The real power of Python is in its library modules. They provide a large and consistent set of application programming interfaces (APIs) to many system libraries (often in a way independent of the operating system).

For example, if you need to use a random number generator, you can do the following:

```
1 >>> import random
2 >>> print random.randint(0, 9)
3 5
```

This prints a random integer in the range of $[0,9]$, 5 in the example. The function `randint` is defined in the module `random`. It is also possible to import an object from a module into the current namespace:

```
1 >>> from random import randint
2 >>> print randint(0, 9)
```

or import all objects from a module into the current namespace:

```
1 >>> from random import *
2 >>> print randint(0, 9)
```

or import everything in a newly defined namespace:

```
1 >>> import random as myrand
2 >>> print myrand.randint(0, 9)
```

In the rest of this book, we will mainly use objects defined in modules `math`, `cmath`, `os`, `sys`, `datetime`, `time`, and `cPickle`. We will also use the `random` module, but we will describe it in a later chapter.

In the following subsections, we consider those modules that are most

useful.

2.6.1 `math` and `cmath`

Here is a sampling of some of the methods available in the `math` and `cmath` packages:

- `math.isinf(x)` returns true if the floating point number `x` is positive or negative infinity
- `math.isnan(x)` returns true if the floating point number `x` is NaN; see Python documentation or IEEE 754 standards for more information
- `math.exp(x)` returns `e**x`
- `math.log(x[, base])` returns the logarithm of `x` to the optional base; if base is not supplied, `e` is assumed
- `math.cos(x)`, `math.sin(x)`, `math.tan(x)` returns the cos, sin, tan of the value of `x`; `x` is in radians
- `math.pi`, `math.e` are the constants for `pi` and `e` to available precision

2.6.2 `os`

This module provides an interface for the operating system API:

```
1 >>> import os
2 >>> os.chdir('.')
3 >>> os.unlink('filename_to_be_deleted')
```

Some of the `os` functions, such as `chdir`, are not thread safe, for example, they should not be used in a multithreaded environment.

`os.path.join` is very useful; it allows the concatenation of paths in an OS-independent way:

```
1 >>> import os
2 >>> a = os.path.join('path', 'sub_path')
3 >>> print a
4 path/sub_path
```

System environment variables can be accessed via

```
1 >>> print os.environ
```

which is a read-only dictionary.

2.6.3 sys

The sys module contains many variables and functions, but used the most is sys.path. It contains a list of paths where Python searches for modules. When we try to import a module, Python searches the folders listed in sys.path. If you install additional modules in some location and want Python to find them, you need to append the path to that location to sys.path:

```
1 >>> import sys
2 >>> sys.path.append('path/to/my/modules')
```

2.6.4 datetime

The use of the datetime module is best illustrated by some examples:

```
1 >>> import datetime
2 >>> print datetime.datetime.today()
3 2008-07-04 14:03:90
4 >>> print datetime.date.today()
5 2008-07-04
```

Occasionally you may need to time stamp data based on the UTC time as opposed to local time. In this case, you can use the following function:

```
1 >>> import datetime
2 >>> print datetime.datetime.utcnow()
3 2008-07-04 14:03:90
```

The datetime module contains various classes: date, datetime, time, and timedelta. The difference between two dates or two datetimes or two time objects is a timedelta:

```
1 >>> a = datetime.datetime(2008, 1, 1, 20, 30)
2 >>> b = datetime.datetime(2008, 1, 2, 20, 30)
3 >>> c = b - a
4 >>> print c.days
5 1
```

We can also parse dates and datetimes from strings:

```

1 >>> s = '2011-12-31'
2 >>> a = datetime.datetime.strptime(s, '%Y-%m-%d') #modified
3 >>> print a.year, a.day, a.month
4 2011 31 12 #modified

```

Notice that “%Y” matches the four-digit year, “%m” matches the month as a number (1–12), “%d” matches the day (1–31), “%H” matches the hour, “%M” matches the minute, and “%S” matches the seconds. Check the Python documentation for more options.

2.6.5 time

The time module differs from date and datetime because it represents time as seconds from the epoch (beginning of 1970):

```

1 >>> import time
2 >>> t = time.time()
3 1215138737.571

```

Refer to the Python documentation for conversion functions between time in seconds and time as a datetime.

2.6.6 urllib and json

The urllib is a module to download data or a web page from a URL:

```

1 >>> import urllib
2 >>> page = urllib.urlopen('http://www.google.com/')
3 >>> html = page.read()

```

Usually urllib is used to download data posted online. The challenge may be parsing the data (converting from the representation used to post it to a proper Python representation).

In the following, we create a simple helper class that can download data from Yahoo! Finance and Google Finance and convert each stock’s historical data into a list of dictionaries. Each list element corresponds to a trading day of history of the stock, and each dictionary stores the data relative to that trading day (date, open, close, volume, adjusted close, arithmetic_return, log_return, etc.):

Listing 2.1: in file: nlib.py

```

1 class YStock:
2     """
3     Class that downloads and stores data from Yahoo Finance
4     Examples:
5     >>> google = YStock('GOOG')
6     >>> current = google.current()
7     >>> price = current['price']
8     >>> market_cap = current['market_cap']
9     >>> h = google.historical()
10    >>> last_adjusted_close = h[-1]['adjusted_close']
11    >>> last_log_return = h[-1]['log_return']
12    previous version of this code user Yahoo for historical data
13    but Yahoo changed API and blocked them, moving to Google finance.
14    """
15    URL_CURRENT = 'http://finance.yahoo.com/d/quotes.csv?s=%(symbol)s&f=%(
        columns)s'
16    URL_HISTORICAL = 'https://www.google.com/finance/historical?output=csv&q=%(
        symbol)s'
17
18    def __init__(self, symbol):
19        self.symbol = symbol.upper()
20
21    def current(self):
22        import urllib
23        FIELDS = (('price', 'l1'),
24                  ('change', 'c1'),
25                  ('volume', 'v'),
26                  ('average_daily_volume', 'a2'),
27                  ('stock_exchange', 'x'),
28                  ('market_cap', 'j1'),
29                  ('book_value', 'b4'),
30                  ('ebitda', 'j4'),
31                  ('dividend_per_share', 'd'),
32                  ('dividend_yield', 'y'),
33                  ('earnings_per_share', 'e'),
34                  ('52_week_high', 'k'),
35                  ('52_week_low', 'j'),
36                  ('50_days_moving_average', 'm3'),
37                  ('200_days_moving_average', 'm4'),
38                  ('price_earnings_ratio', 'r'),
39                  ('price_earnings_growth_ratio', 'r5'),
40                  ('price_sales_ratio', 'p5'),
41                  ('price_book_ratio', 'p6'),
42                  ('short_ratio', 's7'))
43        columns = ''.join([row[1] for row in FIELDS])
44        url = self.URL_CURRENT % dict(symbol=self.symbol, columns=columns)
45        raw_data = urllib.urlopen(url).read().strip().strip('').split(',')

```



```

46     current = dict()
47     for i,row in enumerate(FIELDS):
48         try:
49             current[row[0]] = float(raw_data[i])
50         except:
51             current[row[0]] = raw_data[i]
52     return current
53
54 def historical(self, start=None, stop=None):
55     import datetime, time, urllib, math
56     url = self.URL_HISTORICAL % dict(symbol=self.symbol)
57     # Date,Open,High,Low,Close,Volume,Adj Close
58     lines = urllib.urlopen(url).readlines()
59     if any('CAPTCHA' in line for line in lines):
60         print url
61         raise
62     raw_data = [row.split(',') for row in lines[1:] if 5 <= row.count(',')
63                 <= 6]
64     previous_adjusted_close = 0
65     series = []
66     raw_data.reverse()
67     for row in raw_data:
68         if row[1] == '-': continue
69         date = datetime.datetime.strptime(row[0], '%d-%b-%y')
70         if (start and date<start) or (stop and date>stop): continue
71         open, high, low = float(row[1]), float(row[2]), float(row[3])
72         close, vol = float(row[4]), float(row[5])
73         adjusted_close = float(row[5]) if len(row)>5 else close
74         adjustment = adjusted_close/close
75         if previous_adjusted_close:
76             arithmetic_return = adjusted_close/previous_adjusted_close-1.0
77
78             log_return = math.log(adjusted_close/previous_adjusted_close)
79         else:
80             arithmetic_return = log_return = None
81             previous_adjusted_close = adjusted_close
82             series.append(dict(
83                 date = date,
84                 open = open,
85                 high = high,
86                 low = low,
87                 close = close,
88                 volume = vol,
89                 adjusted_close = adjusted_close,
90                 adjusted_open = open*adjustment,
91                 adjusted_high = high*adjustment,
92                 adjusted_low = low*adjustment,
93                 adjusted_vol = vol/adjustment,
94                 arithmetic_return = arithmetic_return,

```

```

94         log_return = log_return))
95     return series
96
97     @staticmethod
98     def download(symbol='goog', what='adjusted_close', start=None, stop=None):
99         return [d[what] for d in YStock(symbol).historical(start, stop)]

```

Many web services return data in JSON format. JSON is slowly replacing XML as a favorite protocol for data transfer on the web. It is lighter, simpler to use, and more human readable. JSON can be thought of as serialized JavaScript. the JSON data can be converted to a Python object using a library called json:

```

1  >>> import json
2  >>> a = [1,2,3]
3  >>> b = json.dumps(a)
4  >>> print type(b)
5  <type 'str'>
6  >>> c = json.loads(b)
7  >>> a == c
8  True

```

The module json has loads and dumps methods which work very much as cPickle's methods, but they serialize the objects into a string using JSON instead of the pickle protocol.

2.6.7 pickle

This is a very powerful module. It provides functions that can serialize almost any Python object, including self-referential objects. For example, let's build a weird object:

```

1  >>> class MyClass(object): pass
2  >>> myinstance = MyClass()
3  >>> myinstance.x = 'something'
4  >>> a = [1, 2, {'hello': 'world'}, [3, 4, [myinstance]]]

```

and now:

```

1  >>> import cPickle as pickle
2  >>> b = pickle.dumps(a)
3  >>> c = pickle.loads(b)

```

In this example, b is a string representation of a, and c is a copy of a generated by deserializing b. The module pickle can also serialize to and

deserialize from a file:

```
1 >>> pickle.dump(a, open('myfile.pickle', 'wb'))
2 >>> c = pickle.load(open('myfile.pickle', 'rb'))
```

2.6.8 sqlite

The Python dictionary type is very useful, but it lacks persistence because it is stored in RAM (it is lost if a program ends) and cannot be shared by more than one process running concurrently. Moreover, it is not transaction safe. This means that it is not possible to group operations together so that they succeed or fail as one.

Think for example of using the dictionary to store a bank account. The key is the account number and the value is a list of transactions. We want the dictionary to be safely stored on file. We want it to be accessible by multiple processes and applications. We want transaction safety: it should not be possible for an application to fail during a money transfer, resulting in the disappearance of money.

Python provides a module called `shelve` with the same interface as `dict`, which is stored on disk instead of in RAM. One problem with this module is that the file is not locked when accessed. If two processes try to access it concurrently, the data becomes corrupted. This module also does not provide transactional safety.

The proper alternative consists of using a database. There are two types of databases: relational databases (which normally use SQL syntax) and non-relational databases (often referred to as NoSQL). Key-value persistent storage databases usually follow under the latter category. Relational databases excel at storing structured data (in the form of tables), establishing relations between rows of those tables, and searches involving multiple tables linked by references. NoSQL databases excel at storing and retrieving schemaless data and replication of data (redundancy for fail safety).

Python comes with an embedded SQL database called SQLite [9]. All data in the database are stored in one single file. It supports the SQL query

language and transactional safety. It is very fast and allows concurrent read (from multiple processes), although not concurrent write (the file is locked when a process is writing to the file until the transaction is committed). Concurrent write requests are queued and executed in order when the database is unlocked.

Installing and using any of these database systems is beyond the scope of this book and not necessary for our purposes. In particular, we are not concerned with relations, data replications, and speed.

As an exercise, we are going to implement a new Python class called `PersistentDictionary` that exposes an interface similar to a dict but uses the SQLite database for storage. The database file is created if it does not exist. `PersistentDictionary` will use a single table (also called persistence) to store rows containing a key (pkey) and a value (pvalue).

For later convenience, we will also add a method that can generate a UUID key. A UUID is a random string that is long enough to be, most likely, unique. This means that two calls to the same function will return different values, and the probability that the two values will be the same is negligible. Python includes a library to generate UUID strings based on a common industry standard. We use the function `uuid4`, which also uses the time and the IP of the machine to generate the UUID. This means the UUID is unlikely to have conflicts with (be equal to) another UUID generated on other machines. The `uuid` method will be useful to generate random unique keys.

We will also add a method that allows us to search for keys in the database using GLOB patterns (in a GLOB pattern, `***` represents a generic wildcard and `?` is a single-character wildcard).

Here is the code:

Listing 2.2: in file: `nlib.py`

```

1 import os
2 import uuid
3 import sqlite3
4 import cPickle as pickle
5 import unittest
6

```

```

7 class PersistentDictionary(object):
8     """
9     A sqlite based key,value storage.
10    The value can be any pickleable object.
11    Similar interface to Python dict
12    Supports the GLOB syntax in methods keys(),items(), __delitem__()
13
14    Usage Example:
15    >>> p = PersistentDictionary(path='test.sqlite')
16    >>> key = 'test/' + p.uuid()
17    >>> p[key] = {'a': 1, 'b': 2}
18    >>> print p[key]
19    {'a': 1, 'b': 2}
20    >>> print len(p.keys('test/*'))
21    1
22    >>> del p[key]
23    """
24
25    CREATE_TABLE = "CREATE TABLE persistence (pkey, pvalue)"
26    SELECT_KEYS = "SELECT pkey FROM persistence WHERE pkey GLOB ?"
27    SELECT_VALUE = "SELECT pvalue FROM persistence WHERE pkey GLOB ?"
28    INSERT_KEY_VALUE = "INSERT INTO persistence(pkey, pvalue) VALUES (?,?)"
29    UPDATE_KEY_VALUE = "UPDATE persistence SET pvalue = ? WHERE pkey = ?"
30    DELETE_KEY_VALUE = "DELETE FROM persistence WHERE pkey LIKE ?"
31    SELECT_KEY_VALUE = "SELECT pkey,pvalue FROM persistence WHERE pkey GLOB ?"
32
33    def __init__(self,
34                  path='persistence.sqlite',
35                  autocommit=True,
36                  serializer=pickle):
37        self.path = path
38        self.autocommit = autocommit
39        self.serializer = serializer
40        create_table = not os.path.exists(path)
41        self.connection = sqlite3.connect(path)
42        self.connection.text_factory = str # do not use unicode
43        self.cursor = self.connection.cursor()
44        if create_table:
45            self.cursor.execute(self.CREATE_TABLE)
46            self.connection.commit()
47
48    def uuid(self):
49        return str(uuid4())
50
51    def keys(self,pattern='*'):
52        "returns a list of keys filtered by a pattern, * is the wildcard"
53        self.cursor.execute(self.SELECT_KEYS,(pattern,))
54        return [row[0] for row in self.cursor.fetchall()]
55

```

62 ANNOTATED ALGORITHMS IN PYTHON

```
56 def __contains__(self, key):
57     return True if self.get(key) != None else False
58
59 def __iter__(self):
60     for key in self:
61         yield key
62
63 def __setitem__(self, key, value):
64     if key in self:
65         if value is None:
66             del self[key]
67         else:
68             svalue = self.serializer.dumps(value)
69             self.cursor.execute(self.UPDATE_KEY_VALUE, (svalue, key))
70     else:
71         svalue = self.serializer.dumps(value)
72         self.cursor.execute(self.INSERT_KEY_VALUE, (key, svalue))
73     if self.autocommit: self.connection.commit()
74
75 def get(self, key):
76     self.cursor.execute(self.SELECT_VALUE, (key,))
77     row = self.cursor.fetchone()
78     return self.serializer.loads(row[0]) if row else None
79
80 def __getitem__(self, key):
81     self.cursor.execute(self.SELECT_VALUE, (key,))
82     row = self.cursor.fetchone()
83     if not row: raise KeyError
84     return self.serializer.loads(row[0])
85
86 def __delitem__(self, pattern):
87     self.cursor.execute(self.DELETE_KEY_VALUE, (pattern,))
88     if self.autocommit: self.connection.commit()
89
90 def items(self, pattern='*'):
91     self.cursor.execute(self.SELECT_KEY_VALUE, (pattern,))
92     return [(row[0], self.serializer.loads(row[1])) \
93             for row in self.cursor.fetchall()]
94
95 def dumps(self, pattern='*'):
96     self.cursor.execute(self.SELECT_KEY_VALUE, (pattern,))
97     rows = self.cursor.fetchall()
98     return self.serializer.dumps(dict((row[0], self.serializer.loads(row[1]))
99                                     for row in rows))
100
101 def loads(self, raw):
102     data = self.serializer.loads(raw)
103     for key, value in data.iteritems():
```

```
104 self[key] = value
```

This code now allows us to do the following:

- Create a persistent dictionary:

```
1 >>> p = PersistentDictionary(path='storage.sqlite',autocommit=False)
```

- Store data in it:

```
1 >>> p['some/key'] = 'some value'
```

where “some/key” must be a string and “some value” can be any Python pickleable object.

- Generate a UUID to be used as the key:

```
1 >>> key = p.uuid()
2 >>> p[key] = 'some other value'
```

- Retrieve the data:

```
1 >>> data = p['some/key']
```

- Loop over keys:

```
1 >>> for key in p: print key, p[key]
```

- List all keys:

```
1 >>> keys = p.keys()
```

- List all keys matching a pattern:

```
1 >>> keys = p.keys('some/*')
```

- List all key-value pairs matching a pattern:

```
1 >>> for key,value in p.items('some/*'): print key, value
```

- Delete keys matching a pattern:

```
1 >>> del p['some/*']
```

We will now use our persistence storage to download 2011 financial data from the SP100 stocks. This will allow us to later perform various analysis tasks on these stocks:

Listing 2.3: in file: `nlib.py`

```
1 >>> SP100 = ['AA', 'AAPL', 'ABT', 'AEP', 'ALL', 'AMGN', 'AMZN', 'AVP',
2 ... 'AXP', 'BA', 'BAC', 'BAX', 'BHI', 'BK', 'BMY', 'BRK.B', 'CAT', 'C', 'CL',
3 ... 'CMCSA', 'COF', 'COP', 'COST', 'CPB', 'CSCO', 'CVS', 'CVX', 'DD', 'DELL',
4 ... 'DIS', 'DOW', 'DVN', 'EMC', 'ETR', 'EXC', 'F', 'FCX', 'FDX', 'GD', 'GE',
5 ... 'GILD', 'GOOG', 'GS', 'HAL', 'HD', 'HNZ', 'HON', 'HPQ', 'IBM', 'INTC',
```

64 ANNOTATED ALGORITHMS IN PYTHON

```
6 ... 'JNJ', 'JPM', 'KFT', 'KO', 'LMT', 'LOW', 'MA', 'MCD', 'MDT', 'MET',
7 ... 'MMM', 'MO', 'MON', 'MRK', 'MS', 'MSFT', 'NKE', 'NOV', 'NSC', 'NWSA',
8 ... 'NYX', 'ORCL', 'OXY', 'PEP', 'PFE', 'PG', 'PM', 'QCOM', 'RF', 'RTN', 'S',
9 ... 'SLB', 'SLE', 'SO', 'T', 'TGT', 'TWX', 'TXN', 'UNH', 'UPS', 'USB',
10 ... 'UTX', 'VZ', 'WAG', 'WFC', 'WMB', 'WMT', 'WY', 'XOM', 'XRX']
11 >>> from datetime import date
12 >>> storage = PersistentDictionary('sp100.sqlite')
13 >>> for symbol in SP100:
14 ...     key = symbol+'/2011'
15 ...     if not key in storage:
16 ...         storage[key] = YStock(symbol).historical(start=date(2011,1,1),
17 ...                                                    stop=date(2011,12,31))
```

Notice that while storing one item may be slower than storing an individual item in its own files, accessing the file system becomes progressively slower as the number of files increases. Storing data in a database, long term, is a winning strategy as it scales better and it is easier to search for and extract data than it is with multiple flat files. Which type of database is most appropriate depends on the type of data and the type of queries we need to perform on the data.

2.6.9 numpy

The library `numpy` [2] is the Python library for efficient arrays, multidimensional arrays, and their manipulation. `numpy` does not ship with Python and must be installed separately.

On most platforms, this is as easy as typing in the Bash Shell:

```
1 pip install numpy
```

Yet on other platforms, it can be a more lengthy process, and we leave it to the reader to find the best installation procedure.

The basic object in `numpy` is the `ndarray` (n -dimensional array). Here we make a $10 \times 4 \times 3$ array of 64 bits float:

```
1 >>> import numpy
2 >>> a = numpy.ndarray((10,4,3), dtype=numpy.float64)
```

The class `ndarray` is more efficient than Python's list. It takes much less space because their elements have a fixed given type (e.g., `float64`). Other popular available types are: `int8`, `int16`, `int32`, `int64`, `uint8`, `uint16`, `uint32`,

uint64, float16, float32, float64, complex64, and complex128.

We can access elements:

```
1 >>> a[0,0,0] = 1
2 >>> print a[0,0,0]
3 1.0
```

We can query for its size:

```
1 >>> print a.shape
2 (10, 4, 3)
```

We can reshape its elements:

```
1 >>> b = a.reshape((10,12))
2 >>> print a.shape
3 (10, 12)
```

We can map one type into another

```
1 >>> c = b.astype(float32)
```

We can load and save them:

```
1 >>> numpy.save('array.npy',a)
2 >>> b = numpy.load('array.npy')
```

And we can perform operations on them (most operations are element-wise operations):

```
1 >>> a = numpy.array([[1,2],[3,4]]) # converts a list into a ndarray
2 >>> print a
3 [[1 2]
4  [3 4]]
5 >>> print a+1
6 [[2 3]
7  [4 5]]
8 >>> print a+a
9 [[2 4]
10 [6 8]]
11 >>> print a*2
12 [[2 4]
13 [6 8]]
14 >>> print a*a
15 [[ 1  4]
16  [ 9 16]]
17 >>> print numpy.exp(a)
18 [[ 2.71828183  7.3890561 ]
19  [20.08553692 54.59815003]]
```

The numpy module also implements common linear algebra operations:

66 ANNOTATED ALGORITHMS IN PYTHON

```
1 >>> from numpy import dot
2 >>> from numpy.linalg import inv
3 >>> print dot(a,a)
4 [[ 7 10]
5  [15 22]]
6 >>> print inv(a)
7 [[-2.   1. ]
8  [ 1.5 -0.5]]
```

These operations are particularly efficient because they are implemented on top of the BLAS and LAPACK libraries.

There are many other functions in the `numpy` module, and you can read more about it in the official documentation.

2.6.10 matplotlib

Library `matplotlib` [10] is the de facto standard plotting library for Python. It is one of the best and most versatile plotting libraries available. It has two modes of operation. One mode of operation, called `pylab`, follows a MATLAB-like syntax. The other mode follows a more Python-style syntax. Here we use the latter.

You can install `matplotlib` with

```
1 pip install matplotlib
```

and it requires `numpy`. In `matplotlib`, we need to distinguish the following objects:

- **Figure:** a blank grid that can contain pairs of *XY* axes
- **Axes:** a pair of *XY* axes that may contain multiple superimposed plots
- **FigureCanvas:** a binary representation of a figure with everything that it contains
- **plot:** a representation of a data set such as a line plot or a scatter plot

In `matplotlib`, a canvas can be visualized in a window or serialized into an image file. Here we take the latter approach and create two helper functions that take data and configuration parameters and output PNG images.

We start by importing matplotlib and other required libraries:

Listing 2.4: in file: nlib.py

```

1 import math
2 import cmath
3 import random
4 import os
5 import tempfile
6 os.environ['MPLCONFIGDIR'] = tempfile.mkdtemp()

```

Now we define a helper that can plot lines, points with error bars, histograms, and scatter plots on a single canvas:

Listing 2.5: in file: nlib.py

```

1 from cStringIO import StringIO
2 try:
3     from matplotlib.figure import Figure
4     from matplotlib.backends.backend_agg import FigureCanvasAgg
5     from matplotlib.patches import Ellipse
6     HAVE_MATPLOTLIB = True
7 except ImportError:
8     HAVE_MATPLOTLIB = False
9
10 class Canvas(object):
11
12     def __init__(self, title='', xlab='x', ylab='y', xrange=None, yrange=None):
13         self.fig = Figure()
14         self.fig.set_facecolor('white')
15         self.ax = self.fig.add_subplot(111)
16         self.ax.set_title(title)
17         self.ax.set_xlabel(xlab)
18         self.ax.set_ylabel(ylab)
19         if xrange:
20             self.ax.set_xlim(xrange)
21         if yrange:
22             self.ax.set_ylim(yrange)
23         self.legend = []
24
25     def save(self, filename='plot.png'):
26         if self.legend:
27             legend = self.ax.legend([e[0] for e in self.legend],
28                                     [e[1] for e in self.legend])
29             legend.get_frame().set_alpha(0.7)
30         if filename:
31             FigureCanvasAgg(self.fig).print_png(open(filename, 'wb'))
32         else:
33             s = StringIO()
34             FigureCanvasAgg(self.fig).print_png(s)

```

68 ANNOTATED ALGORITHMS IN PYTHON

```
35         return s.getvalue()
36
37     def binary(self):
38         return self.save(None)
39
40     def hist(self, data, bins=20, color='blue', legend=None):
41         q = self.ax.hist(data, bins)
42         #if legend:
43         #    self.legend.append((q[0], legend))
44         return self
45
46     def plot(self, data, color='blue', style='-', width=2,
47              legend=None, xrange=None):
48         if callable(data) and xrange:
49             x = [xrange[0]+0.01*i*(xrange[1]-xrange[0]) for i in xrange(0,101)]
50             y = [data(p) for p in x]
51         elif data and isinstance(data[0],(int,float)):
52             x, y = xrange(len(data)), data
53         else:
54             x, y = [p[0] for p in data], [p[1] for p in data]
55         q = self.ax.plot(x, y, linestyle=style, linewidth=width, color=color)
56         if legend:
57             self.legend.append((q[0], legend))
58         return self
59
60     def errorbar(self, data, color='black', marker='o', width=2, legend=None):
61         x,y,dy = [p[0] for p in data], [p[1] for p in data], [p[2] for p in data
62         ]
63         q = self.ax.errorbar(x, y, yerr=dy, fmt=marker, linewidth=width, color=
64         color)
65         if legend:
66             self.legend.append((q[0], legend))
67         return self
68
69     def ellipses(self, data, color='blue', width=0.01, height=0.01, legend=None)
70     :
71         for point in data:
72             x, y = point[:2]
73             dx = point[2] if len(point)>2 else width
74             dy = point[3] if len(point)>3 else height
75             ellipse = Ellipse(xy=(x, y), width=dx, height=dy)
76             self.ax.add_artist(ellipse)
77             ellipse.set_clip_box(self.ax.bbox)
78             ellipse.set_alpha(0.5)
79             ellipse.set_facecolor(color)
80         if legend:
81             self.legend.append((q[0], legend))
82         return self
```

```

81 def imshow(self, data, interpolation='bilinear'):
82     self.ax.imshow(data).set_interpolation(interpolation)
83     return self

```

Notice we only make one set of axes.

The argument `111` of `figure.add_subplot(111)` indicates that we want a grid of 1×1 axes, and we ask for the first one of them (the only one).

The `linesets` parameter is a list of dictionaries. Each dictionary must have a “data” key corresponding to a list of (x, y) values. Each dictionary is rendered by a line connecting the points. It can have a “label,” a “color,” a “style,” and a “width.”

The `pointsets` parameter is a list of dictionaries. Each dictionary must have a “data” key corresponding to a list of $(x, y, \delta x, \delta y)$ values. Each dictionary is rendered by a set of circles with error bars. It can optionally have a “label,” a “color,” and a “marker” (symbol to replace the circle).

The `histsets` parameter is a list of dictionaries. Each dictionary must have a “data” key corresponding to a list of x values. Each dictionary is rendered by histogram. Each dictionary can optionally have a “label” and a “color.”

The `ellisets` parameter is also a list of dictionaries. Each dictionary must have a “data” key corresponding to a list of $(x, y, \delta x, \delta y)$ values. Each dictionary is rendered by a set of ellipses, one per point. It can optionally have a “color.”

We chose to draw all these types of plots with a single function because it is common to superimpose fitting lines to histograms, points, and scatter plots.

As an example, we can plot the adjusted closing price for AAPL:

Listing 2.6: in file: `nlib.py`

```

1 >>> storage = PersistentDictionary('sp100.sqlite')
2 >>> appl = storage['AAPL/2011']
3 >>> points = [(x, y['adjusted_close']) for (x, y) in enumerate(appl)]
4 >>> Canvas(title='Apple Stock (2011)', xlab='trading day', ylab='adjusted close').
    plot(points, legend='AAPL').save('images/aapl2011.png')

```

Here is an example of a histogram of daily arithmetic returns for the



Figure 2.1: Example of a line plot. Adjusted closing price for the AAPL stock in 2011 (source: Yahoo! Finance).

AAPL stock in 2011:

Listing 2.7: in file: nlib.py

```
1 >>> storage = PersistentDictionary('spl00.sqlite')
2 >>> appl = storage['AAPL/2011'][1:] # skip 1st day
3 >>> points = [day['arithmetic_return'] for day in appl]
4 >>> Canvas(title='Apple Stock (2011)', xlab='arithmetic return', ylab='frequency'
    ).hist(points).save('images/aapl2011hist.png')
```

Here is a scatter plot for random data points:

Listing 2.8: in file: nlib.py

```
1 >>> from random import gauss
2 >>> points = [(gauss(0,1),gauss(0,1),gauss(0,0.2),gauss(0,0.2)) for i in xrange
    (30)]
3 >>> Canvas(title='example scatter plot', xrange=(-2,2), yrange=(-2,2)).ellipses(
    points).save('images/scatter.png')
```

Here is a scatter plot showing the return and variance of the S&P100 stocks:




Figure 2.2: Example of a histogram plot. Distribution of daily arithmetic returns for the AAPL stock in 2011 (source: Yahoo! Finance).

Listing 2.9: in file: nlib.py

```

1 >>> storage = PersistentDictionary('sp100.sqlite')
2 >>> points = []
3 >>> for key in storage.keys('*/2011'):
4 ...     v = [day['log_return'] for day in storage[key][1:]]
5 ...     ret = sum(v)/len(v)
6 ...     var = sum(x**2 for x in v)/len(v) - ret**2
7 ...     points.append((var*math.sqrt(len(v)),ret*len(v),0.0002,0.02))
8 >>> Canvas(title='S&P100 (2011)',xlab='risk',ylab='return',
9 ...         xrange = (min(p[0] for p in points),max(p[0] for p in points)),
10 ...         yrange = (min(p[1] for p in points),max(p[1] for p in points))
11 ...         ).ellipses(points).save('images/sp100rr.png')

```

Notice the daily log returns have been multiplied by the number of days in one year to obtain the annual return. Similarly, the daily volatility has been multiplied by the square root of the number of days in one year to obtain the annual volatility (risk). The reason for this procedure will be explained in a later chapter.

Listing 2.10: in file: nlib.py

```

1 >>> def f(x,y): return (x-1)**2+(y-2)**2

```




Figure 2.3: Example of a scatter plot using some random points.

```
2 >>> points = [[f(0.1*i-3,0.1*j-3) for i in range(61)] for j in range(61)]
3 >>> Canvas(title='example 2d function').imshow(points).save('images/color2d.png')
    )
```

The class `Canvas` is both in `nlib.py` and in the Python module `canvas` [11].

2.6.11 ocl

One of the best features of Python is that it can introspect itself, and this can be used to just-in-time compile Python code into other languages. For example, the Cython [12] and the `ocl` libraries allow decorating Python code and converting it to C code. This makes the decorated functions much faster. Cython is more powerful, and it supports a richer subset of the Python syntax; `ocl` instead supports only a subset of the Python syntax, which can be directly mapped into the C equivalent, but it is easier to use. Moreover, `ocl` can convert Python code to JavaScript and to OpenCL (this is discussed in our last chapter).

Here is a simple example that implements the factorial function:




Figure 2.4: Example of a scatter plot. Risk-return plot for the S&P100 stocks in 2011 (source: Yahoo! Finance).

```

1 from ocl import Compiler
2 c99 = Compiler()
3
4 @c99.define(n='int')
5 def factorial(n):
6     output = 1
7     for k in xrange(1, n + 1):
8         output = output * k
9     return output
10 compiled = c99.compile()
11 print compiled.factorial(10)
12 assert compiled.factorial(10) == factorial(10)

```

The line `@c99.define(n='int')` instructs `ocl` that `factorial` must be converted to `c99` and that `n` is an integer. The `assert` command checks that `compiled.factorial(10)` produces the same output as `factorial(10)`, where the former runs compiled `c99` code, whereas the latter runs Python code.




Figure 2.5: Example of a two-dimensional color plot using for $f(x,y) = (x-1)^2 + (y-2)^2$.

3

Theory of Algorithms

An algorithm is a step-by-step procedure for solving a problem and is typically developed before doing any programming. The word comes from *algorism*, from the mathematician al-Khwarizmi, and was used to refer to the rules of performing arithmetic using Hindu–Arabic numerals and the systematic solution of equations.

In fact, algorithms are independent of any programming language. Efficient algorithms can have a dramatic effect on our problem-solving capabilities.

The basic steps of algorithms are loops (for, conditionals (if), and function calls. Algorithms also make use of arithmetic expressions, logical expressions (not, and, or), and expressions that can be reduced to the other basic components.

The issues that concern us when developing and analyzing algorithms are the following:

1. Correctness: of the problem specification, of the proposed algorithm, and of its implementation in some programming language (we will not worry about the third one; program verification is another subject altogether)
2. Amount of work done: for example, running time of the algorithm in terms of the input size (independent of hardware and programming

language)

3. Amount of space used: here we mean the amount of extra space (system resources) beyond the size of the input (independent of hardware and programming language); we will say that an algorithm is *in place* if the amount of extra space is constant with respect to input size
4. Simplicity, clarity: unfortunately, the simplest is not always the best in other ways
5. Optimality: can we prove that it does as well as or better than any other algorithm?

3.1 Order of growth of algorithms

The *insertion sort* is a simple algorithm in which an array of elements is sorted in place, one entry at a time. It is not the fastest sorting algorithm, but it is simple and does not require extra memory other than the memory needed to store the input array.

The insertion sort works by iterating. Every iteration i of the insertion sort removes one element from the input data and inserts it into the correct position in the already-sorted subarray $A[j]$ for $0 \leq j < i$. The algorithm iterates n times (where n is the total size of the input array) until no input elements remain to be sorted:

```

1 def insertion_sort(A):
2     for i in xrange(1,len(A)):
3         for j in xrange(i,0,-1):
4             if A[j]<A[j-1]:
5                 A[j], A[j-1] = A[j-1], A[j]
6             else: break

```

Here is an example:

```

1 >>> import random
2 >>> a=[random.randint(0,100) for k in xrange(20)]
3 >>> insertion_sort(a)
4 >>> print a
5 [6, 8, 9, 17, 30, 31, 45, 48, 49, 56, 56, 57, 65, 66, 75, 75, 82, 89, 90, 99]

```

One important question is, how long does this algorithm take to run?

How does its running time scale with the input size?

Given any algorithm, we can define three characteristic functions:

- $T_{worst}(n)$: the running time in the worst case
- $T_{best}(n)$: the running time in the best case
- $T_{average}(n)$: the running time in the average case

The best case for an insertion sort is realized when the input is already sorted. In this case, the inner for loop exits (breaks) always at the first iteration, thus only the most outer loop is important, and this is proportional to n ; therefore $T_{best}(n) \propto n$. The worst case for the insertion sort is realized when the input is sorted in reversed order. In this case, we can prove, and we do so subsequently, that $T_{worst}(n) \propto n^2$. For this algorithm, a statistical analysis shows that the worst case is also the average case.

Often we cannot determine exactly the running time function, but we may be able to set bounds to the running time.

We define the following sets:

- $O(g(n))$: the set of functions that grow no faster than $g(n)$ when $n \rightarrow \infty$
- $\Omega(g(n))$: the set of functions that grow no slower than $g(n)$ when $n \rightarrow \infty$
- $\Theta(g(n))$: the set of functions that grow at the same rate as $g(n)$ when $n \rightarrow \infty$
- $o(g(n))$: the set of functions that grow slower than $g(n)$ when $n \rightarrow \infty$
- $\omega(g(n))$: the set of functions that grow faster than $g(n)$ when $n \rightarrow \infty$

We can rewrite the preceding definitions in a more formal way:

$$O(g(n)) \equiv \{f(n) : \exists n_0, c_0, \forall n > n_0, 0 \leq f(n) < c_0 g(n)\} \quad (3.1)$$

$$\Omega(g(n)) \equiv \{f(n) : \exists n_0, c_0, \forall n > n_0, 0 \leq c_0 g(n) < f(n)\} \quad (3.2)$$

$$\Theta(g(n)) \equiv O(g(n)) \cap \Omega(g(n)) \quad (3.3)$$

$$o(g(n)) \equiv O(g(n)) - \Omega(g(n)) \quad (3.4)$$

$$\omega(g(n)) \equiv \Omega(g(n)) - O(g(n)) \quad (3.5)$$

We can also provide a practical rule to determine if a function f belongs to one of the previous sets defined by g .

Compute the limit

$$\lim_{n \rightarrow \infty} \frac{f(n)}{g(n)} = a \quad (3.6)$$

and look up the result in the following table:

a is positive or zero	\implies	$f(n) \in O(g(n))$	\Leftrightarrow	$f \preceq g$	
a is positive or infinity	\implies	$f(n) \in \Omega(g(n))$	\Leftrightarrow	$f \succeq g$	
a is positive	\implies	$f(n) \in \Theta(g(n))$	\Leftrightarrow	$f \sim g$	(3.7)
a is zero	\implies	$f(n) \in o(g(n))$	\Leftrightarrow	$f \prec g$	
a is infinity	\implies	$f(n) \in \omega(g(n))$	\Leftrightarrow	$f \succ g$	

Notice the preceding practical rule assumes the limits exist.

Here is an example:

Given $f(n) = n \log n + 3n$ and $g(n) = n^2$

$$\lim_{n \rightarrow \infty} \frac{n \log n + 3n}{n^2} \xrightarrow{L'H\hat{o}pital} \lim_{n \rightarrow \infty} \frac{1/n}{2} = 0 \quad (3.8)$$

we conclude that $n \log n + 3n$ is in $O(n^2)$.

Given an algorithm A that acts on input of size n , we say that the algorithm is $O(g(n))$ if its worst running time as a function of n is in $O(g(n))$. Similarly, we say that the algorithm is in $\Omega(g(n))$ if its best running time is in $\Omega(g(n))$. We also say that the algorithm is in $\Theta(g(n))$ if both its best running time and its worst running time are in $\Theta(g(n))$.

More formally, we can write the following:

$$T_{\text{worst}}(n) \in O(g(n)) \Rightarrow A \in O(g(n)) \quad (3.9)$$

$$T_{\text{best}}(n) \in \Omega(g(n)) \Rightarrow A \in \Omega(g(n)) \quad (3.10)$$

$$A \in O(g(n)) \text{ and } A \in \Omega(g(n)) \Rightarrow A \in \Theta(g(n)) \quad (3.11)$$

We still have not solved the problem of computing the best, average, and worst running times.

3.1.1 Best and worst running times

The procedure for computing the worst and best running times is similar. It is simple in theory but difficult in practice because it requires an understanding of the algorithm's inner workings.

Consider the following algorithm, which finds the minimum of an array or list A:

```

1 def find_minimum(A):
2     minimum = a[0]
3     for element in A:
4         if element < minimum:
5             minimum = element
6     return minimum

```

To compute the running time in the worst case, we assume that the maximum number of computations is performed. That happens when the if statements are always True. To compute the best running time, we assume that the minimum number of computations is performed. That happens when the if statement is always False. Under each of the two scenarios, we compute the running time by counting how many times the most nested operation is performed.

In the preceding algorithm, the most nested operation is the evaluation of the if statement, and that is executed for each element in A; for example, assuming A has n elements, the if statement will be executed n times.

Therefore both the best and worst running times are proportional to n ,

thus making this algorithm $O(n)$, $\Omega(n)$, and $\Theta(n)$.

More formally, we can observe that this algorithm performs the following operations:

- One assignment (line 2)
- Loops $n = \text{len}(A)$ times (line 3)
- For each loop iteration, performs one comparison (line 4)
- Line 5 is executed only if the condition is true

Because there are no nested loops, the time to execute each loop iteration is about the same, and the running time is proportional to the number of loop iterations.

For a loop iteration that does not contain further loops, the time it takes to compute each iteration, its running time, is constant (therefore equal to 1). For algorithms that contain nested loops, we will have to evaluate nested sums.

Here is the simplest example:

```
1 def loop0(n):
2     for i in xrange(0,n):
3         print i
```

which we can map into

$$T(n) = \sum_{i=0}^{i < n} 1 = n \in \Theta(n) \Rightarrow \text{loop0} \in \Theta(n) \quad (3.12)$$

Here is a similar example where we have a single loop (corresponding to a single sum) that loops n^2 times:

```
1 def loop1(n):
2     for i in xrange(0,n*n):
3         print i
```

and here is the corresponding running time formula:

$$T(n) = \sum_{i=0}^{i < n^2} 1 = n^2 \in \Theta(n^2) \Rightarrow \text{loop1} \in \Theta(n^2) \quad (3.13)$$

The following provides an example of nested loops:

```

1 def loop2(n):
2     for i in xrange(0,n):
3         for j in xrange(0,n):
4             print i,j

```

Here the time for the inner loop is directly determined by n and does not depend on the outer loop's counter; therefore

$$T(n) = \sum_{i=0}^{i < n} \sum_{j=0}^{j < n} 1 = \sum_{i=0}^{i < n} n = n^2 + \dots \in \Theta(n^2) \Rightarrow \text{loop2} \in \Theta(n^2) \quad (3.14)$$

This is not always the case. In the following code, the inner loop does depend on the value of the outer loop:

```

1 def loop3(n):
2     for i in xrange(0,n):
3         for j in xrange(0,i):
4             print i,j

```

Therefore, when we write its running time in terms of a sum, care must be taken that the upper limit of the inner sum is the upper limit of the outer sum:

$$T(n) = \sum_{i=0}^{i < n} \sum_{j=0}^{j < i} 1 = \sum_{i=0}^{i < n} i = \frac{1}{2}n(n-1) \in \Theta(n^2) \Rightarrow \text{loop3} \in \Theta(n^2) \quad (3.15)$$

The appendix of this book provides examples of typical sums that come up in these types of formulas and their solutions.

Here is one more example falling in the same category, although the inner loop depends quadratically on the index of the outer loop:

Example: loop4

```

1 def loop4(n):
2     for i in xrange(0,n):
3         for j in xrange(0,i*i):
4             print i,j

```

Therefore the formula for the running time is more complicated:

$$T(n) = \sum_{i=0}^{i \leq n} \sum_{j=0}^{j < i^2} 1 = \sum_{i=0}^{i \leq n} i^2 = \frac{1}{6}n(n-1)(2n-1) \in \Theta(n^3) \quad (3.16)$$

$$\Rightarrow \text{loop4} \in \Theta(n^3) \quad (3.17)$$

If the algorithm does not contain nested loops, then we need to compute the running time of each loop and take the maximum:

Example: concatenate0

```

1 def concatenate0(n):
2     for i in xrange(n*n):
3         print i
4     for j in xrange(n*n*n):
5         print j

```

$$T(n) = \Theta(\max(n^2, n^3)) \Rightarrow \text{concatenate0} \in \Theta(n^3) \quad (3.18)$$

If there is an if statement, we need to compute the running time for each condition and pick the maximum when computing the worst running time, or the minimum for the best running time:

```

1 def concatenate1(n):
2     if a < 0:
3         for i in xrange(n*n):
4             print i
5     else:
6         for j in xrange(n*n*n):
7             print j

```

$$T_{\text{worst}}(n) = \Theta(\max(n^2, n^3)) \Rightarrow \text{concatenate1} \in O(n^3) \quad (3.19)$$

$$T_{\text{best}}(n) = \Theta(\min(n^2, n^3)) \Rightarrow \text{concatenate1} \in \Omega(n^2) \quad (3.20)$$

This can be expressed more formally as follows:

$$O(f(n)) + \Theta(g(n)) = \Theta(g(n)) \text{ iff } f(n) \in O(g(n)) \quad (3.21)$$

$$\Theta(f(n)) + \Theta(g(n)) = \Theta(g(n)) \text{ iff } f(n) \in O(g(n)) \quad (3.22)$$

$$\Omega(f(n)) + \Theta(g(n)) = \Omega(f(n)) \text{ iff } f(n) \in \Omega(g(n)) \quad (3.23)$$

which we can apply as in the following example:

$$T(n) = \underbrace{[n^2 + n + 3]}_{\Theta(n^2)} + \underbrace{[e^n - \log n]}_{\Theta(e^n)} \in \Theta(e^n) \text{ because } n^2 \in O(e^n) \quad (3.24)$$

3.2 Recurrence relations

The *merge sort* [13] is another sorting algorithm. It is faster than the insertion sort. It was invented by John von Neumann, the physicist credited for inventing also modern computer architecture and game theory.

The merge sort works as follows.

If the input array has length 0 or 1, then it is already sorted, and the algorithm does not perform any other operation.

If the input array has a length greater than 1, it divides the array into two subsets of about half the size. Each subarray is sorted by applying the merge sort recursively (it calls itself!). It then merges the two subarrays back into one sorted array (this step is called *merge*).

Consider the following Python implementation of the merge sort:

```

1 def mergesort(A, p=0, r=None):
2     if r is None: r = len(A)
3     if p < r-1:
4         q = int((p+r)/2)
5         mergesort(A, p, q)
6         mergesort(A, q, r)
7         merge(A, p, q, r)
8
9 def merge(A, p, q, r):
10    B, i, j = [], p, q
11    while True:
12        if A[i] <= A[j]:
13            B.append(A[i])
14            i=i+1
15        else:
16            B.append(A[j])
17            j=j+1
18    if i==q:
19        while j<r:

```

```

20         B.append(A[j])
21         j=j+1
22     break
23     if j==r:
24         while i<q:
25             B.append(A[i])
26             i=i+1
27         break
28     A[p:r]=B

```

Because this algorithm calls itself *recursively*, it is more difficult to compute its running time.

Consider the merge function first. At each step, it increases either i or j , where i is always in between p and q and j is always in between q and r . This means that the running time of the merge is proportional to the total number of values they can span from p to r . This implies that

$$\text{merge} \in \Theta(r - p) \quad (3.25)$$

We cannot compute the running time of the mergesort function using the same direct analysis, but we can assume its running time is $T(n)$, where $n = r - p$ and n is the size of the input data to be sorted and also the difference between its two arguments p and r . We can express this running time in terms of its components:

- It calls itself twice on half of the input data, $2T(n/2)$
- It calls the merge once on the entire data, $\Theta(n)$

We can summarize this into

$$T(n) = 2T(n/2) + n \quad (3.26)$$

This is called a *recurrence relation*. We turned the problem of computing the running time of the algorithm into the problem of solving the recurrence relation. This is now a math problem.

Some recurrence relations can be difficult to solve, but most of them follow in one of these categories:

$$T(n) = aT(n-b) + \Theta(f(n)) \Rightarrow T(n) \in \Theta(\max(a^n, nf(n))) \quad (3.27)$$

$$T(n) = T(b) + T(n-b-a) + \Theta(f(n)) \Rightarrow T(n) \in \Theta(nf(n)) \quad (3.28)$$

$$T(n) = aT(n/b) + \Theta(n^m) \text{ and } a < b^m \Rightarrow T(n) \in \Theta(n^m) \quad (3.29)$$

$$T(n) = aT(n/b) + \Theta(n^m) \text{ and } a = b^m \Rightarrow T(n) \in \Theta(n^m \log n) \quad (3.30)$$

$$T(n) = aT(n/b) + \Theta(n^m) \text{ and } a > b^m \Rightarrow T(n) \in \Theta(n^{\log_b a}) \quad (3.31)$$

$$T(n) = aT(n/b) + \Theta(n^m \log^p n) \text{ and } a < b^m \Rightarrow T(n) \in \Theta(n^m \log^p n) \quad (3.32)$$

$$T(n) = aT(n/b) + \Theta(n^m \log^p n) \text{ and } a = b^m \Rightarrow T(n) \in \Theta(n^m \log^{p+1} n) \quad (3.33)$$

$$T(n) = aT(n/b) + \Theta(n^m \log^p n) \text{ and } a > b^m \Rightarrow T(n) \in \Theta(n^{\log_b a}) \quad (3.34)$$

$$T(n) = aT(n/b) + \Theta(q^n) \Rightarrow T(n) \in \Theta(q^n) \quad (3.35)$$

$$T(n) = aT(n/a-b) + \Theta(f(n)) \Rightarrow T(n) \in \Theta(f(n) \log(n)) \quad (3.36)$$

(they work for $m \geq 0$, $p \geq 0$, and $q > 1$).

These results are a practical simplification of a theorem known as the *master theorem* [14].

3.2.1 Reducible recurrence relations

Other recurrence relations do not immediately fit one of the preceding patterns, but often they can be reduced (transformed) to fit.

Consider the following recurrence relation:

$$T(n) = 2T(\sqrt{n}) + \log n \quad (3.37)$$

We can replace n with $e^k = n$ in eq. (3.37) and obtain

$$T(e^k) = 2T(e^{k/2}) + k \quad (3.38)$$

If we also replace $T(e^k)$ with $S(k) = T(e^k)$, we obtain

$$\underbrace{S(k)}_{T(e^k)} = 2 \underbrace{S(k/2)}_{T(e^{k/2})} + k \quad (3.39)$$

so that we can now apply the master theorem to S . We obtain that $S(k) \in \Theta(k \log k)$. Once we have the order of growth of S , we can determine the order of growth of $T(n)$ by substitution:

$$T(n) = S(\log n) \in \Theta(\underbrace{\log n}_k \log \underbrace{\log n}_k) \quad (3.40)$$

Note that there are recurrence relations that cannot be solved with any of the methods described.

Here are some examples of recursive algorithms and their corresponding recurrence relations with solution:

```

1 def factorial1(n):
2     if n==0:
3         return 1
4     else:
5         return n*factorial1(n-1)

```

$$T(n) = T(n-1) + 1 \Rightarrow T(n) \in \Theta(n) \Rightarrow \text{factorial1} \in \Theta(n) \quad (3.41)$$

```

1 def recursive0(n):
2     if n==0:
3         return 1
4     else:
5         loop3(n)
6         return n*n*recursive0(n-1)

```

$$T(n) = T(n-1) + P_2(n) \Rightarrow T(n) \in \Theta(n^2) \Rightarrow \text{recursive0} \in \Theta(n^3) \quad (3.42)$$

```

1 def recursive1(n):
2     if n==0:
3         return 1
4     else:
5         loop3(n)
6         return n*recursive1(n-1)*recursive1(n-1)

```

$$T(n) = 2T(n-1) + P_2(n) \Rightarrow T(n) \in \Theta(2^n) \Rightarrow \text{recursive1} \in \Theta(2^n) \quad (3.43)$$

```

1 def recursive2(n):
2     if n==0:
3         return 1
4     else:
5         a=factorial0(n)
6         return a*recursive2(n/2)*recursive1(n/2)

```

$$T(n) = 2T(n/2) + P_1(n) \Rightarrow T(n) \in \Theta(n \log n) \Rightarrow \text{recursive2} \in \Theta(n \log n) \quad (3.44)$$

One example of practical interest for us is the binary search below. It finds the location of the element in a sorted input array A :

```

1 def binary_search(A,element):
2     a,b = 0, len(A)-1
3     while b>=a:
4         x = int((a+b)/2)
5         if A[x]<element:
6             a = x+1
7         elif A[x]>element:
8             b = x-1
9         else:
10            return x
11    return None

```

Notice that this algorithm does not appear to be recursive, but in practice, it is because of the apparently infinite while loop. The content of the while loop runs in constant time and then loops again on a problem of half of the original size:

$$T(n) = T(n/2) + 1 \Rightarrow \text{binary_search} \in \Theta(\log n) \quad (3.45)$$

The idea of the `binary_search` is used in the bisection method for solving nonlinear equations.

Do not confuse T notation with Θ notation:

The theta notation can also be used to describe the memory used by an algorithm as a function of the input, T_{memory} , as well as its running time.

Algorithm	Recurrence Relationship	Running time
Binary Search	$T(n) = T(\frac{n}{2}) + \Theta(1)$	$\Theta(\log(n))$
Binary Tree Traversal	$T(n) = 2T(\frac{n}{2}) + \Theta(1)$	$\Theta(n)$
Optimal Sorted Matrix Search	$T(n) = 2T(\frac{n}{2}) + \Theta(\log(n))$	$\Theta(n)$
Merge Sort	$T(n) = T(\frac{n}{2}) + \Theta(n)$	$\Theta(n\log(n))$

3.3 Types of algorithms

Divide-and-conquer is a method of designing algorithms that (informally) proceeds as follows: given an instance of the problem to be solved, split this into several, smaller sub-instances (of the same problem), independently solve each of the sub-instances and then combine the sub-instance solutions to yield a solution for the original instance. This description raises the question, by what methods are the sub-instances to be independently solved? The answer to this question is central to the concept of the divide-and-conquer algorithm and is a key factor in gauging their efficiency. The solution is unique for each problem.

The merge sort algorithm of the previous section is an example of a divide-and-conquer algorithm. In the merge sort, we sort an array by dividing it into two arrays and recursively sorting (conquering) each of the smaller arrays.

Most divide-and-conquer algorithms are recursive, although this is not a requirement.

Dynamic programming is a paradigm that is most often applied in the construction of algorithms to solve a certain class of optimization problems, that is, problems that require the minimization or maximization of some measure. One disadvantage of using divide-and-conquer is that the process of recursively solving separate sub-instances can result in the same computations being performed repeatedly because identical sub-instances may arise. For example, if you are computing the path between two nodes in a graph, some portions of multiple paths will follow the same last few hops. Why compute the last few hops for every path when you would get the same result every time?

The idea behind dynamic programming is to avoid this pathology by obviating the requirement to calculate the same quantity twice. The method usually accomplishes this by maintaining a table of sub-instance results. We say that dynamic programming is a bottom-up technique in which the smallest sub-instances are explicitly solved first and the results of these are used to construct solutions to progressively larger sub-instances. In contrast, we say that the divide-and-conquer is a top-down technique.

We can refactor the mergesort algorithm to eliminate recursion in the algorithm implementation, while keeping the logic of the algorithm unchanged. Here is a possible implementation:

```

1 def mergesort_nonrecursive(A):
2     blocksize, n = 1, len(A)
3     while blocksize < n:
4         for p in xrange(0, n, 2*blocksize):
5             q = p+blocksize
6             r = min(q+blocksize, n)
7             if r > q:
8                 Merge(A, p, q, r)
9             blocksize = 2*blocksize

```

Notice that this has the same running time as the original mergesort because, although it is not recursive, it performs the same operations:

$$T_{best} \in \Theta(n \log n) \quad (3.46)$$

$$T_{average} \in \Theta(n \log n) \quad (3.47)$$

$$T_{worst} \in \Theta(n \log n) \quad (3.48)$$

$$T_{memory} \in \Theta(1) \quad (3.49)$$

Greedy algorithms work in phases. In each phase, a decision is made that appears to be good, without regard for future consequences. Generally, this means that some local optimum is chosen. This “take what you can get now” strategy is the source of the name for this class of algorithms. When the algorithm terminates, we hope that the local optimum is equal to the global optimum. If this is the case, then the algorithm is correct; otherwise, the algorithm has produced a suboptimal solution. If the best answer is not required, then simple greedy algorithms are some-

times used to generate approximate answers, rather than using the more complicated algorithms generally required to generate an exact answer. Even for problems that can be solved exactly by a greedy algorithm, establishing the correctness of the method may be a nontrivial process.

For example, computing change for a purchase in a store is a good case of a greedy algorithm. Assume you need to give change back for a purchase. You would have three choices:

- Give the smallest denomination repeatedly until the correct amount is returned
- Give a random denomination repeatedly until you reach the correct amount. If a random choice exceeds the total, then pick another denomination until the correct amount is returned
- Give the largest denomination less than the amount to return repeatedly until the correct amount is returned

In this case, the third choice is the correct one.

Other types of algorithms do not fit into any of the preceding categories. One is, for example, backtracking. Backtracking is not covered in this course.

3.3.1 Memoization

One case of a top-down approach that is very general and falls under the umbrella of dynamic programming is called *memoization*. Memoization consists of allowing users to write algorithms using a naive divide-and-conquer approach, but functions that may be called more than once are modified so that their output is cached, and if they are called again with the same initial state, instead of the algorithm running again, the output is retrieved from the cache and returned without any computations.

Consider, for example, Fibonacci numbers:

$$\text{Fib}(0) = 0 \quad (3.50)$$

$$\text{Fib}(1) = 1 \quad (3.51)$$

$$\text{Fib}(n) = \text{Fib}(n-1) + \text{Fib}(n-2) \text{ for } n > 1 \quad (3.52)$$

which we can implement using divide-and-conquer as follows:

```

1 def fib(n):
2     return n if n<2 else fib(n-1)+fib(n-2)

```

The recurrence relation for this algorithm is $T(n) = T(n-1) + T(n-2) + 1$, and its solution can be proven to be exponential. This is because this algorithm calls itself more than necessary with the same input values and keeps solving the same subproblem over and over.

Python can implement memoization using the following decorator:

Listing 3.1: in file: nlib.py

```

1 class memoize(object):
2     def __init__(self, f):
3         self.f = f
4         self.storage = {}
5     def __call__(self, *args, **kwargs):
6         key = str((self.f.__name__, args, kwargs))
7         try:
8             value = self.storage[key]
9         except KeyError:
10            value = self.f(*args, **kwargs)
11            self.storage[key] = value
12        return value

```

and simply decorating the recursive function as follows:

Listing 3.2: in file: nlib.py

```

1 @memoize
2 def fib(n):
3     return n if n<2 else fib(n-1)+fib(n-2)

```

which we can call as

Listing 3.3: in file: nlib.py

```

1 >>> print fib(11)
2 89

```

A decorator is a Python function that takes a function and returns a callable object (or a function) to replace the one passed as input. In the previous example, we are using the `@memoize` decorator to replace the `fib` function with the `__call__` argument of the `memoize` class.

This makes the algorithm run much faster. Its running time goes from exponential to linear. Notice that the preceding `memoize` decorator is very general and can be used to decorate any other function.

One more direct dynamic programming approach consists in removing the recursion:

```

1 def fib(n):
2     if n < 2: return n
3     a, b = 0, 1
4     for i in xrange(1,n):
5         a, b = b, a+b
6     return b

```

This also makes the algorithm linear and $T(n) \in \Theta(n)$.

Notice that we easily modify the memoization algorithm to store the partial results in a shared space, for example, on disk using the `PersistentDictionary`:

Listing 3.4: in file: `nlib.py`

```

1 class memoize_persistent(object):
2     STORAGE = 'memoize.sqlite'
3     def __init__(self, f):
4         self.f = f
5         self.storage = PersistentDictionary(memoize_persistent.STORAGE)
6     def __call__(self, *args, **kwargs):
7         key = str((self.f.__name__, args, kwargs))
8         if key in self.storage:
9             value = self.storage[key]
10        else:
11            value = self.f(*args, **kwargs)
12            self.storage[key] = value
13        return value

```

We can use it as we did before, but we can now start and stop the program or run concurrent parallel programs, and as long as they have access to the “`memoize.sqlite`” file, they will share the cache.

3.4 Timing algorithms

The order of growth is a theoretical concept. In practice, we need to time algorithms to check if findings are correct and, more important, to determine the magnitude of the constants in the T functions.

For example, consider this:

```

1 def f1(n):
2     return sum(g1(x) for x in range(n))
3
4 def f2(n):
5     return sum(g2(x) for x in range(n**2))

```

Since f_1 is $\Theta(n)$ and f_2 is $\Theta(n^2)$, we may be led to conclude that the latter is slower. It may very well be that g_1 is 10^6 smaller than g_2 and therefore $T_{f_1}(n) = c_1 n$, $T_{f_2}(n) = c_2 n^2$, but if $c_1 = 10^6 c_2$, then $T_{f_1}(n) > T_{f_2}(n)$ when $n < 10^6$.

To time functions in Python, we can use this simple algorithm:

```

1 def timef(f, ns=1000, dt = 60):
2     import time
3     t = t0 = time.time()
4     for k in xrange(1,ns):
5         f()
6         t = time.time()
7         if t-t0>dt: break
8     return (t-t0)/k

```

This function calls and averages the running time of $f()$ for the minimum between $ns=1000$ iterations and $dt=60$ seconds.

It is now easy, for example, to time the fib function without memoize,

```

1 >>> def fib(n):
2 ...     return n if n<2 else fib(n-1)+fib(n-2)
3 >>> for k in range(15,20):
4 ...     print k,timef(lambda: fib(k))
5 15 0.000315684575338
6 16 0.000576375363706
7 17 0.000936052104732
8 18 0.00135168084153
9 19 0.00217730337912

```

and with memoize,

```

1 >>> @memoize

```

```

2 ... def fib(n):
3 ...     return n if n<2 else fib(n-1)+fib(n-2)
4 >>> for k in range(15,20):
5 ...     print k,timef(lambda:fib(k))
6 15 4.24022311802e-06
7 16 4.02901146386e-06
8 17 4.21922128122e-06
9 18 4.02495429084e-06
10 19 3.73784963552e-06

```

The former shows an exponential behavior; the latter does not.

3.5 Data structures

3.5.1 Arrays

An array is a data structure in which a series of numbers are stored contiguously in memory. The time to access each number (to read or write it) is constant. The time to remove, append, or insert an element may require moving the entire array to a more spacious memory location, and therefore, in the worst case, the time is proportional to the size of the array.

Arrays are the appropriate containers when the number of elements does not change often and when elements have to be accessed in random order.

3.5.2 List

A list is a data structure in which data are not stored contiguously, and each element has knowledge of the location of the next element (and perhaps of the previous element, in a doubly linked list). This means that accessing any element for (read and write) requires finding the element and therefore looping. In the worst case, the time to find an element is proportional to the size of the list. Once an element has been found, any operation on the element, including read, write, delete, and insert, before or after can be done in constant time.

Lists are the appropriate choice when the number of elements can vary

often and when their elements are usually accessed sequentially via iterations.

In Python, what is called a list is actually an array of pointers to the elements.

3.5.3 Stack

A stack data structure is a container, and it is usually implemented as a list. It has the property that the first thing you can take out is the last thing put in. This is commonly known as last-in, first-out, or LIFO. The method to insert or add data to the container is called *push*, and the method to extract data is called *pop*.

In Python, we can implement push by appending an item at the end of a list (Python already has a method for this called `.append`), and we can implement pop by removing the last element of a list and returning it (Python has a method for this called `.pop`).

A simple stack example is as follows:

```

1 >>> stk = []
2 >>> stk.append("One")
3 >>> stk.append("Two")
4 >>> print stk.pop()
5 Two
6 >>> stk.append("Three")
7 >>> print stk.pop()
8 Three
9 >>> print stk.pop()
10 One

```

3.5.4 Queue

A queue data structure is similar to a stack but, whereas the stack returns the most recent item added, a queue returns the oldest item in the list. This is commonly called first-in, first-out, or FIFO. To use Python lists to implement a queue, insert the element to add in the first position of the list as follows:

```

1 >>> que = []
2 >>> que.insert(0, "One")
3 >>> que.insert(0, "Two")
4 >>> print que.pop()
5 One
6 >>> que.insert(0, "Three")
7 >>> print que.pop()
8 Two
9 >>> print que.pop()
10 Three

```

Lists in Python are not an efficient mechanism for implementing queues. Each insertion or removal of an element at the front of a list requires all the elements in the list to be shifted by one. The Python package `collections.deque` is designed to implement queues and stacks. For a stack or queue, you use the same method `.append` to add items. For a stack, `.pop` is used to return the most recent item added, while to build a queue, use `.popleft` to remove the oldest item in the list:

```

1 >>> from collections import deque
2 >>> que = deque([])
3 >>> que.append("One")
4 >>> que.append("Two")
5 >>> print que.popleft()
6 One
7 >>> que.append("Three")
8 >>> print que.popleft()
9 Two
10 >>> print que.popleft()
11 Three

```

3.5.5 Sorting

In the previous sections, we have seen the *insertion sort* and the *merge sort*. Here we consider, as examples, other sorting algorithms: the *quicksort* [13], the *randomized quicksort*, and the *counting sort*:

```

1 def quicksort(A, p=0, r=-1):
2     if r is -1:
3         r=len(A)
4     if p<r-1:
5         q=partition(A, p, r)
6         quicksort(A, p, q)
7         quicksort(A, q+1, r)
8

```



```

9 def partition(A,i,j):
10     x=A[i]
11     h=i
12     for k in xrange(i+1,j):
13         if A[k]<x:
14             h=h+1
15             A[h],A[k] = A[k],A[h]
16     A[h],A[i] = A[i],A[h]
17     return h

```

The running time of the quicksort is given by

$$T_{best} \in \Theta(n \log n) \quad (3.53)$$

$$T_{average} \in \Theta(n \log n) \quad (3.54)$$

$$T_{worst} \in \Theta(n^2) \quad (3.55)$$

The quicksort can also be randomized by picking the pivot, $A[r]$, at random:

```

1 def quicksort(A,p=0,r=-1):
2     if r is -1:
3         r=len(A)
4     if p<r-1:
5         q = random.randint(p,r-1)
6         A[p], A[q] = A[q], A[p]
7         q=partition(A,p,r)
8         quicksort(A,p,q)
9         quicksort(A,q+1,r)

```

In this case, the best and the worst running times do not change, but the average improves when the input is already almost sorted.

The *counting sort* algorithm is special because it only works for arrays of positive integers. This extra requirement allows it to run faster than other sorting algorithms, under some conditions. In fact, this algorithm is linear in the range span by the elements of the input array.

Here is a possible implementation:

```

1 def countingsort(A):
2     if min(A)<0:
3         raise '_counting_sort List Unbound'
4     i, n, k = 0, len(A), max(A)+1
5     C = [0]*k

```

```

6  for j in xrange(n):
7      C[A[j]] = C[A[j]]+1
8  for j in xrange(k):
9      while C[j]>0:
10         (A[i], C[j], i) = (j, C[j]-1, i+1)

```

If we define $k = \max(A) - \min(A) + 1$ and $n = \text{len}(A)$, we see

$$T_{best} \in \Theta(k + n) \tag{3.56}$$

$$T_{average} \in \Theta(k + n) \tag{3.57}$$

$$T_{worst} \in \Theta(k + n) \tag{3.58}$$

$$T_{memory} \in \Theta(k) \tag{3.59}$$

Notice that here we have also computed T_{memory} , for example, the order of growth of memory (not of time) as a function of the input size. In fact, this algorithm differs from the previous ones because it requires a temporary array C .

3.6 Tree algorithms

3.6.1 Heapsort and priority queues

Consider a *complete binary tree* as the one in the following figure:




Figure 3.1: Example of a heap data structure. The number represents not the data in the heap but the numbering of the nodes.

It starts from the top node, called the *root*. Each node has zero, one, or two children. It is called complete because nodes have been added from top to bottom and left to right, filling available slots. We can think of each level of the tree as a generation, where the older generation consists of one node, the next generation of two, the next of four, and so on. We can also number nodes from top to bottom and left to right, as in the image. This allows us to map the elements of a complete binary tree into the elements of an array.

We can implement a complete binary tree using a list, and the child-parent relations are given by the following formulas:

```

1 def heap_parent(i):
2     return int((i-1)/2)
3
4 def heap_left_child(i):
5     return 2*i+1
6
7 def heap_right_child(i):
8     return 2*i+2

```

We can store data (e.g., numbers) in the nodes (or in the corresponding array). If the data are stored in such a way that the value at one node is always greater or equal than the value at its children, the array is called a *heap* and also a *priority queue*.

First of all, we need an algorithm to convert a list into a heap:

```

1 def heapify(A):
2     for i in xrange(int(len(A)/2)-1, -1, -1):
3         heapify_one(A, i)
4
5 def heapify_one(A, i, heapsize=None):
6     if heapsize is None:
7         heapsize = len(A)
8     left = 2*i+1
9     right = 2*i+2
10    if left<heapsize and A[left]>A[i]:
11        largest = left
12    else:
13        largest = i
14    if right<heapsize and A[right]>A[largest]:
15        largest = right
16    if largest!=i:
17        (A[i], A[largest]) = (A[largest], A[i])

```

```
18 heapify_one(A, largest, heapsize)
```

Now we can call `build_heap` on any array or list and turn it into a heap. Because the first element is by definition the smallest, we can use the heap to sort numbers in three steps:

- We turn the array into a heap
- We extract the largest element
- We apply recursion by sorting the remaining elements

Instead of using the preceding divide-and-conquer approach, it is better to use a dynamic programming approach. When we extract the largest element, we swap it with the last element of the array and make the heap one element shorter. The new, shorter heap does not need a full `build_heap` step because the only element out of order is the root node. We can fix this by a single call to `heapify`.

This is a possible implementation for the heapsort [15]:

```
1 def heapsort(A):
2     heapify(A)
3     n = len(A)
4     for i in xrange(n-1, 0, -1):
5         (A[0], A[i]) = (A[i], A[0])
6         heapify_one(A, 0, i)
```

In the average and worst cases, it runs as fast as the quicksort, but in the best case, it is linear:

$$T_{best} \in \Theta(n) \quad (3.60)$$

$$T_{average} \in \Theta(n \log n) \quad (3.61)$$

$$T_{worst} \in \Theta(n \log n) \quad (3.62)$$

$$T_{memory} \in \Theta(1) \quad (3.63)$$

A heap can be used to implement a priority queue, for example, storage from which we can efficiently extract the largest element.

All we need is a function that allows extracting the root element from a heap (as we did in the heapsort and `heapify` of the remaining data) and a

function to push a new value into the heap:

```

1 def heap_pop(A):
2     if len(A)<1:
3         raise RuntimeError('Heap Underflow')
4     largest = A[0]
5     A[0] = A[len(A)-1]
6     del A[len(A)-1]
7     heapify_one(A,0)
8     return largest
9
10 def heap_push(A,value):
11     A.append(value)
12     i = len(A)-1
13     while i>0:
14         j = heap_parent(i)
15         if A[j]<A[i]:
16             (A[i],A[j],i) = (A[j],A[i],j)
17     else:
18         break

```

The running times for heap_pop and heap_push are the same:

$$T_{best} \in \Theta(1) \quad (3.64)$$

$$T_{average} \in \Theta(\log n) \quad (3.65)$$

$$T_{worst} \in \Theta(\log n) \quad (3.66)$$

$$T_{memory} \in \Theta(1) \quad (3.67)$$

Here is an example:

```

1 >>> a = [6,2,7,9,3]
2 >>> heap = []
3 >>> for element in a: heap_push(heap,element)
4 >>> while heap: print heap_pop(heap)
5 9
6 7
7 6
8 3
9 2

```

Heaps find application in many numerical algorithms. In fact, there is a built-in Python module for them called `heapq`, which provides similar functionality to the functions defined here, except that we defined a max heap (pops the max element) while `heapq` is a min heap (pops the mini-

mum):

```

1 >>> from heapq import heappop, heappush
2 >>> a = [6,2,7,9,3]
3 >>> heap = []
4 >>> for element in a: heappush(heap,element)
5 >>> while heap: print heappop(heap)
6 9
7 7
8 6
9 3
10 2

```

Notice heappop instead of heap_pop and heappush instead of heap_push.

3.6.2 Binary search trees

A binary tree is a tree in which each node has at most two children (left and right). A binary tree is called a *binary search tree* if the value of a node is always greater than or equal to the value of its left child and less than or equal to the value of its right child.

A binary search tree is a kind of storage that can efficiently be used for searching if a particular value is in the storage. In fact, if the value for which we are looking is less than the value of the root node, we only have to search the left branch of the tree, and if the value is greater, we only have to search the right branch. Using divide-and-conquer, searching each branch of the tree is even simpler than searching the entire tree because it is also a tree, but smaller.

This means that we can search simply by traversing the tree from top to bottom along some path down the tree. We choose the path by moving down and turning left or right at each node, until we find the element for which we are looking or we find the end of the tree. We can search $T(d)$, where d is the depth of the tree. We will see later that it is possible to build binary trees where $d = \log n$.

To implement it, we need to have a class to represent a binary tree:

```

1 class BinarySearchTree(object):
2     def __init__(self):
3         self.left = self.right = None

```

```

4         self.key = self.value = None
5     def __setitem__(self, key, value):
6         if self.key == None:
7             self.key, self.value = key, value
8         elif key == self.key:
9             self.value = value
10        elif key < self.key:
11            if self.left:
12                self.left[key] = value
13            else:
14                self.left = BinarySearchTree(key, value)
15        else:
16            if self.right:
17                self.right[key] = value
18            else:
19                self.right = BinarySearchTree(key, value)
20    def __getitem__(self, key):
21        if self.key == None:
22            return None
23        elif key == self.key:
24            return self.value
25        elif key < self.key and self.left:
26            return self.left[key]
27        elif key > self.key and self.right:
28            return self.right[key]
29        else:
30            return None
31    def min(self):
32        node = self
33        while node.left:
34            node = self.left
35        return node.key, node.value
36    def max(self):
37        node = self
38        while node.right:
39            node = self.right
40        return node.key, node.value

```

The binary tree can be used as follows:

```

1 >>> root = BinarySearchTree()
2 >>> root[5] = 'aaa'
3 >>> root[3] = 'bbb'
4 >>> root[8] = 'ccc'
5 >>> print root.left.key
6 3
7 >>> print root.left.value
8 bbb
9 >>> print root[3]

```

```

10 bbb
11 >>> print root.max()
12 8 ccc

```

Notice that an empty tree is treated as an exception, where `key = None`.

3.6.3 Other types of trees

There are many other types of trees.

For example, AVL trees are binary search trees that are rebalanced after each insertion or deletion. They are rebalanced in such a way that for each node, the height of the left subtree minus the height of the right subtree is more or less the same. The rebalance operation can be done in $O(\log n)$.

For an AVL tree, the time for inserting or removing an element is given by

$$T_{best} \in \Theta(1) \quad (3.68)$$

$$T_{average} \in \Theta(\log n) \quad (3.69)$$

$$T_{worst} \in \Theta(\log n) \quad (3.70)$$

Until now, we have considered binary trees (each node has two children and stores one value). We can generalize this to k trees, for which each node has k children and stores more than one value.

B-trees are a type of k -tree optimized to read and write large blocks of data. They are normally used to implement database indices and are designed to minimize the amount of data to move when the tree is rebalanced.

3.7 Graph algorithms

A *graph* G is a set of *vertices* V and a set of *links* (also called *edges*) connecting those vertices E . Each link connects one vertex to another.

As an example, you can think of a set of cities connected by roads. The cities are the vertices and the roads are the links.

A link may have attributes. In the case of a road, it could be the name of the road or its length.

In general, a link, indicated with the notation e_{ij} , connecting vertex i with vertex j is called a *directed link*. If the link has no direction $e_{ij} = e_{ji}$, it is called an undirected link. A graph that contains only undirected links is an *undirected graph*; otherwise, it is a *directed graph*.

In the road analogy, some roads can be “one way” (directed links) and some can be “two way” (undirected links).

A *walk* is an alternating sequence of vertices and links, with each link being incident to the vertices immediately preceding and succeeding it in the sequence. A *trail* is a walk with no repeated links.

A *path* is a walk with no repeated vertices. A walk is closed if the initial vertex is also the terminal vertex.

A *cycle* is a closed trail with at least one edge and with no repeated vertices, except that the initial vertex is also the terminal vertex.

A graph that contains no cycles is an *acyclic graph*. Any connected acyclic undirected graph is also a *tree*.

A *loop* is a one-link path connecting a vertex with itself.

A non null graph is *connected* if, for every pair of vertices, there is a walk whose ends are the given vertices. Let us write $i \sim j$ if there is a path from i to j . Then \sim is an equivalence relation. The equivalence classes under \sim are the vertex sets of the connected components of G . A connected graph is therefore a graph with exactly one connected component.

A graph is called *complete* when every pair of vertices is connected by a link (or edge).

A *clique* of a graph is a subset of vertices in which every pair is an edge.

The *degree* of a vertex of a graph is the number of edges incident to it.

If i and j are vertices, the *distance* from i to j , written d_{ij} , is the minimum

length of any path from i to j . In a connected undirected graph, the length of links induces a metric because for every two vertices, we can define their distance as the length of the shortest path connecting them.

The *eccentricity*, $e(i)$, of the vertex i is the maximum value of d_{ij} , where j is allowed to range over all of the vertices of the graph. This gives the largest shortest distance to any connected node in the graph.

The *subgraph* of G induced by a subset W of its vertices V ($W \subseteq V$) is the graph formed by the vertices in W and all edges whose two endpoints are in W .

The graph is the more complex of the data structures considered so far because it includes the tree as a particular case (yes, a tree is also a graph, but in general, a graph is not a tree), and the tree includes a list as a particular case (yes, a list is a tree in which every node has no more than one child); therefore a list is also a particular case of a graph.

The graph is such a general data structure that it can be used to model the brain. Think of neurons as vertices and synapses as links connecting them. We push this analogy later by implementing a simple neural network simulator.

In what follows, we represent a graph in the following way, where links are edges:

```
1 >>> vertices = ['A', 'B', 'C', 'D', 'E']
2 >>> links = [(0,1),(1,2),(1,3),(2,5),(3,4),(3,2)]
3 >>> graph = (vertices, links)
```

Vertices are stored in a list or array and so are links. Each link is a tuple containing the ID of the source vertex, the ID of the target vertex, and perhaps optional parameters. Optional parameters are discussed later, but for now, they may include link details such as length, speed, reliability, or billing rate.

3.7.1 Breadth-first search

The breadth-first search [16] (BFS) is an algorithm designed to visit all vertices in a connected graph. In the cities analogy, we are looking for a

travel strategy to make sure we visit every city reachable by roads, once and only once.

The algorithm begins at one vertex, the origin, and expands out, eventually visiting each node in the graph that is somehow connected to the origin vertex. Its main feature is that it explores the neighbors of the current vertex before moving on to explore remote vertices and their neighbors. It visits other vertices in the same order in which they are discovered.

The algorithm starts by building a table of neighbors so that for each vertex, it knows which other vertices it is connected to. It then maintains two lists, a list of blacknodes (defined as vertices that have been visited) and graynodes (defined as vertices that have been discovered because the algorithm has visited its neighbor). It returns a list of blacknodes in the order in which they have been visited.

Here is the algorithm:

Listing 3.5: in file: nlib.py

```

1 def breadth_first_search(graph,start):
2     vertices, links = graph
3     blacknodes = []
4     graynodes = [start]
5     neighbors = [[] for vertex in vertices]
6     for link in links:
7         neighbors[link[0]].append(link[1])
8     while graynodes:
9         current = graynodes.pop()
10        for neighbor in neighbors[current]:
11            if not neighbor in blacknodes+graynodes:
12                graynodes.insert(0,neighbor)
13        blacknodes.append(current)
14    return blacknodes

```

The BFS algorithm scales as follows:

$$T_{best} \in \Theta(n_E + n_V) \quad (3.71)$$

$$T_{average} \in \Theta(n_E + n_V) \quad (3.72)$$

$$T_{worst} \in \Theta(n_E + n_V) \quad (3.73)$$

$$T_{memory} \in \Theta(n) \quad (3.74)$$

3.7.2 Depth-first search

The depth-first search [17] (DFS) algorithm is very similar to the BFS, but it takes the opposite approach and explores as far as possible along each branch before backtracking.

In the cities analogy, if the BFS was exploring cities in the neighborhood before moving farther away, the DFS does the opposite and brings us first to distant places before visiting other nearby cities.

Here is a possible implementation:

Listing 3.6: in file: nlib.py

```

1 def depth_first_search(graph,start):
2     vertices, links = graph
3     blacknodes = []
4     graynodes = [start]
5     neighbors = [[] for vertex in vertices]
6     for link in links:
7         neighbors[link[0]].append(link[1])
8     while graynodes:
9         current = graynodes.pop()
10        for neighbor in neighbors[current]:
11            if not neighbor in blacknodes+graynodes:
12                graynodes.append(neighbor)
13        blacknodes.append(current)
14    return blacknodes

```

Notice that the BFS and the DFS differ for a single line, which determines whether graynodes is a queue (BSF) or a stack (DFS). When graynodes is a queue, the first vertex discovered is the first visited. When it is a stack, the last vertex discovered is the first visited.

The DFS algorithm goes as follows:

$$T_{best} \in \Theta(n_E + n_V) \quad (3.75)$$

$$T_{average} \in \Theta(n_E + n_V) \quad (3.76)$$

$$T_{worst} \in \Theta(n_E + n_V) \quad (3.77)$$

$$T_{memory} \in \Theta(1) \quad (3.78)$$

3.7.3 Disjoint sets

This is a data structure that can be used to store a set of sets and implements efficiently the join operation between sets. Each element of a set is identified by a representative element. The algorithm starts by placing each element in a set of its own, so there are n initial disjoint sets. Each is represented by itself. When two sets are joined, the representative element of the latter is made to point to the representative element of the former. The set of sets is stored as an array of integers. If at position i the array stores a negative number, this number is interpreted as being the representative element of its own set. If the number stored at position i is instead a nonnegative number j , it means that it belongs to a set that was joined with the set containing j .

Here is the implementation:

Listing 3.7: in file: nlib.py

```

1 class DisjointSets(object):
2     def __init__(self,n):
3         self.sets = [-1]*n
4         self.counter = n
5     def parent(self,i):
6         while True:
7             j = self.sets[i]
8             if j<0:
9                 return i
10            i = j
11    def join(self,i,j):
12        i,j = self.parent(i),self.parent(j)
13        if i!=j:
14            self.sets[i] += self.sets[j]
15            self.sets[j] = i
16            self.counter-=1
17            return True # they have been joined
18        return False  # they were already joined
19    def joined(self,i,j):
20        return self.parent(i) == self.parent(j)
21    def __len__(self):
22        return self.counter

```

Notice that we added a member variable counter that is initialized to the number of disjoint sets and is decreased by one every time two sets are merged. This allows us to keep track of how many disjoint sets exist at

each time. We also override the `__len__` operator so that we can check the value of the counter using the `len` function on a `DisjointSet`.

As an example of application, here is a code that builds a n^d maze. It may be easier to picture it with $d = 2$, a two-dimensional maze. The algorithm works by assuming there is a wall connecting any couple of two adjacent cells. It labels the cells using an integer index. It puts all the cells into a `DisjointSets` data structure and then keeps tearing down walls at random. Two cells on the maze belong to the same set if they are connected, for example, if there is a path that connects them. At the beginning, each cell is its own set because it is isolated by walls. Walls are torn down by being removed from the list `wall` if the wall was separating two disjoint sets of cells. Walls are torn down until all cells belong to the same set, for example, there is a path connecting any cell to any cell:

```

1 def make_maze(n,d):
2     walls = [(i,i+n**j) for i in xrange(n**2) for j in xrange(d) if (i/n**j)%n
3               +1<n]
4     torn_down_walls = []
5     ds = DisjointSets(n**d)
6     random.shuffle(walls)
7     for i,wall in enumerate(walls):
8         if ds.join(wall[0],wall[1]):
9             torn_down_walls.append(wall)
10            if len(ds)==1:
11                break
12    walls = [wall for wall in walls if not wall in torn_down_walls]
13    return walls, torn_down_walls

```

Here is an example of how to use it. This example also draws the walls and the border of the maze:

```

1 >>> walls, torn_down_walls = make_maze(n=20,d=2)

```

The following figure shows a representation of a generated maze:

3.7.4 Minimum spanning tree: Kruskal

Given a connected graph with weighted links (links with a weight or length), a *minimum spanning tree* is a subset of that graph that connects all vertices of the original graph, and the sum of the link weights is minimal. This subgraph is also a tree because the condition of minimal weight




Figure 3.2: Example of a maze as generated using the `DisjointSets` algorithm.

implies that there is only one path connecting each couple of vertices.




Figure 3.3: Example of a minimum spanning tree subgraph of a larger graph. The numbers on the links indicate their weight or length.

One algorithm to build the minimal spanning tree of a graph is the Kruskal [18] algorithm. It works by placing all vertices in a `DisjointSets` structure and looping over links in order of their weight. If the link connects two vertices belonging to different sets, the link is selected to be part of the minimum spanning tree, and the two sets are joined, else the link is ignored. The Kruskal algorithm assumes an undirected graph, for example, all links are bidirectional, and the weight of a link is the same in both directions:

Listing 3.8: in file: `nlib.py`

```

1 def Kruskal(graph):
2     vertices, links = graph
3     A = []
4     S = DisjointSets(len(vertices))
5     links.sort(cmp=lambda a,b: cmp(a[2],b[2]))
6     for source,dest,length in links:
7         if S.join(source,dest):
8             A.append((source,dest,length))
9     return A

```

The Kruskal algorithm goes as follows:

$$T_{worst} \in \Theta(n_E \log n_V) \quad (3.79)$$

$$T_{memory} \in \Theta(n_E) \quad (3.80)$$

We provide an example of application in the next subsection.

3.7.5 Minimum spanning tree: Prim

The Prim [19] algorithm solves the same problem as the Kruskal algorithm, but the Prim algorithm works on a directed graph. It works by placing all vertices in a minimum priority queue where the queue metric for each vertex is the length, or weighted value, of a link connecting the vertex to the closest known neighbor vertex. At each iteration, the algorithm pops a vertex from the priority queue, loops over its neighbors (adjacent links), and, if it finds that one of its neighbors is already in the queue and it is possible to connect it to the current vertex using a shorter link than the one connecting the neighbor to its current closest vertex, the neighbor information is then updated. The algorithm loops until there are no vertices in the priority queue.

The Prim algorithm also differs from the Kruskal algorithm because the former needs a starting vertex, whereas the latter does not. The result when interpreted as a subgraph does not depend on the starting vertex:

Listing 3.9: in file: nlib.py

```

1 class PrimVertex(object):
2     INFINITY = 1e100
3     def __init__(self,id,links):
4         self.id = id

```



```

5     self.closest = None
6     self.closest_dist = PrimVertex.INFINITY
7     self.neighbors = [link[1:] for link in links if link[0]==id]
8     def __cmp__(self,other):
9         return cmp(self.closest_dist, other.closest_dist)
10
11 def Prim(graph, start):
12     from heapq import heappush, heappop, heapify
13     vertices, links = graph
14     P = [PrimVertex(i,links) for i in vertices]
15     Q = [P[i] for i in vertices if not i==start]
16     vertex = P[start]
17     while Q:
18         for neighbor_id,length in vertex.neighbors:
19             neighbor = P[neighbor_id]
20             if neighbor in Q and length<neighbor.closest_dist:
21                 neighbor.closest = vertex
22                 neighbor.closest_dist = length
23         heapify(Q)
24         vertex = heappop(Q)
25     return [(v.id,v.closest.id,v.closest_dist) for v in P if not v.id==start]

```

```

1 >>> vertices = xrange(10)
2 >>> links = [(i,j,abs(math.sin(i+j+1))) for i in vertices for j in vertices]
3 >>> graph = [vertices,links]
4 >>> link = Prim(graph,0)
5 >>> for link in links: print link
6 (1, 4, 0.279...)
7 (2, 0, 0.141...)
8 (3, 2, 0.279...)
9 (4, 1, 0.279...)
10 (5, 0, 0.279...)
11 (6, 2, 0.412...)
12 (7, 8, 0.287...)
13 (8, 7, 0.287...)
14 (9, 6, 0.287...)

```

The Prim algorithm, when using a priority queue for Q , goes as follows:

$$T_{worst} \in \Theta(n_E + n_V \log n_V) \quad (3.81)$$

$$T_{memory} \in \Theta(n_E) \quad (3.82)$$

One important application of the minimum spanning tree is in evolutionary biology. Consider, for example, the DNA for the genes that produce hemoglobin, a molecule responsible for the transport of oxygen in blood. This protein is present in every animal, and the gene is also present in the DNA of every known animal. Yet its DNA structure is a little different.

One can select a pool of animals and, for each two of them, compute the similarity of the DNA of their hemoglobin genes using the `lcs` algorithm discussed later. One can then link each two animals by a metric that represents how similar the two animals are. We can then run the Prim or the Kruskal algorithm to find the minimum spanning tree. The tree represents the most likely evolutionary tree connecting those animal species. Actually, three genes are responsible for hemoglobin (*HBA1*, *HBA2*, and *HBB*). By performing the analysis on different genes and comparing the results, it is possible to establish a consistency check of the results. [20]

Similar studies are performed routinely in evolutionary biology. They can also be applied to viruses to understand how viruses evolved over time. [21]

3.7.6 Single-source shortest paths: Dijkstra

The Dijkstra [22] algorithm solves a similar problem to the Kruskal and Prim algorithms. Given a graph, it computes, for each vertex, the shortest path connecting the vertex to a starting (or source, or root) vertex. The collection of links on all the paths defines the *single-source shortest paths*.

It works, like Prim, by placing all vertices in a min priority queue where the queue metric for each vertex is the length of the path connecting the vertex to the source. At each iteration, the algorithm pops a vertex from the priority queue, loops over its neighbors (adjacent links), and, if it finds that one of its neighbors is already in the queue and it is possible to connect it to the current vertex using a link that makes the path to the source shorter, the neighbor information is updated. The algorithm loops until there are no more vertices in the priority queue.

The implementation of this algorithm is almost identical to the Prim algorithm, except for two lines:

Listing 3.10: in file: `nlib.py`

```

1 def Dijkstra(graph, start):
2     from heapq import heappush, heappop, heapify
3     vertices, links = graph
4     P = [PrimVertex(i,links) for i in vertices]
```

```

5  Q = [P[i] for i in vertices if not i==start]
6  vertex = P[start]
7  vertex.closest_dist = 0
8  while Q:
9      for neighbor_id,length in vertex.neighbors:
10         neighbor = P[neighbor_id]
11         dist = length+vertex.closest_dist
12         if neighbor in Q and dist<neighbor.closest_dist:
13             neighbor.closest = vertex
14             neighbor.closest_dist = dist
15     heapify(Q)
16     vertex = heappop(Q)
17     return [(v.id,v.closest.id,v.closest_dist) for v in P if not v.id==start]

```

Listing 3.11: in file: nlib.py

```

1  >>> vertices = xrange(10)
2  >>> links = [(i,j,abs(math.sin(i+j+1))) for i in vertices for j in vertices]
3  >>> graph = [vertices,links]
4  >>> links = Dijkstra(graph,0)
5  >>> for link in links: print link
6  (1, 2, 0.897...)
7  (2, 0, 0.141...)
8  (3, 2, 0.420...)
9  (4, 2, 0.798...)
10 (5, 0, 0.279...)
11 (6, 2, 0.553...)
12 (7, 2, 0.685...)
13 (8, 0, 0.412...)
14 (9, 0, 0.544...)

```

The Dijkstra algorithm goes as follows:

$$T_{worst} \in \Theta(n_E + n_V \log n_V) \quad (3.83)$$

$$T_{memory} \in \Theta(n_E) \quad (3.84)$$

An application of the Dijkstra is in solving a maze such as the one built when discussing disjoint sets. To use the Dijkstra algorithm, we need to generate a maze, take the links representing torn-down walls, and use them to build an undirected graph. This is done by symmetrizing the links (if i and j are connected, j and i are also connected) and adding to each link a length (1, because all links connect next-neighbor cells):

```

1  >>> n,d = 4, 2
2  >>> walls, links = make_maze(n,d)
3  >>> symmetrized_links = [(i,j,1) for (i,j) in links]+[(j,i,1) for (i,j) in links]

```

```

4 >>> graph = [xrange(n*n),symmetrized_links]
5 >>> links = Dijkstra(graph,0)
6 >>> paths = dict((i,(j,d)) for (i,j,d) in links)

```

Given a maze cell i , `path[i]` gives us a tuple (j, d) where d is the number of steps for the shortest path to reach the origin (o) and j is the ID of the next cell along this path. The following figure shows a generated maze and a reconstructed path connecting an arbitrary cell to the origin:




Figure 3.4: The result shows an application of the Dijkstra algorithm for the single source shortest path applied to solve a maze.

3.8 Greedy algorithms

3.8.1 Huffman encoding

The *Shannon–Fano encoding* [23][24] (also known as **minimal prefix code**) is a lossless data compression algorithm. In this encoding, each character in a string is mapped into a sequence of bits so characters that appear with less frequency are encoded with a longer sequence of bits, whereas characters that appear with more frequency are encoded with a shorter sequence.

The *Huffman encoding* [25] is an implementation of the Shannon–Fano encoding, but the sequence of bits into which each character is mapped is chosen such that the length of the compressed string is minimal. This

choice is constructed in the following way. We associate a tree with each character in the string to compress. Each tree is a trivial tree containing only one node: the root node. We then associate with the root node the frequency of the character representing the tree. We then extract from the list of trees the two trees with rarest or lowest frequency: t_1 and t_2 . We form a new tree, t_3 , we attach t_1 and t_2 to t_3 , and we associate a frequency with t_3 equal to the sum of the frequencies of t_1 and t_2 . We repeat this operation until the list of trees contains only one tree. At this point, we associate a sequence of bits with each node of the tree. Each bit corresponds to one level on the tree. The more frequent characters end up being closer to the root and are encoded with a few bits, while rare characters are far from the root and encoded with more bits.

PKZIP, ARJ, ARC, JPEG, MPEG3 (mp3), MPEG4, and other compressed file formats all use the Huffman coding algorithm for compressing strings. Note that Huffman is a compression algorithm with no information loss. In the JPEG and MPEG compression algorithms, Huffman algorithms are combined with some form or cut of the Fourier spectrum (e.g., MP3 is an audio compression format in which frequencies below 2 KHz are dumped and not compressed because they are not audible). Therefore the JPEG and MPEG formats are referred to as compression with information loss.

Here is a possible implementation of Huffman encoding:

Listing 3.12: in file: nlib.py

```

1 def encode_huffman(input):
2     from heapq import heappush, heappop
3
4     def inorder_tree_walk(t, key, keys):
5         (f,ab) = t
6         if isinstance(ab,tuple):
7             inorder_tree_walk(ab[0],key+'0',keys)
8             inorder_tree_walk(ab[1],key+'1',keys)
9         else:
10            keys[ab] = key
11
12     symbols = {}
13     for symbol in input:
14         symbols[symbol] = symbols.get(symbol,0)+1
15     heap = []
16     for (k,f) in symbols.items():

```

```

17     heappush(heap, (f,k))
18     while len(heap)>1:
19         (f1,k1) = heappop(heap)
20         (f2,k2) = heappop(heap)
21         heappush(heap, (f1+f2, ((f1,k1), (f2,k2))))
22     symbol_map = {}
23     inorder_tree_walk(heap[0], '', symbol_map)
24     encoded = ''.join(symbol_map[symbol] for symbol in input)
25     return symbol_map, encoded
26
27 def decode_huffman(keys, encoded):
28     reversed_map = dict((v,k) for (k,v) in keys.items())
29     i, output = 0, []
30     for j in xrange(1, len(encoded)+1):
31         if encoded[i:j] in reversed_map:
32             output.append(reversed_map[encoded[i:j]])
33             i=j
34     return ''.join(output)

```

We can use it as follows:

Listing 3.13: in file: nlib.py

```

1 >>> input = 'this is a nice day'
2 >>> keys, encoded = encode_huffman(input)
3 >>> print encoded
4 10111001110010001100100011110010101100110100000011111111110
5 >>> decoded = decode_huffman(keys, encoded)
6 >>> print decoded == input
7 True
8 >>> print 1.0*len(input)/(len(encoded)/8)
9 2.57...

```

We managed to compress the original data by a factor 2.57.

We can ask how good is this compression factor. The maximum theoretical best compression factor is given by the Shannon *entropy*, defined as

$$E = - \sum_u w_i \log_2 w_i \quad (3.85)$$

where w_i is the relative frequency of each symbol. In our case, this is easy to compute as

Listing 3.14: in file: nlib.py

```

1 >>> from math import log

```

```

2 >>> input = 'this is a nice day'
3 >>> w = [1.0*input.count(c)/len(input) for c in set(input)]
4 >>> E = -sum(wi*log(wi,2) for wi in w)
5 >>> print E
6 3.23...

```

How could we have done better? Notice for example that the Huffman encoding does not take into account the order in which symbols appear. The original string contains the triple “is” twice, and we could have taken advantage of that pattern, but we did not.

Our choice of using characters as symbols is arbitrary. We could have used a couple of characters as symbols or triplets or any other subsequences of bytes of the original input. We could also have used symbols of different lengths for different parts of the input (we could have used a single symbol for “is”). A different choice would have given a different compression ratio, perhaps better, perhaps worse.

3.8.2 Longest common subsequence

Given two sequences of characters S_1 and S_2 , this is the problem of determining the length of the longest common subsequence (LCS) that is a subsequence of both S_1 and S_2 .

There are several applications for the LCS [26] algorithm:

- **Molecular biology:** DNA sequences (genes) can be represented as sequences of four letters ACGT, corresponding to the four sub-molecules forming DNA. When biologists find a new sequence, they want to find similar sequences or ones that are close. One way of computing how similar two sequences are is to find the length of their LCS.
- **File comparison:** The Unix program `diff` is used to compare two different versions of the same file, to determine what changes have been made to the file. It works by finding a LCS of the lines of the two files and displays the set of lines that have changed. In this instance of the problem, we should think of each line of a file as being a single complicated character.

- **Spelling correction:** If some text contains a word, w , that is not in the dictionary, a “close” word (e.g., one with a small edit distance to w) may be suggested as a correction. Transposition errors are common in written text. A transposition can be treated as a deletion plus an insertion, but a simple variation on the algorithm can treat a transposition as a single point mutation.
- **Speech recognition:** Algorithms similar to the LCS are used in some speech recognition systems—find a close match between a new utterance and one in a library of classified utterances.

Let’s start with some simple observations about the LCS problem. If we have two strings, say, “ATGGCACTACGAT” and “ATCGAGC,” we can represent a subsequence as a way of writing the two so that certain letters line up:

```

1  ATGGCACTACGAT
2  | | | |
3  ATCG AG C

```

From this we can observe the following simple fact: if the two strings start with the same letter, it’s always safe to choose that starting letter as the first character of the subsequence. This is because, if you have some other subsequence, represented as a collection of lines as drawn here, you can “push” the leftmost line to the start of the two strings without causing any other crossings and get a representation of an equally long subsequence that does start this way.

Conversely, suppose that, like in the preceding example, the two first characters differ. Then it is not possible for both of them to be part of a common subsequence. There are three possible choices: remove the first letter from either one of the strings or remove the letter from both strings.

Finally, observe that once we’ve decided what to do with the first characters of the strings, the remaining subproblem is again a LCS problem on two shorter strings. Therefore we can solve it recursively. However, because we don’t know which choice of the three to take, we will take them all and see which choice returns the best result.

Rather than finding the subsequence itself, it turns out to be more efficient

to find the length of the longest subsequence. Then, in the case where the first characters differ, we can determine which subproblem gives the correct solution by solving both and taking the max of the resulting subsequence lengths. Once we turn this into a dynamic programming algorithm, we get the following:

Listing 3.15: in file: `nlib.py`

```

1 def lcs(a, b):
2     previous = [0]*len(a)
3     for i,r in enumerate(a):
4         current = []
5         for j,c in enumerate(b):
6             if r==c:
7                 e = previous[j-1]+1 if i*j>0 else 1
8             else:
9                 e = max(previous[j] if i>0 else 0,
10                        current[-1] if j>0 else 0)
11             current.append(e)
12         previous=current
13     return current[-1]
```

Here is an example:

Listing 3.16: in file: `nlib.py`

```

1 >>> dna1 = 'ATGCTTTAGAGGATGCGTAGATAGCTAAATAGCTCGCTAGA'
2 >>> dna2 = 'GATAGGTACCACAATAATAAGGATAGCTCGCAAATCCTCGA'
3 >>> print lcs(dna1,dna2)
4 26
```

The algorithms can be shown to be $O(nm)$ (where $m = \text{len}(a)$ and $n = \text{len}(b)$).

Another application of this algorithm is in the Unix `diff` utility. Here is a simple example to find the number of common lines between two files:

```

1 >>> a = open('file1.txt').readlines()
2 >>> b = open('file2.txt').readlines()
3 >>> print lcs(a,b)
```

3.8.3 Needleman–Wunsch

With some minor changes to the LCS algorithm, we obtain the Needleman–Wunsch algorithm [27], which solves the problem of *global sequence alignment*. The changes are that, instead of using only two al-

ternating rows (c and d for storing the temporary results, we store all temporary results in an array z; when two matching symbols are found and they are not consecutive, we apply a penalty equal to p^m , where m is the distance between the two matches and is also the size of the gap in the matching subsequence:

Listing 3.17: in file: nlib.py

```

1 def needleman_wunsch(a,b,p=0.97):
2     z=[]
3     for i,r in enumerate(a):
4         z.append([])
5         for j,c in enumerate(b):
6             if r==c:
7                 e = z[i-1][j-1]+1 if i*j>0 else 1
8             else:
9                 e = p*max(z[i-1][j] if i>0 else 0,
10                        z[i][j-1] if j>0 else 0)
11             z[-1].append(e)
12     return z

```

This algorithm can be used to identify common subsequences of DNA between chromosomes (or in general common similar subsequences between any two strings of binary data). Here is an example in which we look for common genes in two randomly generated chromosomes:

Listing 3.18: in file: nlib.py

```

1 >>> bases = 'ATGC'
2 >>> from random import choice
3 >>> genes = [''.join(choice(bases) for k in xrange(10)) for i in xrange(20)]
4 >>> chromosome1 = ''.join(choice(genes) for i in xrange(10))
5 >>> chromosome2 = ''.join(choice(genes) for i in xrange(10))
6 >>> z = needleman_wunsch(chromosome1, chromosome2)
7 >>> Canvas(title='Needleman-Wunsch').imshow(z).save('images/needleman.png')

```

The output of the algorithm is the following image:

The arrow-like patterns in the figure correspond to locations where chromosome1 (Y coordinate) and where chromosome2 (X coordinate) have DNA in common. Those are the places where the sequences are more likely to be aligned for a more detailed comparison.




Figure 3.5: A Needleman and Wunsch plot sequence alignment. The arrow-like patterns indicate the point in the two sequences (represented by the X- and Y-coordinates) where the two sequences are more likely to align.

3.8.4 Continuous Knapsack

Assume you want to fill your knapsack such that you will maximize the value of its contents [28]. However, you are limited by the volume your knapsack can hold. In the continuous knapsack, the amount of each product can vary continuously. In the discrete one, each product has a finite size, and you either carry it or no.

The continuous knapsack problem can be formulated as the problem of maximizing

$$f(x) = a_0x_0 + a_1x_1 + \cdots + a_nx_n \quad (3.86)$$

given the constraint

$$b_0x_0 + b_1x_1 + \cdots + b_nx_n \leq c \quad (3.87)$$

where coefficients a_i , b_i , and c are provided and $x_i \in [0, 1]$ are to be determined.

Using financial terms, we can say that

- The set $\{x_0, x_1, \dots, x_n\}$ forms a portfolio
- b_i is the cost of investment i
- c is the total investment capital available
- a_i is the expected return of investment for investment i
- $f(x)$ is the expected value of our portfolio $\{x_0, x_1, \dots, x_n\}$

Here is the solving algorithm:

Listing 3.19: in file: nlib.py

```

1 def continuum_knapsack(a,b,c):
2     table = [(a[i]/b[i],i) for i in xrange(len(a))]
3     table.sort()
4     table.reverse()
5     f=0.0
6     for (y,i) in table:
7         quantity = min(c/b[i],1)
8         x.append((i,quantity))
9         c = c-b[i]*quantity
10        f = f+a[i]*quantity
11    return (f,x)

```

This algorithm is dominated by the sort; therefore

$$T_{\text{worst}}(x) \in O(n \log n) \quad (3.88)$$

3.8.5 Discrete Knapsack

The discrete Knapsack problem is very similar to the continuous knapsack problem but $x_i \in \{0, 1\}$ (can only be 0 or 1).

Consider the jars of liquids replaced with baskets of objects, say, a basket each of gold bars, silver coins, copper beads, and Rolex watches. How many of each item do you take? The discrete knapsack problem does not consider “baskets of items” but rather all the items together. In this example, dump out all the baskets and you have individual objects to take. Which objects do you take, and which do you leave behind?

In this case, a greedy approach does not apply and the problem is, in

general, NP complete. This concept is defined formally later but it means that there is no known algorithm that can solve this problem and that its order of growth is a polynomial. The best known algorithm has an exponential running time.

This kind of problem is unsolvable for large input.

If we assume that c and b_i are all multiples of a finite factor ε , then it is possible to solve the problem in $O(c/\varepsilon)$. Even when there is not a finite factor ε , we can always round c and b_i to some finite precision ε , and we can conclude that, for any finite precision ε , we can solve the problem in linear time. The algorithm that solves this problem follows a dynamic programming approach.

We can reformulate the problem in terms of a simple capital budgeting problem. We have to invest \$5M. We assume $\varepsilon = \$1\text{M}$. We are in contact with three investment firms. Each offers a number of investment opportunities characterized by an investment cost $c[i, j]$ and an expected return of investment $r[i, j]$. The index i labels the investment firm and the index j labels the different investment opportunities offered by the firm. We have to build a portfolio that maximizes the return of investment. We cannot select more than one investment for each firm, and we cannot select fractions of investments.

Without loss of generality, we will assume that

$$c[i, j] \leq c[i, j + 1] \text{ and } r[i, j] \leq r[i, j + 1] \quad (3.89)$$

which means that investment opportunities for each firm are sorted according to their cost.

Consider the following explicit case:

	Firm	$i = 0$	Firm	$i = 1$	Firm	$i = 2$
proposal	$c[0, j]$	$r[0, j]$	$c[1, j]$	$r[1, j]$	$c[2, j]$	$r[2, j]$
$j = 0$	0	0	0	0	0	0
$j = 1$	1	5	2	8	1	4
$j = 2$	2	6	3	9	-	-
$j = 3$	-	-	4	12	-	-

(Table 1)

(table values are always multiples of $\varepsilon = \$1\text{M}$).

Notice that we can label each possible portfolio by a triplet $\{j_0, j_1, j_2\}$.

A straightforward way to solve this is to try all possibilities and choose the best. In this case, there are only $3 \times 4 \times 2 = 24$ possible portfolios. Many of these are infeasible (e.g., portfolio $\{2, 3, 0\}$ costs \$6M and we cannot afford it). Other portfolios are feasible but very poor (like portfolio $\{0, 0, 1\}$, which is feasible but returns only \$4M).

Here are some disadvantages of total enumeration:

- For larger problems, the enumeration of all possible solutions may not be computationally feasible.
- Infeasible combinations may not be detectable a priori, leading to inefficiency.
- Information about previously investigated combinations is not used to eliminate inferior or infeasible combinations (unless we use memoization, but in this case the algorithm would grow polynomially in memory space).

We can, instead, use a dynamic programming approach.

We break the problem into three stages, and at each stage, we fill a table of optimal investments for each discrete amount of money. At each stage i , we only consider investments from firm i and the table during the previous stage.

So stage 0 represents the money allocated to firm 0, stage 1 the money to firm 1, and stage 2 the money to firm 2.

STAGE ZERO: we maximize the return of investment considering only offers from firm 0. We fill a table $f[0, k]$ with the maximum return of investment if we invest k million dollars in firm 0:

$$f[0, k] = \max_{j | c[0, j] \leq k} r[0, j] \quad (3.90)$$

k	$f[0, k]$
0	0
1	5
2*	6*
3	6
4	6
5	6

(3.91)

STAGE TWO: we maximize the return of investment considering offers from firm 1 and the prior table. We fill a table $f[1, k]$ with the maximum return of investment if we invest k million dollars in firm 0 and firm 1:

$$f[1, k] = \max_{j | c[1, j] \leq k} r[1, j] + f[0, k - c[0, j]] \quad (3.92)$$

k	$c[2, j]$	$f[0, k - c[0, j]]$	$f[1, k]$
0	0	0	0
1	0	1	5
2	2	0	8
3	2	1	9
4	3	1	13
5*	4*	1*	18*

(3.93)

STAGE THREE: we maximize the return of investment considering offers from firm 2 and the preceding table. We fill a table $f[2, k]$ with the maximum return of investment if we invest k million dollars in firm 0, firm 1, and firm 2:

$$f[2, k] = \max_{j | c[2, j] \leq k} r[2, j] + f[1, k - c[1, j]] \quad (3.94)$$

k	$c[2, j]$	$f[1, k - c[1, j]]$	$f[2, k]$
0	0	0	0
1	0	1	5
2	2	0	8
3	2	1	9
4	1	3	13
5*	2*	3*	18*

(3.95)

The maximum return of investment with \$5M is therefore \$18M. It can be achieved by investing \$2M in firm 2 and \$3M in firms 0 and 1. The optimal choice is marked with a star in each table. Note that to determine how much money has to be allocated to maximize the return of investment requires storing past tables to be able to look up the solution to subproblems.

We can generalize eq.(3.92) and eq.(3.94) for any number of investment firms (decision stages):

$$f[i, k] = \max_{j | c[i, j] < k} r[i, j] + f[i - 1, k - c[i - 1, j]] \quad (3.96)$$

3.9 Artificial intelligence and machine learning

3.9.1 Clustering algorithms

There are many algorithms available to cluster data [29]. They are all based on empirical principles because the cluster themselves are defined by the algorithm used to identify them. Normally we distinguish three categories:

- *Hierarchical clustering*: These algorithms start by considering each point a cluster of its own. At each iteration, the two clusters closest to each other are joined together, forming a larger cluster. Hierarchical clustering algorithms differ from each other about the rule used to determine the distance between clusters. The algorithm returns a tree representing the clusters that are joined, called a *dendrogram*.
- *Centroid-based clustering*: These algorithms require that each point be represented by a vector and each cluster also be represented by a vector (centroid of the cluster). With each iteration, a better estimation for the centroids is given. An example of centroid-based clustering is *k-means* clustering. These algorithms require an a priori knowledge of the number of clusters and return the position of the centroids as well the set of points belonging to each cluster.

- *Distribution-based clustering*: These algorithms are based on statistics (more than the other two categories). They assume the points are generated from a distribution (which must be known a priori) and determine the parameters of the distribution. It provides clustering because the distribution may be a sum of more than one localized distribution (each being a cluster).

Both k -means and distribution-based clustering assume an a priori knowledge about the data that often defies the purpose of using clustering: learn something we do now know about the data using an empirical algorithm. They also require that the points be represented by vectors in a Euclidean space, which is not always the case. Consider the case of clustering DNA sequences or financial time series. Technically the latter can be presented as vectors, but their dimensionality can be very large, thus making the algorithms impractical.

Hierarchical clustering only requires the notion of a distance between points, for some of the points.

Phylogenetic Tree of Life

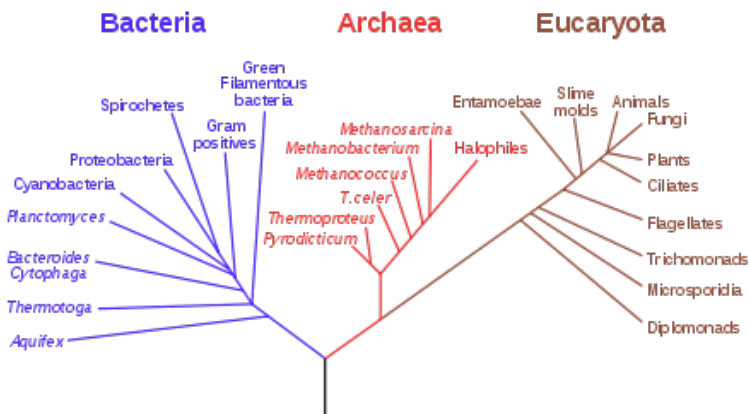


Figure 3.6: Example of a dendrogram.

The following algorithm is a hierarchical clustering algorithm with the following characteristics:

- Individual points do not need to be vectors (although they can be).
- Points may have a weight used to determine their relative importance in identifying the characteristics of the cluster (think of clustering financial assets based on the time series of their returns; the weight could be the average traded volume).
- The distance between points is computed by a metric function provided by the user. The metric can return `None` if there is no known connection between two points.
- The algorithm can be used to build the entire *dendrogram*, or it can stop for a given value of k , a target number of clusters.
- For points that are vectors and a given k , the result is similar to the result of the k -means clustering.

The algorithm works like any other hierarchical clustering algorithm. At the beginning, all-to-all distances are computed and stored in a list `d`. Each point is its own cluster. At each iteration, the two clusters closer together are merged to form one bigger cluster. The distance between each other cluster and the merged cluster is computed by performing a weighted average of the distances between the other cluster and the two merged clusters. The weight factors are provided as input. This is equivalent to what the k -means algorithm does by computing the position of a centroid based on the vectors of the member points.

The algorithm `self.q` implements disjointed sets representing the set of clusters. The algorithm `self.q` is a dictionary. If `self.q[i]` is a list, then i is its own cluster, and the list contains the IDs of the member points. If `self.q[i]` is an integer, then cluster i is no longer its own cluster as it was merged to the cluster represented by the integer.

At each point in time, each cluster is represented by one element, which can be found recursively by `self.parent(i)`. This function returns the ID of the cluster containing element i and returns a list of IDs of all points in the same cluster:

Listing 3.20: in file: `nlib.py`

```
1 class Cluster(object):
```

```

2  def __init__(self, points, metric, weights=None):
3      self.points, self.metric = points, metric
4      self.k = len(points)
5      self.w = weights or [1.0]*self.k
6      self.q = dict((i,[i]) for i,e in enumerate(points))
7      self.d = []
8      for i in xrange(self.k):
9          for j in xrange(i+1,self.k):
10             m = metric(points[i],points[j])
11             if not m is None:
12                 self.d.append((m,i,j))
13      self.d.sort()
14      self.dd = []
15  def parent(self,i):
16      while isinstance(i,int): (parent, i) = (i, self.q[i])
17      return parent, i
18  def step(self):
19      if self.k>1:
20          # find new clusters to join
21          (self.r,i,j),self.d = self.d[0],self.d[1:]
22          # join them
23          i,x = self.parent(i) # find members of cluster i
24          j,y = self.parent(j) # find members if cluster j
25          x += y               # join members
26          self.q[j] = i        # make j cluster point to i
27          self.k -= 1          # decrease cluster count
28          # update all distances to new joined cluster
29          new_d = [] # links not related to joined clusters
30          old_d = {} # old links related to joined clusters
31          for (r,h,k) in self.d:
32              if h in (i,j):
33                  a,b = old_d.get(k,(0.0,0.0))
34                  old_d[k] = a+self.w[k]*r,b+self.w[k]
35              elif k in (i,j):
36                  a,b = old_d.get(h,(0.0,0.0))
37                  old_d[h] = a+self.w[h]*r,b+self.w[h]
38              else:
39                  new_d.append((r,h,k))
40          new_d += [(a/b,i,k) for k,(a,b) in old_d.items()]
41          new_d.sort()
42          self.d = new_d
43          # update weight of new cluster
44          self.w[i] = self.w[i]+self.w[j]
45          # get new list of cluster members
46          self.v = [s for s in self.q.values() if isinstance(s,list)]
47          self.dd.append((self.r,len(self.v)))
48      return self.r, self.v
49
50  def find(self,k):

```

```

51     # if necessary start again
52     if self.k < k: self.__init__(self.points, self.metric)
53     # step until we get k clusters
54     while self.k > k: self.step()
55     # return list of cluster members
56     return self.r, self.v

```

Given a set of points, we can determine the most likely number of clusters representing the data, and we can make a plot of the number of clusters versus distance and look for a plateau in the plot. In correspondence with the plateau, we can read from the y -coordinate the number of clusters. This is done by the function `cluster` in the preceding algorithm, which returns the average distance between clusters and a list of clusters.

For example:

Listing 3.21: in file: `nlib.py`

```

1 >>> def metric(a,b):
2 ...     return math.sqrt(sum((x-b[i])**2 for i,x in enumerate(a)))
3 >>> points = [[random.gauss(i % 5,0.3) for j in xrange(10)] for i in xrange(200)
4 ]
5 >>> c = Cluster(points,metric)
6 >>> r, clusters = c.find(1) # cluster all points until one cluster only
7 >>> Canvas(title='clustering example',xlab='distance',ylab='number of clusters'
8 ...         ).plot(c.dd[150:]).save('clustering1.png')
9 >>> Canvas(title='clustering example (2d projection)',xlab='p[0]',ylab='p[1]'
10 ...         ).ellipses([p[:2] for p in points]).save('clustering2.png')

```

With our sample data, we obtain the following plot (“clustering1.png”):

and the location where the curve bends corresponds to five clusters. Although our points live in 10 dimensions, we can try to project them into two dimensions and see the five clusters (“clustering2.png”):

3.9.2 Neural network

An artificial *neural network* is an electrical circuit (usually simulated in software) that mimics the functionality of the neurons in the animal (and human) brain [30]. It is usually employed in pattern recognition. The network consists of a set of simulated neurons, connected by links (synapses). Some links connect the neurons with each other, some connect the neurons with the input and some with the output. Neurons are usually or-




Figure 3.7: Number of clusters found as a function of the distance cutoff.

ganized in the layers with one *input layer* of neurons connected only with the input and the next layer. Another one, the *output layer*, comprises neurons connected only with the output and previous layers, or many *hidden layers* of neurons connected only with other neurons. Each neuron is characterized by input links and output links. Each output of a neuron is a function of its inputs. The exact shape of that function depends on the network and on parameters that can be adjusted. Usually this function is chosen to be a monotonic increasing function on the sum of the inputs, where both the inputs and the outputs take values in the $[0,1]$ range. The inputs can be thought as electrical signals reaching the neuron. The output is the electrical signal emitted by the neuron. Each neuron is defined by a set of parameters a which determined the relative weight of the input signals. A common choice for this characteristic function is:

$$\text{output}_{ij} = \tanh\left(\sum_k a_{ijk} \text{input}_{ik}\right) \quad (3.97)$$

where i labels the neuron, j labels the output, k labels the input, and a_{ijk}




Figure 3.8: Visual representation of the clusters where the points coordinates are projected in 2D.

are characteristic parameters describing the neurons.

The network is trained by providing an input and adjusting the characteristics a_{ijk} of each neuron k to produce the expected output. The network is trained iteratively until its parameters converge (if they converge), and then it is ready to make predictions. We say the network has learned from the training data set.

Listing 3.22: in file: nlib.py

```

1 class NeuralNetwork:
2     """
3     Back-Propagation Neural Networks
4     Placed in the public domain.
5     Original author: Neil Schemenauer <nas@arctrix.com>
6     Modified by: Massimo Di Pierro
7     Read more: http://www.ibm.com/developerworks/library/l-neural/
8     """
9
10    @staticmethod
11    def rand(a, b):
12        """ calculate a random number where: a <= rand < b """

```




Figure 3.9: Example of a minimalist neural network.

```

13         return (b-a)*random.random() + a
14
15     @staticmethod
16     def sigmoid(x):
17         """ our sigmoid function, tanh is a little nicer than the standard 1/(1+
18             e^-x) """
19         return math.tanh(x)
20
21     @staticmethod
22     def dsigmoid(y):
23         """ # derivative of our sigmoid function, in terms of the output """
24         return 1.0 - y**2
25
26     def __init__(self, ni, nh, no):
27         # number of input, hidden, and output nodes
28         self.ni = ni + 1 # +1 for bias node
29         self.nh = nh
30         self.no = no
31
32         # activations for nodes
33         self.ai = [1.0]*self.ni
34         self.ah = [1.0]*self.nh
35         self.ao = [1.0]*self.no
36
37         # create weights
38         self.wi = Matrix(self.ni, self.nh, fill=lambda r,c: self.rand(-0.2, 0.2)
39             )
39         self.wo = Matrix(self.nh, self.no, fill=lambda r,c: self.rand(-2.0, 2.0)

```

```

    )
39
40     # last change in weights for momentum
41     self.ci = Matrix(self.ni, self.nh)
42     self.co = Matrix(self.nh, self.no)
43
44     def update(self, inputs):
45         if len(inputs) != self.ni-1:
46             raise ValueError('wrong number of inputs')
47
48         # input activations
49         for i in xrange(self.ni-1):
50             self.ai[i] = inputs[i]
51
52         # hidden activations
53         for j in xrange(self.nh):
54             s = sum(self.ai[i] * self.wi[i,j] for i in xrange(self.ni))
55             self.ah[j] = self.sigmoid(s)
56
57         # output activations
58         for k in xrange(self.no):
59             s = sum(self.ah[j] * self.wo[j,k] for j in xrange(self.nh))
60             self.ao[k] = self.sigmoid(s)
61         return self.ao[:]
62
63     def back_propagate(self, targets, N, M):
64         if len(targets) != self.no:
65             raise ValueError('wrong number of target values')
66
67         # calculate error terms for output
68         output_deltas = [0.0] * self.no
69         for k in xrange(self.no):
70             error = targets[k]-self.ao[k]
71             output_deltas[k] = self.dsigmoid(self.ao[k]) * error
72
73         # calculate error terms for hidden
74         hidden_deltas = [0.0] * self.nh
75         for j in xrange(self.nh):
76             error = sum(output_deltas[k]*self.wo[j,k] for k in xrange(self.no))
77             hidden_deltas[j] = self.dsigmoid(self.ah[j]) * error
78
79         # update output weights
80         for j in xrange(self.nh):
81             for k in xrange(self.no):
82                 change = output_deltas[k]*self.ah[j]
83                 self.wo[j,k] = self.wo[j,k] + N*change + M*self.co[j,k]
84                 self.co[j,k] = change
85                 #print N*change, M*self.co[j,k]
86

```



```

87     # update input weights
88     for i in xrange(self.ni):
89         for j in xrange(self.nh):
90             change = hidden_deltas[j]*self.ai[i]
91             self.wi[i,j] = self.wi[i,j] + N*change + M*self.ci[i,j]
92             self.ci[i,j] = change
93
94     # calculate error
95     error = sum(0.5*(targets[k]-self.ao[k])**2 for k in xrange(len(targets))
96                )
97     return error
98
99 def test(self, patterns):
100     for p in patterns:
101         print p[0], '->', self.update(p[0])
102
103 def weights(self):
104     print 'Input weights:'
105     for i in xrange(self.ni):
106         print self.wi[i]
107     print
108     print 'Output weights:'
109     for j in xrange(self.nh):
110         print self.wo[j]
111
112 def train(self, patterns, iterations=1000, N=0.5, M=0.1, check=False):
113     # N: learning rate
114     # M: momentum factor
115     for i in xrange(iterations):
116         error = 0.0
117         for p in patterns:
118             inputs = p[0]
119             targets = p[1]
120             self.update(inputs)
121             error = error + self.back_propagate(targets, N, M)
122         if check and i % 100 == 0:
123             print 'error %-14f' % error

```

In the following example, we teach the network the XOR function, and we create a network with two inputs, two intermediate neurons, and one output. We train it and check what it learned:

Listing 3.23: in file: nlib.py

```

1 >>> pat = [[0,0], [0]], [[0,1], [1]], [[1,0], [1]], [[1,1], [0]]]
2 >>> n = NeuralNetwork(2, 2, 1)
3 >>> n.train(pat)
4 >>> n.test(pat)
5 [0, 0] -> [0.00...]

```

```

6 [0, 1] -> [0.98...]
7 [1, 0] -> [0.98...]
8 [1, 1] -> [-0.00...]

```

Now, we use our neural network to learn patterns in stock prices and predict the next day return. We then check what it has learned, comparing the sign of the prediction with the sign of the actual return for the same days used to train the network:

Listing 3.24: in file: test.py

```

1 >>> storage = PersistentDictionary('sp100.sqlite')
2 >>> v = [day['arithmetic_return']*300 for day in storage['AAPL/2011'][1:]]
3 >>> pat = [[v[i:i+5],v[i+5]] for i in xrange(len(v)-5)]
4 >>> n = NeuralNetwork(5, 5, 1)
5 >>> n.train(pat)
6 >>> predictions = [n.update(item[0]) for item in pat]
7 >>> success_rate = sum(1.0 for i,e in enumerate(predictions)
8 ...                    if e[0]*v[i+5]>0)/len(pat)

```

The learning process depends on the random number generator; therefore, sometimes, for this small training data set, the network succeeds in predicting the sign of the next day arithmetic return of the stock with more than 50% probability, and sometimes it does not. We leave it to the reader to study the significance of this result but using a different subset of the data for the training of the network and for testing its success rate.

3.9.3 Genetic algorithms

Here we consider a simple example of genetic algorithms [31].

We have a population of chromosomes in which each chromosome is just a data structure, in our example, a string of random “ATGC” characters.

We also have a metric to measure the fitness of each chromosome.

At each iteration, only the top-ranking chromosomes in the population survive. The top 10 mate with each other, and their offspring constitute the population for the next iteration. When two members of the population mate, the newborn member of the population has a new DNA sequence, half of which comes from the father and half from the mother, with two randomly mutated DNA basis.

The algorithm stops when we reach a maximum number of generations or we find a chromosome of the population with maximum fitness.

In the following example, the fitness is measured by the similarity between a chromosome and a random target chromosome. The population evolves to approximate better and better that one random target chromosome:

```

1 from random import randint, choice
2
3 class Chromosome:
4     alphabet = 'ATGC'
5     size = 32
6     mutations = 2
7     def __init__(self, father=None, mother=None):
8         if not father or not mother:
9             self.dna = [choice(self.alphabet) for i in xrange(self.size)]
10        else:
11            self.dna = father.dna[:self.size/2]+mother.dna[self.size/2:]
12            for mutation in xrange(self.mutations):
13                self.dna[randint(0,self.size-1)] = choice(self.alphabet)
14        def fitness(self, target):
15            return sum(1 for i,c in enumerate(self.dna) if c==target.dna[i])
16
17 def top(population, target, n=10):
18     table = [(chromo.fitness(target), chromo) for chromo in population]
19     table.sort(reverse = True)
20     return [row[1] for row in table][:n]
21
22 def oneof(population):
23     return population[randint(0, len(population)-1)]
24
25 def main():
26     GENERATIONS = 10000
27     OFFSPRING = 20
28     SEEDS = 20
29     TARGET = Chromosome()
30
31     population = [Chromosome() for i in xrange(SEEDS)]
32     for i in xrange(GENERATIONS):
33         print '\n\nGENERATION:', i
34         print 0, TARGET.dna
35         fittest = top(population, TARGET)
36         for chromosome in fittest: print i, chromosome.dna
37         if max(chromo.fitness(TARGET) for chromo in fittest)==Chromosome.size:
38             print 'SOLUTION FOUND'
39             break

```

```

40     population = [Chromosome(father=oneof(fittest),mother=oneof(fittest)) \
41                     for i in xrange(OFFSPRING)]
42
43 if __name__=='__main__': main()

```

Notice that this algorithm can easily be modified to accommodate other fitness metrics and DNA that consists of a data structure other than a sequence of “ATGC” symbols. The only trickery is finding a proper mating algorithm that preserves some of the fitness features of the parents in the DNA of their offspring. If this does not happen, each next generation loses the fitness properties gained by its parents, thus causing the algorithm not to converge. In our case, it works because if the parents are “close” to the target, then half of the DNA of each parent is also close to the corresponding half of the target DNA. Therefore the DNA of the offspring is as fit as the average of their parents. On top of this, the two random mutations allow the algorithm to further explore the space of all possible DNA sequences.

3.10 Long and infinite loops

3.10.1 P, NP, and NPC

We say a problem is in P if it can be solved in polynomial time: $T_{worst} \in O(n^\alpha)$ for some α .

We say a problem is in NP if an input string can be verified to be a solution in polynomial time: $T_{worst} \in O(n^\alpha)$ for some α .

We say a problem is in co-NP if an input string can be verified not to be a solution in polynomial time: $T_{worst} \in O(n^\alpha)$ for some α .

We say a problem is in NPH (NP Hard) if it is harder than any other problem in NP.

We say a problem is in NPC (NP Complete) if it is in NP and in NPH. Consequences:

$$\text{if } \exists x \mid x \in NPC \text{ and } x \in P \Rightarrow \forall y \in NP, y \in P \quad (3.98)$$

There are a number of open problems about the relations among these sets. Is the set co-NP equivalent to NP? Or perhaps is the intersection between co-NP and NP equal to P? Are NP and NPC the same set? These questions are very important in computer science because if, for example, NP turns out to be the same set as NPC, it means that it must be possible to find algorithms that solve in polynomial time problems that currently do not have a polynomial time solution. Conversely, if one could prove that NP is not equivalent to NPC, we would know that a polynomial time solution to NPC problems does not exist [32].

3.10.2 Cantor's argument

Cantor proved that the real numbers in any interval (e.g., in $[0,1)$) are more than the integer numbers, therefore real numbers are uncountable [33]. The proof proceeds as follows:

1. Consider the real numbers in the interval $[0,1)$ not including 1.
2. Assume that these real numbers are countable. Therefore it is possible to associate each of them to an integer

$$\begin{array}{lll}
 1 & \longleftrightarrow & 0.xxxxxxxxxx\dots \\
 2 & \longleftrightarrow & 0.xxxxxxxxxx\dots \\
 3 & \longleftrightarrow & 0.xxxxxxxxxx\dots \\
 4 & \longleftrightarrow & 0.xxxxxxxxxx\dots \\
 5 & \longleftrightarrow & 0.xxxxxxxxxx\dots \\
 \dots & \dots & \dots
 \end{array} \tag{3.99}$$

(here x represent a decimal digit of a real numbers)

3. Now construct a number $\alpha = 0.yyyyyyyy\dots$ where the first decimal digit differs from the first decimal digit of the first real number of table 3.99, the second decimal digit differs from the second decimal digit of the second real number of table 3.99, and so on and so on for all the

infinite decimal digits:

1

\longleftrightarrow

$0.\overline{x}xxxxxxxxxx\dots$

2

\longleftrightarrow

$0.x\overline{x}xxxxxxxxxx\dots$

3

\longleftrightarrow

$0.xx\overline{x}xxxxxxxxxx\dots$

4

\longleftrightarrow

$0.xxx\overline{x}xxxxxxxxxx\dots$

5

\longleftrightarrow

$0.xxxx\overline{x}xxxxxxxxxx\dots$

...

...

(3.100)

4. The new number α is a real number, and by construction, it is not in the table. In fact, it differs with each item by at least one decimal digit. Therefore the existence of α disproves the assumption that all real numbers in the interval $[0, 1)$ are listed in the table.

There is a very practical consequence of this argument. In fact, in chapter 2, we have seen the distinction between type `float` and class `Decimal`. We have seen about pitfalls of `float` and how `Decimal` can represent floating point numbers with arbitrary precision (assuming we have the memory to do so). Cantor’s argument tells us there are numbers that cannot even be represented as `Decimal` because they would require an infinite amount of storage; π and e are examples of these numbers.

3.10.3 Gödel’s theorem

Gödel used a similar diagonal argument to prove that there are as many problems (or theorems) as real numbers and as many algorithms (or proofs) as natural numbers [33]. Because there is more of the former than the latter, it follows that there are problems for which there is no corresponding solving algorithm. Another interpretation of Gödel’s theorem is that, in any formal language, for example, mathematics, there are theorems that cannot be proved.

Another consequence of Gödel’s theorem is the following: it is impossible to write a computer program to test if a given algorithm stops or enters into an infinite loop.

Consider the following code:

```
1 def next(i):  
2     while len(set(str(i*i))) > 2:  
3         i=i+2  
4     print i  
5  
6 next(81621)
```

This code check searches for a number equal or greater than 81621 which square is comprised of only two digits. Nobody knows whether such number exists, therefore nobody knows if this code stops.

Although one day this problem may be solved, there are many other problems that are still unsolved; actually, there are an infinite number of them.

4

Numerical Algorithms

4.1 Well-posed and stable problems

Numerical algorithms deal mostly with well-posed and stable problems.

A problem is well posed if

- The solution exists and is unique
- The solution has a continuous dependence on input data (a small change in the input causes a small change in the output)

Most physical problems are well posed, except at *critical points*, where any infinitesimal variation in one of the input parameters of the system can cause a large change in the output and therefore in the behavior of the system. This is called *chaos*.

Consider the case of dropping a ball on a triangular-shaped mountain. Let the input of the problem be the horizontal position where the drop occurs and the output the horizontal position of the ball after a fixed amount of time. Almost anywhere the ball is dropped, it will roll down the mountain following deterministic and classical laws of physics, thus the position is calculable and a continuous function of the input position. This is true everywhere, except when the ball is dropped on top of the peak of the mountain. In this case, a minor infinitesimal variation to

the right or to the left can make the ball roll to the right or to the left, respectively. Therefore this is not a well posed problem.

A problem is said to be *stable* if the solution is not just continuous but also weakly sensitive to input data. This means that the change of the output (in percent) is smaller than the change in the input (in percent).

Numerical algorithms work best with stable problems.

We can quantify this as follows. Let x be an input and y be the output of a function:

$$y = f(x) \tag{4.1}$$

We define the condition number of f in x as

$$\text{cond}(f, x) \equiv \frac{|dy/y|}{|dx/x|} = |xf'(x)/f(x)| \tag{4.2}$$

(the latter equality only holds if f is differentiable in x).

A problem with a low condition number is said to be well-conditioned, while a problem with a high condition number is said to be ill-conditioned. XXX

We say that a problem characterized by a function f is well conditioned in a domain D if the condition number is less than 1 for every input in the domain. We also say that a problem is stable if it is well conditioned.

In this book, we are mostly concerned with stable (well-conditioned) problems. If a problem is well-conditioned in for all input in a domain, it is also stable.

4.2 Approximations and error analysis

Consider a physical quantity, for example, the length of a nail. Given one nail, we can measure its length by choosing a measuring instrument. Whatever instrument we choose, we will be able to measure the length of the nail within the resolution of the instrument. For example, with a tape measure with a resolution of 1 mm, we will only be able to determine the

length of the nail within 1 mm of resolution. Repeated measurements performed at different times, by different people, using different instruments may bring different results. We can choose a more precise instrument, but it would not change the fact that different measures will bring different values compatible with the resolution of the instrument. Eventually one will have to face the fact that there may not be such a thing as the length of a nail. For example, the length varies with the temperature and the details of how the measurement is performed. In fact, a nail (as everything else) is made out of atoms, which are made of protons, neutrons, and electrons, which determine an electromagnetic cloud that fluctuates in space and time and depends on the surrounding objects and interacts with the instrument of measure. The length of the nail is the result of a measure.

For each measure there is a result, but the results of multiple measurements are not identical. The results of many measurements performed with the same resolution can be summarized in a distribution of results. This distribution will have a mean \bar{x} and a standard deviation δx , which we call uncertainty. From now on, unless otherwise specified, we assume that the distribution of results is Gaussian so that \bar{x} can be interpreted as the mean and δx as the standard deviation.

Now let us consider a system that, given an input x , produces the output y ; x and y are physical quantities that we can measure, although only with a finite resolution. We can model the system with a function f such that $y = f(x)$ and, in general, f is not known.

We have to make various approximations:

- We can replace the “true” value for the input with our best estimate, \bar{x} , and its associated uncertainty, δx .
- We can replace the “true” value for the output with our best estimate, \bar{y} , and its associated uncertainty, δy .
- Even if we know there is a “true” function f describing the system, our implementation for the function is always an approximation, \bar{f} . In fact, we may not have a single approximation but a series of approxi-

mations of increasing precision, f_n , which become more and more accurate (usually) as n increases. If we are lucky, up to precision errors, as n increases, our approximations will become closer and closer to f , but this will take an infinite amount of time. We have to stop at some finite n .

With the preceding definition, we can define the following types of errors:

- **Data error:** the difference between x and \bar{x} .
- **Computational error:** the difference between $\bar{f}(\bar{x})$ and y . Computational error includes two parts systematic error and statistical error.
- **Statistical error:** due to the fact that, often, the computation of $\bar{f}(x) = \lim_{n \rightarrow \infty} f_n(x)$ is too computationally expensive and we must approximate $\bar{f}(x)$ with $f_n(x)$. This error can be estimated and controlled.
- **Systematic error:** due to the fact that $\bar{f}(x) = \lim_{n \rightarrow \infty} f_n(x) \neq f(x)$. This is for two reasons: modeling errors (we do not know $f(x)$) and rounding errors (we do not implement $f(x)$ with arbitrary precision arithmetics).
- **Total error:** defined as the computational error + the propagated data error and in a formula:

$$\delta y = |f(\bar{x}) - f_n(\bar{x})| + |f'_n(\bar{x})|\delta x \quad (4.3)$$

The first term is the computational error (we use f_n instead of the true f), and the second term is the propagated data error (δx , the uncertainty in x , propagates through f_n).

4.2.1 Error propagation

When a variable x has a finite Gaussian uncertainty δx , how does the uncertainty propagate through a function f ? Assuming the uncertainty is small, we can always expand using a Taylor series:

$$y + \delta y = f(x + \delta x) = f(x) + f'(x)\delta x + O(\delta x^2) \quad (4.4)$$

And because we interpret δy as the width of the distribution y , it should be positive:

$$\delta y = |f'(x)|\delta x \quad (4.5)$$

We have used this formula before for the propagated data error. For functions of two variables $z = f(x, y)$ and assuming the uncertainties in x and y are independent,

$$\delta z = \sqrt{\left| \frac{\partial f(x, y)}{\partial x} \right|^2 \delta x^2 + \left| \frac{\partial f(x, y)}{\partial y} \right|^2 \delta y^2} \quad (4.6)$$

which for simple arithmetic operations reduces to

$$\begin{aligned} z = x + y & \quad \delta z = \sqrt{\delta x^2 + \delta y^2} \\ z = x - y & \quad \delta z = \sqrt{\delta x^2 + \delta y^2} \\ z = x * y & \quad \delta z = |x * y| \sqrt{(\delta x/x)^2 + (\delta y/y)^2} \\ z = x/y & \quad \delta z = |x/y| \sqrt{(\delta x/x)^2 + (\delta y/y)^2} \end{aligned}$$

Notice that when $z = x - y$ approaches zero, the uncertainty in z is larger than the uncertainty in x and y and can overwhelm the result. Also notice that if $z = x/y$ and y is small compared to x , then the uncertainty in z can be large. Bottom line: try to avoid differences between numbers that are in proximity of each other and try to avoid dividing by small numbers.

4.2.2 buckingham

Buckingham is a Python library that implements error propagation and unit conversion. It defines a single class called `Number`, and a number object has value, an uncertainty, and a dimensionality (e.g., length, volume, mass).

Here is an example:

```
1 >>> from buckingham import *
2 >>> globals().update(allunits())
3 >>> L = (4 + pm(0.5)) * meter
4 >>> v = 5 * meter/second
5 >>> t = L/v
6 >>> print t)
```

```

7 (8.00 +/- 1.00)/10
8 >>> print t.units()
9 second
10 >>> print t.convert('hour')
11 (2.222 +/- 0.278)/10^4

```

Notice how adding an uncertainty to a numeric value with `pm(...)` or adding units to a numeric value (integer or floating point) transforms the float number into a `Number` object. A `Number` object behaves like a floating point but propagates its uncertainty and its units. Internally, all units are converted to the International System, unless an explicit conversion is specified.

4.3 Standard strategies

Here are some strategies that are normally employed in numerical algorithms:

- Approximate a continuous system with a discrete system
- Replace integrals with sums
- Replace derivatives with finite differences
- Replace nonlinear with linear + corrections
- Transform a problem into a different one
- Approach the true result by iterations

Here are some examples of each of the strategies.

4.3.1 Approximate continuous with discrete

Consider a ball in a one-dimensional box of size L , and let x be the position of the ball in the box. Instead of treating x as a continuous variable, we can assume a finite resolution of $h = L/n$ (where h is the minimum distance we can distinguish without instruments and n is the maximum number of distinct discrete points we can discriminate), and set $x \equiv hi$, where i is an integer in between 0 and n ; $x = 0$ when $i = 0$ and $x = L$

when $i = n$.

4.3.2 Replace derivatives with finite differences

Computing $df(x)/dx$ analytically is only possible when the function f is expressed in simple analytical terms. Computing it analytically is not possible when $f(x)$ is itself implemented as a numerical algorithm. Here is an example:

```

1 def f(x):
2     (s,t) = (1.0,1.0)
3     for i in xrange(1,10): (s, t) = (s+t, t*x/i)
4     return s

```

What is the derivative of $f(x)$?

The most common ways to define a derivative are the right derivative

$$\frac{df^+(x)}{dx} = \lim_{h \rightarrow 0} \frac{f(x+h) - f(x)}{h} \quad (4.7)$$

the left derivative

$$\frac{df^-(x)}{dx} = \lim_{h \rightarrow 0} \frac{f(x) - f(x-h)}{h} \quad (4.8)$$

and the average of the two

$$\frac{df(x)}{dx} = \frac{1}{2} \left(\frac{df^+(x)}{dx} + \frac{df^-(x)}{dx} \right) = \lim_{h \rightarrow 0} \frac{f(x+h) - f(x-h)}{2h} \quad (4.9)$$

If the function is differentiable in x , then, by definition of “differentiable,” the left and right definitions are equal, and the three prior definitions are equivalent. We can pick one or the other, and the difference will be a systematic error.

If the limit exists, then it means that

$$\frac{df(x)}{dx} = \frac{f(x+h) - f(x-h)}{2h} + O(h) \quad (4.10)$$

where $O(h)$ indicates a correction that, at most, is proportional to h .

The three definitions are equivalent for functions that are differentiable in x , and the latter is preferable because it is more symmetric.

Notice that even more definitions are possible as long as they agree in the limit $h \rightarrow 0$. Definitions that converge faster as h goes to zero are referred to as “improvement.”

We can easily implement the concept of a numerical derivative in code by creating a *functional* D that takes a function f and returns the function $\frac{df(x)}{dx}$ (a functional is a function that returns another function):

Listing 4.1: in file: `nlib.py`

```
1 def D(f,h=1e-6): # first derivative of f
2     return lambda x,f=f,h=h: (f(x+h)-f(x-h))/2/h
```

We can do the same with the second derivative:

$$\frac{d^2 f(x)}{dx^2} = \frac{f(x+h) - 2f(x) - f(x-h)}{h^2} + O(h) \quad (4.11)$$

Listing 4.2: in file: `nlib.py`

```
1 def DD(f,h=1e-6): # second derivative of f
2     return lambda x,f=f,h=h: (f(x+h)-2*f(x)+f(x-h))/(h*h)
```

Here is an example:

Listing 4.3: in file: `nlib.py`

```
1 >>> def f(x): return x*x-5.0*x
2 >>> print f(0)
3 0.0
4 >>> f1 = D(f) # first derivative
5 >>> print f1(0)
6 -5.0
7 >>> f2 = DD(f) # second derivative
8 >>> print f2(0)
9 2.00000...
10 >>> f2 = D(f1) # second derivative
11 >>> print f2(0)
12 1.99999...
```

Notice how composing the first derivative twice or computing the second derivative directly yields a similar result.

We could easily derive formulas for higher-order derivatives and imple-

ment them, but they are rarely needed.

4.3.3 Replace nonlinear with linear

Suppose we are interested in the values of $f(x) = \sin(x)$ for values of x between 0 and 0.1:

```

1 >>> from math import sin
2 >>> points = [0.01*i for i in xrange(0,11)]
3 >>> for x in points:
4 ...     print x, sin(x), "%.2f" % (abs(x-sin(x))/sin(x)*100)
5 0.01 0.009999833... 0.00
6 0.02 0.019998666... 0.01
7 0.03 0.029995500... 0.02
8 0.04 0.039989334... 0.03
9 0.05 0.049979169... 0.04
10 0.06 0.059964006... 0.06
11 0.07 0.069942847... 0.08
12 0.08 0.079914693... 0.11
13 0.09 0.089878549... 0.14
14 0.1 0.0998334166... 0.17

```

Here the first column is the value of x , the second column is the corresponding $\sin(x)$, and the third column is the relative difference (in percent) between x and $\sin(x)$. The difference is always less than 20%; therefore, if we are happy with this precision, then we can replace $\sin(x)$ with x .

This works because any function $f(x)$ can be expanded using a Taylor series. The first order of the Taylor expansion is linear. For values of x sufficiently close to the expansion point, the function can therefore be approximated with its Taylor expansion.

Expanding on the previous example, consider the following code:

```

1 >>> from math import sin
2 >>> points = [0.01*i for i in xrange(0,11)]
3 >>> for x in points:
4 ...     s = x - x*x*x/6
5 ...     print x, math.sin(x), s, ``%.6f`` % (abs(s-sin(x))/(sin(x))*100)
6 0.01 0.009999833... 0.009999... 0.000000
7 0.02 0.019998666... 0.019998... 0.000000
8 0.03 0.029995500... 0.029995... 0.000001
9 0.04 0.039989334... 0.039989... 0.000002

```

```

10 0.05 0.049979169... 0.049979... 0.000005
11 0.06 0.059964006... 0.059964... 0.000011
12 0.07 0.069942847... 0.069942... 0.000020
13 0.08 0.079914693... 0.079914... 0.000034
14 0.09 0.089878549... 0.089878... 0.000055
15 0.1 0.0998334166... 0.099833... 0.000083

```

Notice that the third column $s = x - x^3/6$ is very close to $\sin(x)$. In fact, the difference is less than one part in 10,000 (fourth column). Therefore, for $x \in [-1, +1]$, it is possible to replace the $\sin(x)$ function with the $x - x^3/6$ polynomial. Here we just went one step further in the Taylor expansion, replacing the first order with the third order. The error committed in this approximation is very small.

4.3.4 Transform a problem into a different one

Continuing with the previous example, the polynomial approximation for the \sin function works when x is smaller than 1 but fails when x is greater than or equal to 1. In this case, we can use the following relations to reduce the computation of $\sin(x)$ for large x to $\sin(x)$ for $0 < x < 1$. In particular, we can use

$$\sin(x) = -\sin(-x) \text{ when } x < 0 \quad (4.12)$$

to reduce the domain to $x \in [0, \infty]$. We can then use

$$\sin(x) = \sin(x - 2k\pi) \quad k \in \mathbb{N} \quad (4.13)$$

to reduce the domain to $x \in [0, 2\pi)$

$$\sin(x) = -\sin(2\pi - x) \quad (4.14)$$

to reduce the domain to $x \in [0, \pi)$

$$\sin(x) = \sin(\pi - x) \quad (4.15)$$

to reduce the domain to $x \in [0, \pi/2)$, and

$$\sin(x) = \sqrt{1 - \sin(\pi/2 - x)^2} \quad (4.16)$$

to reduce the domain to $x \in [0, \pi/4)$, where the latter is a subset of $[0, 1)$.

4.3.5 Approximate the true result via iteration

The approximations $\sin(x) \simeq x$ and $\sin(x) \simeq x - x^3/6$ came from linearizing the function $\sin(x)$ and adding a correction to the previous approximation, respectively. In general, we can iterate the process of finding corrections and approximating the true result.

Here is an example of a general iterative algorithm:

```

1 result=guess
2 loop:
3     compute correction
4     result=result+correction
5     if result sufficiently close to true result:
6         return result

```

For the sin function:

```

1 def mysin(x):
2     (s,t) = (0.0,x)
3     for i in xrange(3,10,2): (s, t) = (s+t, -t*x*x/i/(i-1))
4     return s

```

Where do these formulas come from? How do we decide how many iterations we need? We address these problems in the next section.

4.3.6 Taylor series

A function $f(x) : \mathbb{R} \rightarrow \mathbb{R}$ is said to be a *real analytical* in \bar{x} if it is continuous in $x = \bar{x}$ and all its derivatives exist and are continuous in $x = \bar{x}$.

When this is the case, the function can be locally approximated with a

local power series:

$$f(x) = f(\bar{x}) + f^{(0)}(\bar{x})(x - \bar{x}) + \cdots + \frac{f^{(k)}(\bar{x})}{k!}(x - \bar{x})^k + R_k \quad (4.17)$$

The remainder R_k can be proven to be (Taylor's theorem):

$$R_k = \frac{f^{(k+1)}(\xi)}{(k+1)!}(x - \bar{x})^{k+1} \quad (4.18)$$

where ξ is a point in between x and \bar{x} . Therefore, if $f^{(k+1)}$ exists and is limited within a neighborhood $D = \{x \text{ for } |x - \bar{x}| < \epsilon\}$, then

$$|R_k| < \left| \max_{x \in D} f^{(k+1)} \right| |(x - \bar{x})^{k+1}| \quad (4.19)$$

If we stop the Taylor expansion at a finite value of k , the preceding formula gives us the statistical error part of the computational error.

Some Taylor series are very easy to compute:

Exponential for $\bar{x} = 0$:

$$f(x) = e^x \quad (4.20)$$

$$f^{(1)}(x) = e^x \quad (4.21)$$

$$\dots = \dots \quad (4.22)$$

$$f^{(k)}(x) = e^x \quad (4.23)$$

$$e^x = 1 + x + \frac{1}{2}x^2 + \cdots + \frac{1}{k!}x^k + \dots \quad (4.24)$$

Sin for $\bar{x} = 0$:

$$f(x) = \sin(x) \quad (4.25)$$

$$f^{(1)}(x) = \cos(x) \quad (4.26)$$

$$f^{(2)}(x) = -\sin(x) \quad (4.27)$$

$$f^{(3)}(x) = -\cos(x) \quad (4.28)$$

$$\dots = \dots \quad (4.29)$$

$$\sin(x) = x - \frac{1}{3!}x^3 + \dots + \frac{(-1)^n}{(2k+1)!}x^{(2k+1)} + \dots \quad (4.30)$$

We can show the effects of the various terms:

Listing 4.4: in file: nlib.py

```

1 >>> X = [0.03*i for i in xrange(200)]
2 >>> c = Canvas(title='sin(x) approximations')
3 >>> c.plot([(x,math.sin(x)) for x in X],legend='sin(x)')
4 <...>
5 >>> c.plot([(x,x) for x in X[:100]],legend='Taylor 1st')
6 <...>
7 >>> c.plot([(x,x-x**3/6) for x in X[:100]],legend='Taylor 5th')
8 <...>
9 >>> c.plot([(x,x-x**3/6+x**5/120) for x in X[:100]],legend='Taylor 5th')
10 <...>
11 >>> c.save('images/sin.png')
```

Notice that we can very well expand in Taylor around any other point, for example, $\bar{x} = \pi/2$, and we get

$$\sin(x) = 1 - \frac{1}{2}\left(x - \frac{\pi}{2}\right)^2 + \dots + \frac{(-1)^n}{(2k)!}\left(x - \frac{\pi}{2}\right)^{(2k)} + \dots \quad (4.31)$$

and a plot would show:

Listing 4.5: in file: nlib.py

```

1 >>> a = math.pi/2
2 >>> X = [0.03*i for i in xrange(200)]
3 >>> c = Canvas(title='sin(x) approximations')
4 >>> c.plot([(x,math.sin(x)) for x in X],legend='sin(x)')
5 <...>
6 >>> c.plot([(x,1-(x-a)**2/2) for x in X[:150]],legend='Taylor 2nd')
7 <...>
```




Figure 4.1: The figure shows the sin function and its approximation using the Taylor expansion around $x = 0$ at different orders.

```

8 >>> c.plot([(x,1-(x-a)**2/2+(x-a)**4/24) for x in X[:150]], legend='Taylor 4th')
9 <...>
10 >>> c.plot([(x,1-(x-a)**2/2+(x-a)**4/24-(x-a)**6/720) for x in X[:150]], legend='
    Taylor 6th')
11 <...>
12 >>> c.save('images/sin2.png')

```

Similarly we can expand the cos function around $\bar{x} = 0$. Not accidentally, we would get the same coefficients as the Taylor expansion of the sin function around $\bar{x} = \pi/2$. In fact, $\sin(x) = \cos(x - \pi/2)$:




Figure 4.2: The figure shows the sin function and its approximation using the Taylor expansion around $x = \pi/2$ at different orders.

$$f(x) = \cos(x) \quad (4.32)$$

$$f^{(1)}(x) = -\sin(x) \quad (4.33)$$

$$f^{(2)}(x) = -\cos(x) \quad (4.34)$$

$$f^{(3)}(x) = \sin(x) \quad (4.35)$$

$$\dots = \dots \quad (4.36)$$

$$\cos(x) = 1 - \frac{1}{2}x^2 + \dots + \frac{(-1)^n}{(2k)!}x^{(2k)} + \dots \quad (4.37)$$

With a simple replacement, it is easy to prove that

$$e^{ix} = \cos(x) + i \sin(x) \quad (4.38)$$

which will be useful when we talk about Fourier and Laplace transforms.

Now let's consider the k th term in Taylor expansion of e^x . It can be re-

ranged as a function of the previous $(k-1)$ – *th* term:

$$T_k(x) = \frac{1}{k!} x^n = \frac{x}{k} \frac{1}{(k-1)!} x^{k-1} = \frac{x}{k} T_{k-1}(x) \quad (4.39)$$

For $x < 0$, the terms in the sign have alternating sign and are decreasing in magnitude; therefore, for $x < 0$, $R_k < T_{k+1}(1)$. This allows for an easy implementation of the Taylor expansion and its stopping condition:

Listing 4.6: in file: nlib.py

```

1 def myexp(x, precision=1e-6, max_steps=40):
2     if x==0:
3         return 1.0
4     elif x>0:
5         return 1.0/myexp(-x, precision, max_steps)
6     else:
7         t = s = 1.0 # first term
8         for k in xrange(1, max_steps):
9             t = t*x/k # next term
10            s = s + t # add next term
11            if abs(t)<precision: return s
12            raise ArithmeticError('no convergence')
```

This code presents all the features of many of the algorithms we see later in the chapter:

- It deals with the special case $e^0 = 1$.
- It reduces difficult problems to easier problems (exponential of a positive number to the exponential of a negative number via $e^x = 1/e^{-x}$).
- It approximates the “true” solution by iterations.
- The max number of iterations is limited.
- There is a stopping condition.
- It detects failure to converge.

Here is a test of its convergence:

Listing 4.7: in file: nlib.py

```

1 >>> for i in xrange(10):
2     ...     x= 0.1*i
3     ...     assert abs(myexp(x) - math.exp(x)) < 1e-4
```


We can do the same for the sin function:

$$T_k(x) = -\frac{x^2}{(2k)(2k+1)}T_{k-1}(x) \quad (4.40)$$

In this case, the residue is always limited by

$$|R_k| < |x^{2k+1}| \quad (4.41)$$

because the derivatives of sin are always sin and cos and their image is always between $[-1,1]$.

Also notice that the stopping condition is only true when $0 \leq x < 1$. Therefore, for other values of x , we must use trigonometric relations to reduce the problem to a domain where the Taylor series converges.

Hence we write:

Listing 4.8: in file: nlib.py

```

1 def mysin(x,precision=1e-6,max_steps=40):
2     pi = math.pi
3     if x==0:
4         return 0
5     elif x<0:
6         return -mysin(-x)
7     elif x>2.0*pi:
8         return mysin(x % (2.0*pi))
9     elif x>pi:
10        return -mysin(2.0*pi - x)
11    elif x>pi/2:
12        return mysin(pi-x)
13    elif x>pi/4:
14        return sqrt(1.0-mysin(pi/2-x)**2)
15    else:
16        t = s = x                # first term
17        for k in xrange(1,max_steps):
18            t = t*(-1.0)*x*x/(2*k)/(2*k+1) # next term
19            s = s + t                # add next term
20            r = x**(2*k+1)           # estimate residue
21            if r<precision: return s # stopping condition
22        raise ArithmeticError('no convergence')
```

Here we test it:

Listing 4.9: in file: nlib.py

```

1 >>> for i in xrange(10):
2 ...     x= 0.1*i
3 ...     assert abs(mysin(x) - math.sin(x)) < 1e-4

```

Finally, we can do the same for the cos function:

Listing 4.10: in file: nlib.py

```

1 def mycos(x,precision=1e-6,max_steps=40):
2     pi = math.pi
3     if x==0:
4         return 1.0
5     elif x<0:
6         return mycos(-x)
7     elif x>2.0*pi:
8         return mycos(x % (2.0*pi))
9     elif x>pi:
10        return mycos(2.0*pi - x)
11    elif x>pi/2:
12        return -mycos(pi-x)
13    elif x>pi/4:
14        return sqrt(1.0-mycos(pi/2-x)**2)
15    else:
16        t = s = 1 # first term
17        for k in xrange(1,max_steps):
18            t = t*(-1.0)*x*x/(2*k)/(2*k-1) # next term
19            s = s + t # add next term
20            r = x**(2*k) # estimate residue
21            if r<precision: return s # stopping condition
22        raise ArithmeticError('no convergence')

```

Here is a test of convergence:

Listing 4.11: in file: nlib.py

```

1 >>> for i in xrange(10):
2 ...     x = 0.1*i
3 ...     assert abs(mycos(x) - math.cos(x)) < 1e-4

```

4.3.7 Stopping Conditions

To implement a stopping condition, we have two options. We can look at the absolute error, defined as

$$[\text{absolute error}] = [\text{approximate value}] - [\text{true value}] \quad (4.42)$$

or we can look at the relative error

$$[\text{relative error}] = [\text{absolute error}]/[\text{true value}] \quad (4.43)$$

or better, we can consider both. Here is an example of pseudo-code:

```

1 result = guess
2 loop:
3     compute correction
4     result = result+correction
5     compute remainder
6     if |remainder| < target_absolute_precision return result
7     if |remainder| < target_relative_precision*|result| return result

```

In the code, we use the computed result as an estimate of the [true value] and, occasionally, the computed correction is an estimate of the [absolute error]. The target absolute precision is an input value that we use as an upper limit for the absolute error. The target relative precision is an input value we use as an upper limit for the relative error. When absolute error falls below the target absolute precision or the relative error falls below the target relative precision, we stop looping and assume the result is sufficiently precise:

```

1 def generic_looping_function(guess, ap, rp, ns):
2     result = guess
3     for k in xrange(ns):
4         correction = ...
5         result = result+correction
6         remainder = ...
7         if norm(remainder) < max(ap,norm(result)*rp): return result
8     raise ArithmeticError('no convergence')

```

In the preceding code,

- ap is the target absolute precision.
- rp is the target relative precision.
- ns is the maximum number of steps.

From now on, we will adopt this naming convention.

Consider, for example, a financial algorithm that outputs a dollar amount. If it converges to a number very close to 1 or 0, the concept of relative precision loses significance for a result equal to zero, and the algorithm

never detects convergence. In this case, setting an absolute precision of \$1 or 1c is the right thing to do. Conversely, if the algorithm converges to a very large dollar amount, setting a precision of \$1 or 1c may be a too strong requirement, and the algorithm will take too long to converge. In this case, setting a relative precision of 1% or 0.1% is the correct thing to do.

Because in general we do not know in advance the output of the algorithm, we should use both stopping conditions. We should also detect which of the two conditions causes the algorithm to stop looping and return, so that we can estimate the uncertainty in the result.

4.4 Linear algebra

In this section, we consider the following algorithms:

- Arithmetic operation among matrices
- Gauss–Jordan elimination for computing the inverse of a matrix A
- Cholesky decomposition for factorizing a symmetric positive definite matrix A into LL^t , where L is a lower triangular matrix
- The Jacobi algorithms for finding eigenvalues
- Fitting algorithms based on linear least squares

We will provide examples of applications.

4.4.1 Linear systems

In mathematics, a system described by a function f is linear if it is additive:

$$f(x + y) = f(x) + f(y) \quad (4.44)$$

and if it is homogeneous,

$$f(\alpha x) = \alpha f(x) \quad (4.45)$$

In simpler words, we can say that the output is proportional to the input.

As discussed in the introduction to this chapter, one of the simplest techniques for approximating any unknown system consists of approximating it with a linear system (and this approximation will be correct for some system and not for others).

When we try to model a new system, approximating the system with a linear system is often the first step in describing it in a quantitative way, even if it may turn out that this is not a good approximation.

This is the same as approximating the function f describing the system with the first-order Taylor expansions $f(x+h) - f(x) = f'(x)h$.

For a multidimensional system with input \mathbf{x} (now a vector) and output \mathbf{y} (also a vector, not necessarily of the same size as \mathbf{x}), we can still approximate $\mathbf{y} = f(\mathbf{x})$ with $f(\mathbf{y} + \mathbf{h}) - \mathbf{y} \simeq A\mathbf{h}$, yet we need to clarify what this latter equation means.

Given

$$\mathbf{x} \equiv \begin{pmatrix} x_0 \\ x_1 \\ \dots \\ x_{n-1} \end{pmatrix} \quad \mathbf{y} \equiv \begin{pmatrix} y_0 \\ y_1 \\ \dots \\ y_{m-1} \end{pmatrix} \quad (4.46)$$

$$A \equiv \begin{pmatrix} a_{00} & a_{01} & \dots & a_{0,n-1} \\ a_{10} & a_{11} & \dots & a_{1,n-1} \\ \dots & \dots & \dots & \dots \\ a_{m-1,0} & a_{m-1,1} & \dots & a_{m-1,n-1} \end{pmatrix} \quad (4.47)$$

the following equation means

$$\mathbf{y} = f(\mathbf{x}) \simeq A\mathbf{x} \quad (4.48)$$

which means

$$y_0 = f_0(\mathbf{x}) \simeq a_{00}x_0 + a_{01}x_1 + \cdots + a_{0,n-1}x_{n-1} \quad (4.49)$$

$$y_1 = f_1(\mathbf{x}) \simeq a_{10}x_0 + a_{11}x_1 + \cdots + a_{1,n-1}x_{n-1} \quad (4.50)$$

$$y_2 = f_2(\mathbf{x}) \simeq a_{20}x_0 + a_{21}x_1 + \cdots + a_{2,n-1}x_{n-1} \quad (4.51)$$

$$\dots = \dots \simeq \dots \quad (4.52)$$

$$y_{m-1} = f_{m-1}(\mathbf{x}) \simeq a_{m-1,0}x_0 + a_{m-1,1}x_1 + \cdots a_{m-1,n-1}x_{n-1} \quad (4.53)$$

which says that every output variable y_j is approximated with a function proportional to each of the input variables x_i .

A system is linear if the \simeq relations turn out to be exact and can be replaced by $=$ symbols.

As a corollary of the basic properties of a linear system discussed earlier, linear systems have one nice additional property. If we combine two linear systems $y = Ax$ and $z = By$, the combined system is also a linear system $z = (BA)x$.

Elementary algebra is defined as a set of numbers (e.g., real numbers) endowed with the ordinary four elementary operations ($+$, $-$, \times , $/$).

Abstract algebra is a generalization of the concept of elementary algebra to other sets of objects (not necessarily numbers) by definition operations among them such as addition and multiplication.

Linear algebra is the extension of elementary algebra to matrices (and vectors, which can be seen as special types of matrices) by defining the four elementary operations among them.

We will implement them in code using Python. In Python, we can implement a matrix as a list of lists, as follows:

```
>>> A = [[1,2,3],[4,5,6],[7,8,9]]
```

But such an object (a list of lists) does not have the mathematical properties we want, so we have to define them.

First, we define a class representing a matrix:

Listing 4.12: in file: `nlib.py`

```

1 class Matrix(object):
2     def __init__(self, rows, cols=1, fill=0.0):
3         """
4         Constructor a zero matrix
5         Examples:
6         A = Matrix([[1,2],[3,4]])
7         A = Matrix([1,2,3,4])
8         A = Matrix(10,20)
9         A = Matrix(10,20,fill=0.0)
10        A = Matrix(10,20,fill=lambda r,c: 1.0 if r==c else 0.0)
11        """
12        if isinstance(rows,list):
13            if isinstance(rows[0],list):
14                self.rows = [[e for e in row] for row in rows]
15            else:
16                self.rows = [[e] for e in rows]
17        elif isinstance(rows,int) and isinstance(cols,int):
18            xrows, xcols = xrange(rows), xrange(cols)
19            if callable(fill):
20                self.rows = [[fill(r,c) for c in xcols] for r in xrows]
21            else:
22                self.rows = [[fill for c in xcols] for r in xrows]
23        else:
24            raise RuntimeError("Unable to build matrix from %s" % repr(rows))
25        self.nrows = len(self.rows)
26        self.ncols = len(self.rows[0])

```

Notice that the constructor takes the number of rows and columns (cols) of the matrix but also a fill value, which can be used to initialize the matrix elements and defaults to zero. It can be callable in case we need to initialize the matrix with row,col dependent values.

The actual matrix elements are stored as a list or array into the data member variable. If optimize=True, the data are stored in an array of double precision floating point numbers ("d"). This optimization will prevent you from building matrices of complex numbers or matrices of arbitrary precision decimal numbers.

Now we define a getter method, a setter method, and a string representation for the matrix elements:

Listing 4.13: in file: nlib.py

```

1 def __getitem__(A,coords):
2     " x = A[0,1]"
3     i,j = coords

```

```

4         return A.rows[i][j]
5
6     def __setitem__(A, coords, value):
7         " A[0,1] = 3.0 "
8         i,j = coords
9         A.rows[i][j] = value
10
11    def tolist(A):
12        " assert(Matrix([[1,2],[3,4]]).tolist() == [[1,2],[3,4]]) "
13        return A.rows
14
15    def __str__(A):
16        return str(A.rows)
17
18    def flatten(A):
19        " assert(Matrix([[1,2],[3,4]]).flatten() == [1,2,3,4]) "
20        return [A[r,c] for r in xrange(A.nrows) for c in xrange(A.ncols)]
21
22    def reshape(A,n,m):
23        " assert(Matrix([[1,2],[3,4]]).reshape(1,4).tolist() == [[1,2,3,4]]) "
24        if n*m != A.nrows*A.ncols:
25            raise RuntimeError("Impossible reshape")
26        flat = A.flatten()
27        return Matrix(n,m,fill=lambda r,c,m=m,flat=flat: flat[r*m+c])
28
29    def swap_rows(A,i,j):
30        " assert(Matrix([[1,2],[3,4]]).swap_rows(1,0).tolist() == [[3,4],[1,2]]) "
31        A.rows[i], A.rows[j] = A.rows[j], A.rows[i]

```

We also define some convenience functions for constructing the identity matrix (given its size) and a diagonal matrix (given the diagonal elements). We make these methods static because they do not act on an existing matrix.

Listing 4.14: in file: nlib.py

```

1    @staticmethod
2    def identity(rows=1,e=1.0):
3        return Matrix(rows,rows,lambda r,c,e=e: e if r==c else 0.0)
4
5    @staticmethod
6    def diagonal(d):
7        return Matrix(len(d),len(d),lambda r,c,d=d:d[r] if r==c else 0.0)

```

Now we are ready to define arithmetic operations among matrices. We start with addition and subtraction:

Listing 4.15: in file: nlib.py

```

1  def __add__(A,B):
2      """
3      Adds A and B element by element, A and B must have the same size
4      Example
5      >>> A = Matrix([[4,3.0], [2,1.0]])
6      >>> B = Matrix([[1,2.0], [3,4.0]])
7      >>> C = A + B
8      >>> print C
9      [[5, 5.0], [5, 5.0]]
10     """
11     n, m = A.nrows, A.ncols
12     if not isinstance(B,Matrix):
13         if n==m:
14             B = Matrix.identity(n,B)
15         elif n==1 or m==1:
16             B = Matrix([[B for c in xrange(m)] for r in xrange(n)])
17     if B.nrows!=n or B.ncols!=m:
18         raise ArithmeticError('incompatible dimensions')
19     C = Matrix(n,m)
20     for r in xrange(n):
21         for c in xrange(m):
22             C[r,c] = A[r,c]+B[r,c]
23     return C
24
25  def __sub__(A,B):
26      """
27      Adds A and B element by element, A and B must have the same size
28      Example
29      >>> A = Matrix([[4.0,3.0], [2.0,1.0]])
30      >>> B = Matrix([[1.0,2.0], [3.0,4.0]])
31      >>> C = A - B
32      >>> print C
33      [[3.0, 1.0], [-1.0, -3.0]]
34     """
35     n, m = A.nrows, A.ncols
36     if not isinstance(B,Matrix):
37         if n==m:
38             B = Matrix.identity(n,B)
39         elif n==1 or m==1:
40             B = Matrix(n,m,fill=B)
41     if B.nrows!=n or B.ncols!=m:
42         raise ArithmeticError('Incompatible dimensions')
43     C = Matrix(n,m)
44     for r in xrange(n):
45         for c in xrange(m):
46             C[r,c] = A[r,c]-B[r,c]
47     return C

```

```

48 def __radd__(A,B): #B+A
49     return A+B
50 def __rsub__(A,B): #B-A
51     return (-A)+B
52 def __neg__(A):
53     return Matrix(A.nrows,A.ncols,fill=lambda r,c:-A[r,c])

```

With the preceding definitions, we can add matrices to matrices, subtract matrices from matrices, but also add and subtract scalars to and from matrices and vectors (scalars are interpreted as diagonal matrices when added to square matrices and as constant vectors when added to vectors).

Here are some examples:

Listing 4.16: in file: nlib.py

```

1 >>> A = Matrix([[1.0,2.0],[3.0,4.0]])
2 >>> print A + A      # calls A.__add__(A)
3 [[2.0, 4.0], [6.0, 8.0]]
4 >>> print A + 2      # calls A.__add__(2)
5 [[3.0, 2.0], [3.0, 6.0]]
6 >>> print A - 1      # calls A.__add__(1)
7 [[0.0, 2.0], [3.0, 3.0]]
8 >>> print -A         # calls A.__neg__()
9 [[-1.0, -2.0], [-3.0, -4.0]]
10 >>> print 5 - A      # calls A.__rsub__(5)
11 [[4.0, -2.0], [-3.0, 1.0]]
12 >>> b = Matrix([[1.0],[2.0],[3.0]])
13 >>> print b + 2      # calls b.__add__(2)
14 [[3.0], [4.0], [5.0]]

```

The class `Matrix` works with complex numbers as well:

Listing 4.17: in file: nlib.py

```

1 >>> A = Matrix([[1,2],[3,4]])
2 >>> print A + 1j
3 [[(1+1j), (2+0j)], [(3+0j), (4+1j)]]

```

Now we implement multiplication. We are interested in three types of multiplication: multiplication of a scalar by a matrix (`__rmul__`), multiplication of a matrix by a matrix (`__mul__`), and scalar product between two vectors (also handled by `__mul__`):

Listing 4.18: in file: nlib.py

```

1 def __rmul__(A,x):
2     "multiplies a number of matrix A by a scalar number x"
3     import copy

```

```

4     M = copy.deepcopy(A)
5     for r in xrange(M.nrows):
6         for c in xrange(M.ncols):
7             M[r,c] *= x
8     return M
9
10    def __mul__(A,B):
11        "multiplies a number of matrix A by another matrix B"
12        if isinstance(B,(list,tuple)):
13            return (A*Matrix(len(B),1,fill=lambda r,c:B[r])).nrows
14        elif not isinstance(B,Matrix):
15            return B*A
16        elif A.ncols == 1 and B.ncols==1 and A.nrows == B.nrows:
17            # try a scalar product ;-)
18            return sum(A[r,0]*B[r,0] for r in xrange(A.nrows))
19        elif A.ncols!=B.nrows:
20            raise ArithmeticError('Incompatible dimension')
21        M = Matrix(A.nrows,B.ncols)
22        for r in xrange(A.nrows):
23            for c in xrange(B.ncols):
24                for k in xrange(A.ncols):
25                    M[r,c] += A[r,k]*B[k,c]
26        return M

```

This allows us the following operations:

Listing 4.19: in file: nlib.py

```

1 >>> A = Matrix([[1.0,2.0],[3.0,4.0]])
2 >>> print(2*A)      # scalar * matrix
3 [[2.0, 4.0], [6.0, 8.0]]
4 >>> print(A*A)      # matrix * matrix
5 [[7.0, 10.0], [15.0, 22.0]]
6 >>> b = Matrix([[1],[2],[3]])
7 >>> print(b*b)      # scalar product
8 14

```

4.4.2 Examples of linear transformations

In this section, we try to provide an intuitive understanding of two-dimensional linear transformations.

In the following code, we consider an image (a set of points) containing a circle and two orthogonal axes. We then apply the following linear transformations to it:

- A_1 , which scales the X-axis
- A_2 , which scales the Y-axis
- S , which scales both axes
- B_1 , which scales the X-axis and then rotates (R) the image 0.5 rad
- B_2 , which is neither a scaling nor a rotation; as it can be seen from the image, it does not preserve angles

Listing 4.20: in file: nlib.py

```

1 >>> points = [(math.cos(0.0628*t),math.sin(0.0628*t)) for t in xrange(200)]
2 >>> points += [(0.02*t,0) for t in xrange(50)]
3 >>> points += [(0,0.02*t) for t in xrange(50)]
4 >>> Canvas(title='Linear Transformation',xlab='x',ylab='y',
5 ...       xrange=(-1,1), yrange=(-1,1)).ellipses(points).save('la1.png')
6 >>> def f(A,points,filename):
7 ...     data = [(A[0,0]*x+A[0,1]*y,A[1,0]*x+A[1,1]*y) for (x,y) in points]
8 ...     Canvas(title='Linear Transformation',xlab='x',ylab='y'
9 ...           ).ellipses(points).ellipses(data).save(filename)
10 >>> A1 = Matrix([[0.2,0],[0,1]])
11 >>> f(A1, points, 'la2.png')
12 >>> A2 = Matrix([[1,0],[0,0.2]])
13 >>> f(A2, points, 'la3.png')
14 >>> S = Matrix([[0.3,0],[0,0.3]])
15 >>> f(S, points, 'la4.png')
16 >>> s, c = math.sin(0.5), math.cos(0.5)
17 >>> R = Matrix([[c,-s],[s,c]])
18 >>> B1 = R*A1
19 >>> f(B1, points, 'la5.png')
20 >>> B2 = Matrix([[0.2,0.4],[0.5,0.3]])
21 >>> f(B2, points, 'la6.png')

```

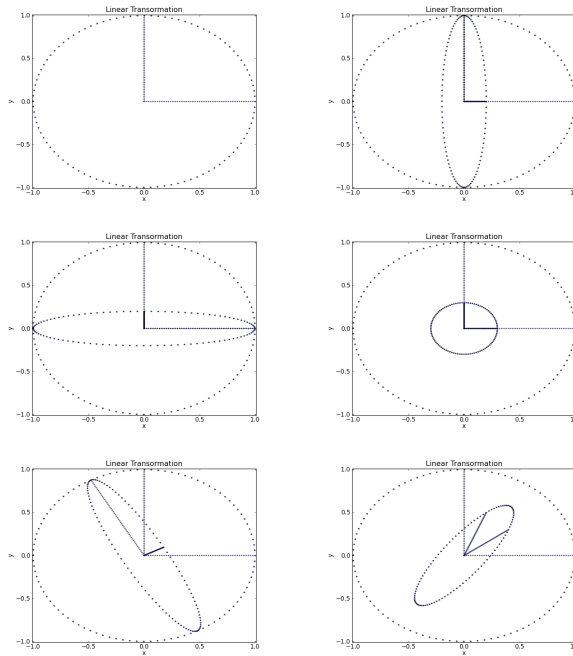


Figure 4.3: Example of the effect of different linear transformations on the same set of points. From left to right, top to bottom, they show stretching along both the X - and Y -axes, scaling across both axes, a rotation, and a generic transformation that does not preserve angles.

4.4.3 Matrix inversion and the Gauss–Jordan algorithm

Implementing the inverse of the multiplication (division) is a more challenging task.

We define A^{-1} , the inverse of the square matrix A , as that matrix such that for every vector b , $A(x) = b$ implies $(x) = A^{-1}b$. The Gauss–Jordan algorithm computes A^{-1} given A .

To implement it, we must first understand how it works. Consider the following equation:

$$Ax = b \quad (4.54)$$

We can rewrite it as:

$$A\mathbf{x} = B\mathbf{b} \quad (4.55)$$

where $B = 1$, the identity matrix. This equation remains true if we multiply both terms by a nonsingular matrix S_0 :

$$S_0 A \mathbf{x} = S_0 B \mathbf{b} \quad (4.56)$$

The trick of the Gauss-Jordan elimination consists in finding a series of matrices S_0, S_1, \dots, S_{n-1} so that

$$S_{n-1} \dots S_1 S_0 A \mathbf{x} = S_{n-1} \dots S_1 S_0 B \mathbf{b} = \mathbf{x} \quad (4.57)$$

Because the preceding expression must be true for every b and because x is the solution of $Ax = b$, by definition, $S_{n-1} \dots S_1 S_0 B \equiv A^{-1}$.

The Gauss-Jordan algorithm works exactly this way: given A , it computes A^{-1} :

Listing 4.21: in file: nlib.py

```

1  def __rdiv__(A,x):
2      """Computes x/A using Gauss-Jordan elimination where x is a scalar"""
3      import copy
4      n = A.ncols
5      if A.nrows != n:
6          raise ArithmeticError('matrix not squared')
7      indexes = xrange(n)
8      A = copy.deepcopy(A)
9      B = Matrix.identity(n,x)
10     for c in indexes:
11         for r in xrange(c+1,n):
12             if abs(A[r,c])>abs(A[c,c]):
13                 A.swap_rows(r,c)
14                 B.swap_rows(r,c)
15         p = 0.0 + A[c,c] # trick to make sure it is not integer
16         for k in indexes:
17             A[c,k] = A[c,k]/p
18             B[c,k] = B[c,k]/p
19         for r in range(0,c)+range(c+1,n):
20             p = 0.0 + A[r,c] # trick to make sure it is not integer
21             for k in indexes:
22                 A[r,k] -= A[c,k]*p
23                 B[r,k] -= B[c,k]*p
24         # if DEBUG: print A, B
25     return B

```

```

26
27     def __div__(A,B):
28         if isinstance(B,Matrix):
29             return A*(1.0/B) # matrix/matrix
30         else:
31             return (1.0/B)*A # matrix/scalar

```

Here is an example, and we will see many more applications later:

Listing 4.22: in file: nlib.py

```

1 >>> A = Matrix([[1,2],[4,9]])
2 >>> print 1/A
3 [[9.0, -2.0], [-4.0, 1.0]]
4 >>> print A/A
5 [[1.0, 0.0], [0.0, 1.0]]
6 >>> print A/2
7 [[0.5, 1.0], [2.0, 4.5]]

```

4.4.4 Transposing a matrix

Another operation that we will need is transposition:

Listing 4.23: in file: nlib.py

```

1 @property
2 def T(A):
3     """Transposed of A"""
4     return Matrix(A.ncols,A.nrows, fill=lambda r,c: A[c,r])

```

Notice the new matrix is defined with the number of rows and columns switched from matrix A. Notice that in Python, a property is a method that is called like an attribute, therefore without the () notation. This can be used as follows:

Listing 4.24: in file: nlib.py

```

1 >>> A = Matrix([[1,2],[3,4]])
2 >>> print A.T
3 [[1, 3], [2, 4]]

```

For later use, we define two functions to check whether a matrix is symmetrical or zero within a given precision.

Another typical transformation for matrices of complex numbers is the Hermitian operation, which is a transposition combined with complex conjugation of the elements:

Listing 4.25: in file: nlib.py

```

1  @property
2  def H(A):
3      """Hermitian of A"""
4      return Matrix(A.ncols,A.nrows, fill=lambda r,c: A[c,r].conj())

```

In later algorithms we will need to check whether a matrix is symmetrical (or almost symmetrical given precision) or zero (or almost zero):

Listing 4.26: in file: nlib.py

```

1  def is_almost_symmetric(A, ap=1e-6, rp=1e-4):
2      if A.nrows != A.ncols: return False
3      for r in xrange(A.nrows):
4          for c in xrange(r):
5              delta = abs(A[r,c]-A[c,r])
6              if delta>ap and delta>max(abs(A[r,c]),abs(A[c,r]))*rp:
7                  return False
8      return True
9
10 def is_almost_zero(A, ap=1e-6, rp=1e-4):
11     for r in xrange(A.nrows):
12         for c in xrange(A.ncols):
13             delta = abs(A[r,c]-A[c,r])
14             if delta>ap and delta>max(abs(A[r,c]),abs(A[c,r]))*rp:
15                 return False
16     return True

```

4.4.5 Solving systems of linear equations

Linear algebra is fundamental for solving systems of linear equations such as the following:

$$x_0 + 2x_1 + 2x_2 = 3 \quad (4.58)$$

$$4x_0 + 4x_1 + 2x_2 = 6 \quad (4.59)$$

$$4x_0 + 6x_1 + 4x_2 = 10 \quad (4.60)$$

This can be rewritten using the equivalent linear algebra notation:

$$Ax = b \quad (4.61)$$

where

$$A = \begin{pmatrix} 1 & 2 & 2 \\ 4 & 4 & 2 \\ 4 & 6 & 4 \end{pmatrix} \quad \text{and} \quad b = \begin{pmatrix} 3 \\ 6 \\ 10 \end{pmatrix} \quad (4.62)$$

The solution of the equation can now be written as

$$x = A^{-1}b \quad (4.63)$$

We can easily solve the system with our Python library:

Listing 4.27: in file: nlib.py

```
1 >>> A = Matrix([[1,2,2],[4,4,2],[4,6,4]])
2 >>> b = Matrix([[3],[6],[10]])
3 >>> x = (1/A)*b
4 >>> print x
5 [[-1.0], [3.0], [-1.0]]
```

Notice that b is a column vector and therefore

```
1 >>> b = Matrix([[3],[6],[10]])
```

but not

```
1 >>> b = Matrix([[3,6,10]]) # wrong
```

We can also obtain a column vector by performing a transposition of a row vector:

```
1 >>> b = Matrix([[3,6,10]]).T
```

4.4.6 Norm and condition number again

By norm of a vector, we often refer to the 2-norm defined using the Pythagoras theorem:

$$\|x\|_2 = \sqrt{\sum_i x_i^2} \quad (4.64)$$

For a vector, we can define the p -norm as a generalization of the 2-norm:

$$\|x\|_p \equiv \left(\sum_i \text{abs}(x_i)^p \right)^{\frac{1}{p}} \quad (4.65)$$

We can extend the notation of a norm to any function that maps a vector into a vector, as follows:

$$\|f\|_p \equiv \max_x \|f(x)\|_p / \|x\|_p \quad (4.66)$$

An immediate application is to functions implemented as linear transformations:

$$\|A\|_p \equiv \max_x \|Ax\|_p / \|x\|_p \quad (4.67)$$

This can be difficult to compute in the general case, but it reduces to a simple formula for the 1-norm:

$$\|A\|_1 \equiv \max_j \sum_i \text{abs}(A_{ij}) \quad (4.68)$$

The 2-norm is difficult to compute for a matrix, but the 1-norm provides an approximation. It is computed by adding up the magnitude of the elements per each column and finding the maximum sum.

This allows us to define a generic function to compute the norm of lists, matrices/vectors, and scalars:

Listing 4.28: in file: `nlib.py`

```

1 def norm(A,p=1):
2     if isinstance(A,(list,tuple)):
3         return sum(abs(x)**p for x in A)**(1.0/p)
4     elif isinstance(A,Matrix):
5         if A.nrows==1 or A.ncols==1:
6             return sum(norm(A[r,c])**p \
7                 for r in xrange(A.nrows) \
8                 for c in xrange(A.ncols))**(1.0/p)
9         elif p==1:
10            return max([sum(norm(A[r,c]) \
11                for r in xrange(A.nrows)) \
12                for c in xrange(A.ncols)])
13        else:
14            raise NotImplementedError
15    else:
16        return abs(A)

```

Now we can implement a function that computes the condition number for ordinary functions as well as for linear transformations represented by a matrix:

Listing 4.29: in file: nlib.py

```

1 def condition_number(f,x=None,h=1e-6):
2     if callable(f) and not x is None:
3         return D(f,h)(x)*x/f(x)
4     elif isinstance(f,Matrix): # if is the Matrix
5         return norm(f)*norm(1/f)
6     else:
7         raise NotImplementedError

```

Here are some examples:

Listing 4.30: in file: nlib.py

```

1 >>> def f(x): return x*x-5.0*x
2 >>> print condition_number(f,1)
3 0.74999...
4 >>> A = Matrix([[1,2],[3,4]])
5 >>> print condition_number(A)
6 21.0

```

Having the norm for matrices also allows us to extend the definition of convergence of a Taylor series to a series of matrices:

Listing 4.31: in file: nlib.py

```

1 def exp(x,ap=1e-6,rp=1e-4,ns=40):
2     if isinstance(x,Matrix):
3         t = s = Matrix.identity(x.ncols)
4         for k in xrange(1,ns):
5             t = t*x/k # next term
6             s = s + t # add next term
7             if norm(t)<max(ap,norm(s)*rp): return s
8             raise ArithmeticError('no convergence')
9     elif type(x)==type(1j):
10        return cmath.exp(x)
11    else:
12        return math.exp(x)

```

Listing 4.32: in file: nlib.py

```

1 >>> A = Matrix([[1,2],[3,4]])
2 >>> print exp(A)
3 [[51.96..., 74.73...], [112.10..., 164.07...]]

```

4.4.7 Cholesky factorization

A matrix is said to be positive definite if $x^t Ax > 0$ for every $x \neq 0$.

If a matrix is symmetric and positive definite, then there exists a lower triangular matrix L such that $A = LL^t$. A lower triangular matrix is a matrix that has zeros above the diagonal elements.

The Cholesky algorithm takes a matrix A as input and returns the matrix L :

Listing 4.33: in file: nlib.py

```

1 def Cholesky(A):
2     import copy, math
3     if not is_almost_symmetric(A):
4         raise ArithmeticError('not symmetric')
5     L = copy.deepcopy(A)
6     for k in xrange(L.ncols):
7         if L[k,k]<=0:
8             raise ArithmeticError('not positive definite')
9         p = L[k,k] = math.sqrt(L[k,k])
10        for i in xrange(k+1,L.nrows):
11            L[i,k] /= p
12        for j in xrange(k+1,L.nrows):
13            p=float(L[j,k])
14            for i in xrange(k+1,L.nrows):
15                L[i,j] -= p*L[i,k]
16    for i in xrange(L.nrows):
17        for j in xrange(i+1,L.ncols):
18            L[i,j]=0
19    return L

```

Here we provide an example and a check that indeed $A = LL^t$:

Listing 4.34: in file: nlib.py

```

1 >>> A = Matrix([[4,2,1],[2,9,3],[1,3,16]])
2 >>> L = Cholesky(A)
3 >>> print is_almost_zero(A - L*L.T)
4 True

```

The Cholesky algorithm fails if and only if the input matrix is not symmetric or not positive definite, therefore it can be used to check whether a symmetric matrix is positive definite.

Consider for example a generic covariance matrix A . It is supposed to be

positive definite, but sometimes it is not, because it is computed incorrectly by taking different subsets of the data to compute A_{ij} , A_{jk} , and A_{ik} . The Cholesky algorithm provides an algorithm to check whether a matrix is positive definite:

Listing 4.35: in file: nlib.py

```

1 def is_positive_definite(A):
2     if not is_almost_symmetric(A):
3         return False
4     try:
5         Cholesky(A)
6         return True
7     except Exception:
8         return False

```

Another application of the Cholesky is in generating vectors \mathbf{x} with probability distribution

$$p(\mathbf{x}) \propto \exp\left(-\frac{1}{2}\mathbf{x}^t A^{-1} \mathbf{x}\right) \quad (4.69)$$

where A is a symmetric and positive definite matrix. In fact, if $A = LL^t$, then

$$p(\mathbf{x}) \propto \exp\left(-\frac{1}{2}(L^{-1}\mathbf{x})^t(L^{-1}\mathbf{x})\right) \quad (4.70)$$

and with a change of variable $\mathbf{u} = L^{-1}\mathbf{x}$, we obtain

$$p(\mathbf{x}) \propto \exp\left(-\frac{1}{2}\mathbf{u}^t \mathbf{u}\right) \quad (4.71)$$

and

$$p(\mathbf{x}) \propto e^{-\frac{1}{2}u_0^2} e^{-\frac{1}{2}u_1^2} e^{-\frac{1}{2}u_2^2} \dots \quad (4.72)$$

Therefore the u_i components are Gaussian random variables.

In summary, given a covariance matrix A , we can generate random vectors \mathbf{x} or random numbers with the same covariance simply by doing

```

1 def RandomList(A):
2     L = Cholesky(A)
3     while True:
4         u = Matrix([[random.gauss(0,1)] for c in xrange(L.nrows)])
5         yield (L*u).flatten()

```

Here is an example of how to use it:

```

1 >>> A = Matrix([[1.0,0.1],[0.2,3.0]])
2 >>> for k, v in enumerate(RandomList(A)):
3 ...     print v

```

The `RandomList` is a generator. You can iterate over it. The `yield` keyword is used like `return`, except the function will return a generator.

4.4.8 Modern portfolio theory

Modern portfolio theory [34] is an investment approach that tries to maximize return given a fixed risk. Many different metrics have been proposed. One of them is the *Sharpe ratio*.

For a stock or a portfolio with an average return r and risk σ , the Sharpe ratio is defined as

$$\text{Sharpe}(r, \sigma) \equiv (r - \bar{r}) / \sigma \quad (4.73)$$

Here \bar{r} is the current risk-free investment rate. Usually the risk is measured as the standard deviation of its daily (or monthly or yearly) return.

Consider the stock price p_{it} of stock i at time t and its arithmetic daily return $r_{it} = (p_{i,t+1} - p_{it}) / p_{it}$ given a risk-free interest equal to \bar{r} .

For each stock, we can compute the average return and average risk (variance of daily returns) and display it in a risk-return plot as we did in chapter 2.

We can try to build arbitrary portfolios by investing in multiple stocks at the same time. Modern portfolio theory states that there is a maximum Sharpe ratio and there is one portfolio that corresponds to it. It is called the tangency portfolio.

A portfolio is identified by fractions of \$1 invested in each stock in the portfolio. Our goal is to determine the tangent portfolio.

If we assume that daily returns for the stocks are Gaussian, then the solving algorithm is simple.

All we need is to compute the average return for each stock, defined as

$$r_i = 1/T \sum_t r_{it} \quad (4.74)$$

and the covariance matrix

$$A_{ij} = \frac{1}{T} \sum_t (r_{it} - r_i)(r_{jt} - r_j) \quad (4.75)$$

Modern portfolio theory tells us that the tangent portfolio is given by

$$\mathbf{x} = A^{-1}(\mathbf{r} - \bar{r}\mathbf{1}) \quad (4.76)$$

The inputs of the formula are the covariance matrix (A), a vector of arithmetic returns for the assets in the portfolio (r), the risk free rate (\bar{r}). The output is a vector (x) whose elements are the percentages to be invested in each asset to obtain a tangency portfolio. Notice that some elements of x can be negative and this corresponds to short position (sell, not buy, the asset).

Here is the algorithm:

Listing 4.36: in file: nlib.py

```

1 def Markowitz(mu, A, r_free):
2     """Assess Markowitz risk/return.
3     Example:
4     >>> cov = Matrix([[0.04, 0.006, 0.02],
5     ...               [0.006, 0.09, 0.06],
6     ...               [0.02, 0.06, 0.16]])
7     >>> mu = Matrix([[0.10], [0.12], [0.15]])
8     >>> r_free = 0.05
9     >>> x, ret, risk = Markowitz(mu, cov, r_free)
10    >>> print x
11    [0.556634..., 0.275080..., 0.1682847...]
12    >>> print ret, risk
13    0.113915... 0.186747...
14    """
15    x = Matrix([[0.0] for r in xrange(A.nrows)])
16    x = (1/A)*(mu - r_free)
17    x = x/sum(x[r,0] for r in xrange(x.nrows))
18    portfolio = [x[r,0] for r in xrange(x.nrows)]
19    portfolio_return = mu*x
20    portfolio_risk = sqrt(x*(A*x))
21    return portfolio, portfolio_return, portfolio_risk

```

Here is an example. We consider three assets (0,1,2) with the following covariance matrix:

```
1 >>> cov = Matrix([[0.04, 0.006, 0.02],
2 ...               [0.006, 0.09, 0.06],
3 ...               [0.02, 0.06, 0.16]])
```

and the following expected returns (arithmetic returns, not log returns, because the former are additive, whereas the latter are not):

```
1 >>> mu = Matrix([[0.10], [0.12], [0.15]])
```

Given the risk-free interest rate

```
1 >>> r_free = 0.05
```

we compute the tangent portfolio (highest Sharpe ratio), its return, and its risk with one function call:

```
1 >>> x, ret, risk = Markowitz(mu, cov, r_free)
2 >>> print x
3 [0.5566343042071198, 0.27508090614886727, 0.16828478964401297]
4 >>> print ret, risk
5 0.113915857605 0.186747095412
6 >>> print (ret-r_free).risk
7 0.34225891152
8 >>> for r in xrange(3): print (mu[r,0]-r_free)/sqrt(cov[r,r])
9 0.25
10 0.233333333333
11 0.25
```

Investing 55% in asset 0, 27% in asset 1, and 16% in asset 2, the resulting portfolio has an expected return of 11.39% and a risk of 18.67%, which corresponds to a Sharpe ratio of 0.34, much higher than 0.25, 0.23, and 0.23 for the individual assets.

Notice that the tangency portfolio is not the only one with the highest Sharpe ratio (return for unit of risk). In fact, any linear combination of the tangency portfolio with a risk-free asset (putting money in the bank) has the same Sharpe ratio. For any target risk, one can find a linear combination of the risk-free asset and the tangent portfolio that has a better Sharpe ratio than any other possible portfolio comprising the same assets.

If we call α the fraction of the money to invest in the tangency portfolio and $1 - \alpha$ the fraction to keep in the bank at the risk free rate, the resulting

combined portfolio has return:

$$\alpha \mathbf{x} \cdot \mathbf{r} + (1 - \alpha) \bar{r} \quad (4.77)$$

and risk

$$\alpha \sqrt{\mathbf{x}^t A \mathbf{x}} \quad (4.78)$$

We can determine α by deciding how much risk we are willing to take, and these formulas tell us the optimal portfolio for that amount of risk.

4.4.9 Linear least squares, χ^2

Consider a set of data points $(x_j, y_j) = (t_j, o_j \pm do_j)$. We want to fit them with a linear combination of linear independent functions f_i so that

$$c_0 f_0(t_0) + c_1 f_1(t_0) + c_2 f_2(t_0) + \dots = e_0 \simeq o_0 \pm do_0 \quad (4.79)$$

$$c_0 f_0(t_1) + c_1 f_1(t_1) + c_2 f_2(t_1) + \dots = e_1 \simeq o_1 \pm do_1 \quad (4.80)$$

$$c_0 f_0(t_2) + c_1 f_1(t_2) + c_2 f_2(t_2) + \dots = e_2 \simeq o_2 \pm do_2 \quad (4.81)$$

$$\dots = \dots \quad (4.82)$$

We want to find the $\{c_i\}$ that minimizes the sum of the squared distances between the actual “observed” data o_j and the predicted “expected” data e_j , in units of do_j . This metric is called χ^2 in general [35]. An algorithm that minimizes the χ^2 and is linear in the c_i coefficients (our case here) is called *linear least squares* or *linear regression*.

$$\chi^2 = \sum_j \left| \frac{e_j - o_j}{do_j} \right|^2 \quad (4.83)$$

If we define the matrix A and B as

$$A = \begin{pmatrix} \frac{f_0(t_0)}{do_0} & \frac{f_1(t_0)}{do_0} & \frac{f_2(t_0)}{do_0} & \cdots \\ \frac{f_0(t_1)}{do_1} & \frac{f_1(t_1)}{do_1} & \frac{f_2(t_1)}{do_1} & \cdots \\ \frac{f_0(t_2)}{do_2} & \frac{f_1(t_2)}{do_2} & \frac{f_2(t_2)}{do_2} & \cdots \\ \cdots & \cdots & \cdots & \cdots \end{pmatrix} \quad b = \begin{pmatrix} o_0 \\ do_0 \\ o_1 \\ do_1 \\ o_2 \\ do_2 \\ \cdots \end{pmatrix} \quad (4.84)$$

then the problem is reduced to

$$\min_{\mathbf{c}} \chi^2 = \min_{\mathbf{c}} |\mathbf{Ac} - \mathbf{b}|^2 \quad (4.85)$$

$$= \min_{\mathbf{c}} (\mathbf{Ac} - \mathbf{b})^t (\mathbf{Ac} - \mathbf{b}) \quad (4.86)$$

$$= \min_{\mathbf{c}} (\mathbf{c}^t \mathbf{A}^t \mathbf{Ax} - 2\mathbf{b}^t \mathbf{Ac} + \mathbf{b}^t \mathbf{b}) \quad (4.87)$$

This is the same as solving the following equation:

$$\nabla_{\mathbf{c}} (\mathbf{c}^t \mathbf{A}^t \mathbf{Ax} - 2\mathbf{b}^t \mathbf{Ac} + \mathbf{b}^t \mathbf{b}) = 0 \quad (4.88)$$

$$\mathbf{A}^t \mathbf{Ac} - \mathbf{A}^t \mathbf{b} = 0 \quad (4.89)$$

Its solution is

$$\mathbf{c} = (\mathbf{A}^t \mathbf{A})^{-1} (\mathbf{A}^t \mathbf{b}) \quad (4.90)$$

The following algorithm implements a fitting function based on the preceding procedure. It takes as input a list of functions f_i and a list of points $p_j = (t_j, o_j, do_j)$ and returns three objects—a list with the c coefficients, the value of χ^2 for the best fit, and the fitting function:

Listing 4.37: in file: nlib.py

```

1 def fit_least_squares(points, f):
2     """
3     Computes c_j for best linear fit of y[i] \pm dy[i] = fitting_f(x[i])
4     where fitting_f(x[i]) is \sum_j c_j f[j](x[i])
5 
```

```

6  parameters:
7  - a list of fitting functions
8  - a list with points (x,y,dy)
9
10 returns:
11 - column vector with fitting coefficients
12 - the chi2 for the fit
13 - the fitting function as a lambda x: ....
14 """
15 def eval_fitting_function(f,c,x):
16     if len(f)==1: return c*f[0](x)
17     else: return sum(func(x)*c[i,0] for i,func in enumerate(f))
18 A = Matrix(len(points),len(f))
19 b = Matrix(len(points))
20 for i in xrange(A.nrows):
21     weight = 1.0/points[i][2] if len(points[i])>2 else 1.0
22     b[i,0] = weight*float(points[i][1])
23     for j in xrange(A.ncols):
24         A[i,j] = weight*f[j](float(points[i][0]))
25 c = (1.0/(A.T*A))*(A.T*b)
26 chi = A*c-b
27 chi2 = norm(chi,2)**2
28 fitting_f = lambda x, c=c, f=f, q=eval_fitting_function: q(f,c,x)
29 if isinstance(c,Matrix): return c.flatten(), chi2, fitting_f
30 else: return c, chi2, fitting_f
31
32 # examples of fitting functions
33 def POLYNOMIAL(n):
34     return [(lambda x, p=p: x**p) for p in xrange(n+1)]
35 CONSTANT = POLYNOMIAL(0)
36 LINEAR = POLYNOMIAL(1)
37 QUADRATIC = POLYNOMIAL(2)
38 CUBIC = POLYNOMIAL(3)
39 QUARTIC = POLYNOMIAL(4)

```

As an example, we can use it to perform a polynomial fit: given a set of points, we want to find the coefficients of a polynomial that best approximate those points.

In other words, we want to find the c_i such that, given t_j and o_j ,

$$c_0 + c_1 t_0^1 + c_2 t_0^2 + \dots = e_0 \simeq o_0 \pm do_0 \quad (4.91)$$

$$c_0 + c_1 t_1^1 + c_2 t_1^2 + \dots = e_1 \simeq o_1 \pm do_1 \quad (4.92)$$

$$c_0 + c_1 t_2^1 + c_2 t_2^2 + \dots = e_2 \simeq o_2 \pm do_2 \quad (4.93)$$

$$\dots = \dots \quad (4.94)$$

Here is how we can generate some random points and solve the problem for a polynomial of degree 2 (or quadratic fit):

Listing 4.38: in file: nlib.py

```

1 >>> points = [(k,5+0.8*k+0.3*k*k+math.sin(k),2) for k in xrange(100)]
2 >>> a,chi2,fitting_f = fit_least_squares(points,QUADRATIC)
3 >>> for p in points[-10:]:
4 ...     print p[0], round(p[1],2), round(fitting_f(p[0]),2)
5 90 2507.89 2506.98
6 91 2562.21 2562.08
7 92 2617.02 2617.78
8 93 2673.15 2674.08
9 94 2730.75 2730.98
10 95 2789.18 2788.48
11 96 2847.58 2846.58
12 97 2905.68 2905.28
13 98 2964.03 2964.58
14 99 3023.5 3024.48
15 >>> Canvas(title='polynomial fit',xlab='t',ylab='e(t),o(t)'
16 ...     ).errorbar(points[:10],legend='o(t)'
17 ...     ).plot([(p[0],fitting_f(p[0])) for p in points[:10]],legend='e(t)'
18 ...     ).save('images/polynomialfit.png')

```

Fig. 4.4.9 is a plot of the first 10 points compared with the best fit:




Figure 4.4: Random data with their error bars and the polynomial best fit.

We can also define a $\chi_{dof}^2 = \chi^2/(N-1)$ where N is the number of c parameters determined by the fit. A value of $\chi_{dof}^2 \simeq 1$ indicates a good fit. In general, the smaller χ_{dof}^2 , the better the fit. A large value of χ_{dof}^2 is a symptom of poor modeling (the assumptions of the fit are wrong), whereas a value χ_{dof}^2 much smaller than 1 is a symptom of an overestimate of the errors do_j (or perhaps manufactured data).

4.4.10 Trading and technical analysis

In finance, *technical analysis* is an empirical discipline that consists of forecasting the direction of prices through the study of patterns in historical data (in particular, price and volume). As an example, we implement a simple strategy that consists of the following steps:

- We fit the adjusted closing price for the previous seven days and use our fitting function to predict the adjusted close for the next day.
- If we have cash and predict the price will go up, we buy the stock.
- If we hold the stock and predict the price will go down, we sell the stock.

Listing 4.39: in file: nlib.py

```

1 class Trader:
2     def model(self, window):
3         "the forecasting model"
4         # we fit last few days quadratically
5         points = [(x, y['adjusted_close']) for (x, y) in enumerate(window)]
6         a, chi2, fitting_f = fit_least_squares(points, QUADRATIC)
7         # and we extrapolate tomorrow's price
8         tomorrow_prediction = fitting_f(len(points))
9         return tomorrow_prediction
10
11     def strategy(self, history, ndays=7):
12         "the trading strategy"
13         if len(history) < ndays:
14             return
15         else:
16             today_close = history[-1]['adjusted_close']
17             tomorrow_prediction = self.model(history[-ndays:])
18             return 'buy' if tomorrow_prediction > today_close else 'sell'
19 
```

```

20 def simulate(self, data, cash=1000.0, shares=0.0, daily_rate=0.03/360):
21     "find fitting parameters that optimize the trading strategy"
22     for t in xrange(len(data)):
23         suggestion = self.strategy(data[:t])
24         today_close = data[t-1]['adjusted_close']
25         # and we buy or sell based on our strategy
26         if cash>0 and suggestion=='buy':
27             # we keep track of finances
28             shares_bought = int(cash/today_close)
29             shares += shares_bought
30             cash -= shares_bought*today_close
31         elif shares>0 and suggestion=='sell':
32             cash += shares*today_close
33             shares = 0.0
34             # we assume money in the bank also gains an interest
35             cash*=math.exp(daily_rate)
36         # we return the net worth
37     return cash+shares*data[-1]['adjusted_close']

```

Now we back test the strategy using financial data for AAPL for the year 2011:

Listing 4.40: in file: nlib.py

```

1 >>> from datetime import date
2 >>> data = YStock('aapl').historical(
3 ...     start=date(2011,1,1), stop=date(2011,12,31))
4 >>> print Trader().simulate(data, cash=1000.0)
5 1120...
6 >>> print 1000.0*math.exp(0.03)
7 1030...
8 >>> print 1000.0*data[-1]['adjusted_close']/data[0]['adjusted_close']
9 1228...

```

Our strategy did considerably better than the risk-free return of 3% but not as well as investing and holding AAPL shares over the same period.

Of course, we can always engineer a strategy based on historical data that will outperform holding the stock, but *past performance is never a guarantee of future performance*.

According to the definition from investopedia.com, “technical analysts believe that the historical performance of stocks and markets is an indication of future performance.”

The efficacy of both technical and fundamental analysis is disputed by the efficient-market hypothesis, which states that stock market prices are

essentially unpredictable [36].

It is easy to extend the previous class to implement other strategies and back test them.

4.4.11 Eigenvalues and the Jacobi algorithm

Given a matrix A , an eigenvector is defined as a vector \mathbf{x} such that $A\mathbf{x}$ is proportional to \mathbf{x} . The proportionality factor is called an eigenvalue, e . One matrix may have many eigenvectors \mathbf{x}_i and associated eigenvalues e_i :

$$A\mathbf{x}_i = e_i\mathbf{x}_i \quad (4.95)$$

For example:

$$A = \begin{pmatrix} 1 & -2 \\ 1 & 4 \end{pmatrix} \quad \text{and} \quad \mathbf{x}_i = \begin{pmatrix} -1 \\ 1 \end{pmatrix} \quad (4.96)$$

$$\begin{pmatrix} 1 & -2 \\ 1 & 4 \end{pmatrix} * \begin{pmatrix} -1 \\ 1 \end{pmatrix} = 3 * \begin{pmatrix} -1 \\ 1 \end{pmatrix} \quad (4.97)$$

In this case, \mathbf{x}_i is an eigenvector and the corresponding eigenvalue is $e = 3$.

Some eigenvalues may be zero ($e_i = 0$), which means the matrix A is singular. A matrix is singular if it maps a nonzero vector into zero.

Given a square matrix A , if the space generated by the linear independent eigenvectors has the same dimensionality as the number of rows (or columns) of A , then its eigenvalues are real and the matrix can be written as

$$A = UDU^t \quad (4.98)$$

where D is a diagonal matrix with eigenvalues on the diagonal $D_{ii} = e_i$ and U is a matrix whose column i is the \mathbf{x}_i eigenvector.

The following algorithm is called the Jacobi algorithm. It takes as input a symmetric matrix A and returns the matrix U and a list of corresponding eigenvalues e , sorted from smallest to largest:

Listing 4.41: in file: nlib.py

```

1 def sqrt(x):
2     try:
3         return math.sqrt(x)
4     except ValueError:
5         return cmath.sqrt(x)
6
7 def Jacobi_eigenvalues(A,checkpoint=False):
8     """Returns U and e so that  $A=U \cdot \text{diagonal}(e) \cdot \text{transposed}(U)$ 
9         where  $i$ -column of  $U$  contains the eigenvector corresponding to
10        the eigenvalue  $e[i]$  of  $A$ .
11
12        from http://en.wikipedia.org/wiki/Jacobi\_eigenvalue\_algorithm
13    """
14    def maxind(M,k):
15        j=k+1
16        for i in xrange(k+2,M.ncols):
17            if abs(M[k,i])>abs(M[k,j]):
18                j=i
19        return j
20    n = A.nrows
21    if n!=A.ncols:
22        raise ArithmeticError('matrix not squared')
23    indexes = xrange(n)
24    S = Matrix(n,n, fill=lambda r,c: float(A[r,c]))
25    E = Matrix.identity(n)
26    state = n
27    ind = [maxind(S,k) for k in indexes]
28    e = [S[k,k] for k in indexes]
29    changed = [True for k in indexes]
30    iteration = 0
31    while state:
32        if checkpoint: checkpoint('rotating vectors (%i) ...' % iteration)
33        m=0
34        for k in xrange(1,n-1):
35            if abs(S[k,ind[k]])>abs(S[m,ind[m]]): m=k
36            pass
37        k,h = m,ind[m]
38        p = S[k,h]
39        y = (e[h]-e[k])/2
40        t = abs(y)+sqrt(p*p+y*y)
41        s = sqrt(p*p+t*t)
42        c = t/s

```



```

43     s = p/s
44     t = p*p/t
45     if y<0: s,t = -s,-t
46     S[k,h] = 0
47     y = e[k]
48     e[k] = y-t
49     if changed[k] and y==e[k]:
50         changed[k],state = False,state-1
51     elif (not changed[k]) and y!=e[k]:
52         changed[k],state = True,state+1
53     y = e[h]
54     e[h] = y+t
55     if changed[h] and y==e[h]:
56         changed[h],state = False,state-1
57     elif (not changed[h]) and y!=e[h]:
58         changed[h],state = True,state+1
59     for i in xrange(k):
60         S[i,k],S[i,h] = c*S[i,k]-s*S[i,h],s*S[i,k]+c*S[i,h]
61     for i in xrange(k+1,h):
62         S[k,i],S[i,h] = c*S[k,i]-s*S[i,h],s*S[k,i]+c*S[i,h]
63     for i in xrange(h+1,n):
64         S[k,i],S[h,i] = c*S[k,i]-s*S[h,i],s*S[k,i]+c*S[h,i]
65     for i in indexes:
66         E[k,i],E[h,i] = c*E[k,i]-s*E[h,i],s*E[k,i]+c*E[h,i]
67     ind[k],ind[h]=maxind(S,k),maxind(S,h)
68     iteration+=1
69     # sort vectors
70     for i in xrange(1,n):
71         j=i
72         while j>0 and e[j-1]>e[j]:
73             e[j],e[j-1] = e[j-1],e[j]
74             E.swap_rows(j,j-1)
75         j-=1
76     # normalize vectors
77     U = Matrix(n,n)
78     for i in indexes:
79         norm = sqrt(sum(E[i,j]**2 for j in indexes))
80         for j in indexes: U[j,i] = E[i,j]/norm
81     return U,e

```

Here is an example that shows, for a particular case, the relation between the input, A , of the output of the U, e of the Jacobi algorithm:

Listing 4.42: in file: nlib.py

```

1 >>> import random
2 >>> A = Matrix(4,4)
3 >>> for r in xrange(A.nrows):
4 ...     for c in xrange(r,A.ncols):

```

```

5 ...     A[r,c] = A[c,r] = random.gauss(10,10)
6 >>> U,e = Jacobi_eigenvalues(A)
7 >>> print is_almost_zero(U*Matrix.diagonal(e)*U.T-A)
8 True

```

Eigenvalues can be used to filter noise out of data and find hidden dependencies in data. Following are some examples.

4.4.12 Principal component analysis

One important application of the Jacobi algorithm is for principal component analysis (PCA). This is a mathematical procedure that converts a set of observations of possibly correlated vectors into a set of uncorrelated vectors called *principal components*.

Here we consider, as an example, the time series of the adjusted arithmetic returns for the S&P100 stocks that we downloaded and stored in chapter 2.

Each time series is a vector. We know they are not independent because there are correlations. Our goal is to model each time series and a vector plus noise where the vector is the same for all series. We also want find that vector that has maximal superposition with the individual time series, the principal component.

First, we compute the correlation matrix for all the stocks. This is a non-trivial task because we have to make sure that we only consider those days when all stocks were traded:

Listing 4.43: in file: nlib.py

```

1 def compute_correlation(stocks, key='arithmetic_return'):
2     "The input must be a list of YStock(...).historical() data"
3     # find trading days common to all stocks
4     days = set()
5     nstocks = len(stocks)
6     iter_stocks = xrange(nstocks)
7     for stock in stocks:
8         if not days: days=set(x['date'] for x in stock)
9         else: days=days.intersection(set(x['date'] for x in stock))
10    n = len(days)
11    v = []
12    # filter out data for the other days

```

```

13  for stock in stocks:
14      v.append([x[key] for x in stock if x['date'] in days])
15      # compute mean returns (skip first day, data not reliable)
16      mus = [sum(v[i][k] for k in xrange(1,n))/n for i in iter_stocks]
17      # fill in the covariance matrix
18      var = [sum(v[i][k]**2 for k in xrange(1,n))/n - mus[i]**2 for i in
              iter_stocks]
19      corr = Matrix(nstocks,nstocks,fill=lambda i,j: \
20                  (sum(v[i][k]*v[j][k] for k in xrange(1,n))/n - mus[i]*mus[j])/ \
21                  math.sqrt(var[i]*var[j]))
22  return corr

```

We use the preceding function to compute the correlation and pass it as input to the Jacobi algorithm and plot the output eigenvalues:

Listing 4.44: in file: nlib.py

```

1  >>> storage = PersistentDictionary('sp100.sqlite')
2  >>> symbols = storage.keys('*/*2011')[:20]
3  >>> stocks = [storage[symbol] for symbol in symbols]
4  >>> corr = compute_correlation(stocks)
5  >>> U,e = Jacobi_eigenvalues(corr)
6  >>> Canvas(title='SP100 eigenvalues',xlab='i',ylab='e[i]'
7  ...        ).plot([(i,ei) for i,ei in enumerate(e)])
8  ...        ).save('images/sp100eigen.png')

```

The image shows that one eigenvalue, the last one, is much larger than the others. It tells us that the data series have something in common. In fact, the arithmetic returns for stock j at time t can be written as

$$r_{it} = \beta_i p_t + \alpha_{it} \quad (4.99)$$

where p is the principal component given by

$$p_t = \sum_i U_{n-1,j} r_{jt} \quad (4.100)$$

$$\beta_i = \sum_t r_{it} p_t \quad (4.101)$$

$$\alpha_{it} = r_{it} - \beta_i p_t \quad (4.102)$$

Here \mathbf{p} is the vector of adjusted arithmetic returns that better correlates with the returns of the individual assets and therefore best represents the




Figure 4.5: Eigenvalues of the correlation matrix for 20 of the S&P100 stocks, sorted by their magnitude.

market. The β_i coefficient tells us how much \mathbf{r}_i overlaps with \mathbf{p} ; α , at first approximation, measures leftover noise.

4.5 Sparse matrix inversion

Sometimes we have to invert matrices that are very large, and the Gauss-Jordan algorithms fails. Yet, if the matrix is sparse, in the sense that most of its elements are zeros, than two algorithms come to our rescue: the *minimum residual* and the *biconjugate gradient* (for which we consider a variant called the *stabilized bi-conjugate gradient*).

We will also assume that the matrix to be inverted is given in some implicit algorithmic as $\mathbf{y} = f(\mathbf{x})$ because this is always the case for sparse matrices. There is no point to storing all its elements because most of them are zero.

4.5.1 Minimum residual

Given a linear operator f , the Krylov space spanned by a vector x is defined as

$$K(f, y, i) = \{y, f(y), f(f(y)), f(f(f(y))), (f^i)(y)\} \quad (4.103)$$

The *minimum residual* [37] algorithm works by solving $x = f^{-1}(y)$ iteratively. At each iteration, it computes a new orthogonal basis vector q_i for the Krylov space $K(f, y, i)$ and computes the coefficients α_i that project x_i into component i of the Krylov space:

$$x_i = y + \alpha_1 q_1 + \alpha_2 q_2 + \cdots + \alpha_i q_i \in K(f, y, i + 1) \quad (4.104)$$

which minimizes the norm of the residue defined as:

$$r = f(x_i) - y \quad (4.105)$$

Therefore $\lim_{i \rightarrow \infty} f(x_i) = y$. If a solution to the original problem exists, ignoring precision issues, the minimum residual converges to it, and the residue decreases at each iteration.

Notice that in the following code, x and y are exchanged because we adopt the convention that y is the output and x is the input:

Listing 4.45: in file: `nlib.py`

```

1 def invert_minimum_residual(f,x,ap=1e-4,rp=1e-4,ns=200):
2     import copy
3     y = copy.copy(x)
4     r = x-1.0*f(x)
5     for k in xrange(ns):
6         q = f(r)
7         alpha = (q*r)/(q*q)
8         y = y + alpha*r
9         r = r - alpha*q
10        residue = sqrt((r*r)/r.nrows)
11        if residue<max(ap,norm(y)*rp): return y
12    raise ArithmeticError('no convergence')
```

4.5.2 Stabilized biconjugate gradient

The stabilized biconjugate gradient [38] method is also based on constructing a Krylov subspace and minimizing the same residue as in the minimum residual algorithm, yet it is faster than the minimum residual and has a smoother convergence than other conjugate gradient methods:

Listing 4.46: in file: `nlib.py`

```

1 def invert_bicgstab(f,x,ap=1e-4,rp=1e-4,ns=200):
2     import copy
3     y = copy.copy(x)
4     r = x - 1.0*f(x)
5     q = r
6     p = 0.0
7     s = 0.0
8     rho_old = alpha = omega = 1.0
9     for k in xrange(ns):
10        rho = q*r
11        beta = (rho/rho_old)*(alpha/omega)
12        rho_old = rho
13        p = beta*p + r - (beta*omega)*s
14        s = f(p)
15        alpha = rho/(q*s)
16        r = r - alpha*s
17        t = f(r)
18        omega = (t*r)/(t*t)
19        y = y + omega*r + alpha*p
20        residue=sqrt((r*r)/r.nrows)
21        if residue<max(ap,norm(y)*rp): return y
22        r = r - omega*t
23    raise ArithmeticError('no convergence')
```

Notice that the minimum residual and the stabilized biconjugate gradient, if they converge, converge to the same value.

As an example, consider the following. We take a picture using a camera, but we take the picture out of focus. The image is represented by a set of m^2 pixels. The defocusing operation can be modeled as a first approximation with a linear operator acting on the “true” image, x , and turning it into an “out of focus” image, y . We can store the pixels in a one-dimensional vector (both for x and y) as opposed to a matrix by mapping the pixel (r, c) into vector component i using the relation $(r, c) = (i/m, i\%m)$.

Hence we can write

$$\mathbf{y} = A\mathbf{x} \quad (4.106)$$

Here the linear operator A represents the effects of the lens, which transforms one set of pixels into another.

We can model the lens as a sequence of β smearing operators:

$$A = S^\beta \quad (4.107)$$

where a smearing operator is a next neighbor interaction among pixels:

$$S_{ij} = (1 - \alpha/4)\delta_{i,j} + \alpha\delta_{i,j\pm 1} + \alpha\delta_{i,j\pm m} \quad (4.108)$$

Here α and β are smearing coefficients. When $\alpha = 0$ or $\beta = 0$, the lens has no effect, and $A = I$. The value of α controls how much the value of light at point i is averaged with the value at its four neighbor points: left ($j - 1$), right ($j + 1$), top ($j + m$), and bottom ($j - m$). The coefficient β determines the width of the smearing radius. The larger the values of β and α , the more out of focus is the original image.

In the following code, we generate an image x and filter it through a lens operator `smear`, obtaining a smeared image y . We then use the sparse matrix inverter to reconstruct the original image x given the smeared image y . We use the `color2d` plotting function to represent the images:

Listing 4.47: in file: `nlib.py`

```

1 >>> m = 30
2 >>> x = Matrix(m*m,1,fill=lambda r,c:(r//m in(10,20) or r%m in(10,20)) and 1. or
   0.)
3 >>> def smear(x):
4 ...     alpha, beta = 0.4, 8
5 ...     for k in xrange(beta):
6 ...         y = Matrix(x.nrows,1)
7 ...         for r in xrange(m):
8 ...             for c in xrange(m):
9 ...                 y[r*m+c,0] = (1.0-alpha/4)*x[r*m+c,0]
10 ...                 if c<m-1: y[r*m+c,0] += alpha * x[r*m+c+1,0]
```

```

11 ...         if c>0: y[r*m+c,0] += alpha * x[r*m+c-1,0]
12 ...         if r<m-1: y[r*m+c,0] += alpha * x[r*m+c+m,0]
13 ...         if c>0: y[r*m+c,0] += alpha * x[r*m+c-m,0]
14 ...     x = y
15 ...     return y
16 >>> y = smear(x)
17 >>> z = invert_minimum_residual(smear,y,ns=1000)
18 >>> y = y.reshape(m,m)
19 >>> Canvas(title="Defocused image").imshow(y.tolist()).save('images/defocused.
    png')
20 >>> Canvas(title="refocus image").imshow(z.tolist()).save('images/refocused.png'
    )

```




Figure 4.6: An out-of-focus image (left) and the original image (image) computed from the out-of-focus one, using sparse matrix inversion.

When the Hubble telescope was first put into orbit, its mirror was not installed properly and caused the telescope to take pictures out of focus. Until the defect was physically corrected, scientists were able to fix the images using a similar algorithm.

4.6 Solvers for nonlinear equations

In this chapter, we are concerned with the problem of solving in x the equation of one variable:

$$f(x) = 0 \quad (4.109)$$

4.6.1 Fixed-point method

It is always possible to reformulate $f(x) = 0$ as $g(x) = x$ using, for example, one of the following definitions:

- $g(x) = f(x)/c + x$ for some constant c
- $g(x) = f(x)/q(x) + x$ for some $q(x) > 0$ at the solution of $f(x) = 0$

We start at x_0 , an arbitrary point in the domain, and close to the solution we seek. We compute

$$x_1 = g(x_0) \quad (4.110)$$

$$x_2 = g(x_1) \quad (4.111)$$

$$x_3 = g(x_2) \quad (4.112)$$

$$\dots = \dots \quad (4.113)$$

We can compute the distance between x_i and x as

$$|x_i - x| = |g(x_{i-1}) - g(x)| \quad (4.114)$$

$$= |g(x) + g'(\xi)(x_{i-1} - x) - g(x)| \quad (4.115)$$

$$= |g'(\xi)||x_{i-1} - x| \quad (4.116)$$

where we use *L'Hôpital's rule* and ξ is a point in between x and x_{i-1} .

If the magnitude of the first derivative of g , $|g'|$, is less than 1 in a neighborhood of x , and if x_0 is in such a neighborhood, then

$$|x_i - x| = |g'(\xi)||x_{i-1} - x| < |x_{i-1} - x| \quad (4.117)$$

The x_i series will get closer and closer to the solution x .

Here is the process implemented into an algorithm:

Listing 4.48: in file: nlib.py

```

1 def solve_fixed_point(f, x, ap=1e-6, rp=1e-4, ns=100):
2     def g(x): return f(x)+x # f(x)=0 <=> g(x)=x
3     Dg = D(g)
4     for k in xrange(ns):
5         if abs(Dg(x)) >= 1:
```

```

6         raise ArithmeticError('error  $D(g)(x) \geq 1$ ')
7         (x_old, x) = (x, g(x))
8         if k>2 and norm(x_old-x)<max(ap,norm(x)*rp):
9             return x
10        raise ArithmeticError('no convergence')

```

And here is an example:

Listing 4.49: in file: nlib.py

```

1 >>> def f(x): return (x-2)*(x-5)/10
2 >>> print round(solve_fixed_point(f,1.0,rp=0),4)
3 2.0

```

4.6.2 Bisection method

The goal of the bisection [39] method is to solve $f(x) = 0$ when the function is continuous and it is known to change sign in between $x = a$ and $x = b$. The bisection method is the continuous equivalent of the binary search algorithm seen in chapter 3. The algorithm iteratively finds the middle point of the domain $x = (b + a)/2$, evaluates the function there, and decides whether the solution is on the left or the right, thus reducing the size of the domain from (a, b) to (a, x) or (x, b) , respectively:

Listing 4.50: in file: nlib.py

```

1 def solve_bisection(f, a, b, ap=1e-6, rp=1e-4, ns=100):
2     fa, fb = f(a), f(b)
3     if fa == 0: return a
4     if fb == 0: return b
5     if fa*fb > 0:
6         raise ArithmeticError('f(a) and f(b) must have opposite sign')
7     for k in xrange(ns):
8         x = (a+b)/2
9         fx = f(x)
10        if fx==0 or norm(b-a)<max(ap,norm(x)*rp): return x
11        elif fx * fa < 0: (b,fb) = (x, fx)
12        else: (a,fa) = (x, fx)
13    raise ArithmeticError('no convergence')

```

Here is how to use it:

Listing 4.51: in file: nlib.py

```

1 >>> def f(x): return (x-2)*(x-3)
2 >>> print round(solve_bisection(f,1.0,3.0),4)
3 2.0

```

4.6.3 Newton's method

The Newton [40] algorithm also solves $f(x) = 0$. It is faster (on average) than the bisection method because it makes the additional assumption that the function is also differentiable. This algorithm starts from an arbitrary point x_0 and approximates the function at that point with its first-order Taylor expansion

$$f(x) \simeq f(x_0) + f'(x_0)(x - x_0) \quad (4.118)$$

and solves it exactly:

$$f(x) = 0 \rightarrow x = x_0 - \frac{f(x_0)}{f'(x_0)} \quad (4.119)$$

thus finding a new and better estimate for the solution. The algorithm iterates the preceding equation, and when it converges, it approximates the exact solution better and better:

Listing 4.52: in file: nlib.py

```

1 def solve_newton(f, x, ap=1e-6, rp=1e-4, ns=20):
2     x = float(x) # make sure it is not int
3     for k in xrange(ns):
4         (fx, Dfx) = (f(x), D(f)(x))
5         if norm(Dfx) < ap:
6             raise ArithmeticError('unstable solution')
7         (x_old, x) = (x, x-fx/Dfx)
8         if k>2 and norm(x-x_old)<max(ap,norm(x)*rp): return x
9         raise ArithmeticError('no convergence')
```

The algorithm is guaranteed to converge if $|f'(x)| > 1$ in some neighborhood of the solution and if the starting point is in this neighborhood. It may also converge if this condition is not true. It is likely to fail when $|f'(x)| \simeq 0$ is in the neighborhood of the solution or the starting point because the terms fx/Dfx would become very large.

Here is an example:

Listing 4.53: in file: nlib.py

```

1 >>> def f(x): return (x-2)*(x-5)
2 >>> print round(solve_newton(f,1.0),4)
3 2.0
```

4.6.4 Secant method

The secant method is very similar to the Newton's method, except that $f'(x)$ is replaced by a numerical estimate computed using the current point x and the previous point visited by the algorithm:

$$f'(x_i) = \frac{f(x_i) - f(x_{i-1})}{x_i - x_{i-1}} \quad (4.120)$$

$$x_{i+1} = x_i - \frac{f(x_i)}{f'(x_i)} \quad (4.121)$$

As the algorithm approaches the exact solution, the numerical derivative becomes a better and better approximation for the derivative:

Listing 4.54: in file: nlib.py

```

1 def solve_secant(f, x, ap=1e-6, rp=1e-4, ns=20):
2     x = float(x) # make sure it is not int
3     (fx, Dfx) = (f(x), D(f)(x))
4     for k in xrange(ns):
5         if norm(Dfx) < ap:
6             raise ArithmeticError('unstable solution')
7         (x_old, fx_old, x) = (x, fx, x-fx/Dfx)
8         if k>2 and norm(x-x_old)<max(ap,norm(x)*rp): return x
9         fx = f(x)
10        Dfx = (fx-fx_old)/(x-x_old)
11    raise ArithmeticError('no convergence')

```

Here is an example:

Listing 4.55: in file: nlib.py

```

1 >>> def f(x): return (x-2)*(x-5)
2 >>> print round(solve_secant(f,1.0),4)
3 2.0

```

4.7 Optimization in one dimension

While a solver is an algorithm that finds x such that $f(x) = 0$, an optimization algorithm is one that finds the maximum or minimum of the function $f(x)$. If the function is differentiable, this is achieved by solving $f'(x) = 0$.

For this reason, if the function is differentiable twice, we can simply re-name all previous solvers and replace $f(x)$ with $f'(x)$ and $f'(x)$ with $f''(x)$.

4.7.1 Bisection method

Listing 4.56: in file: nlib.py

```
1 def optimize_bisection(f, a, b, ap=1e-6, rp=1e-4, ns=100):
2     return solve_bisection(D(f), a, b, ap, rp, ns)
```

Here is an example:

Listing 4.57: in file: nlib.py

```
1 >>> def f(x): return (x-2)*(x-5)
2 >>> print round(optimize_bisection(f,2.0,5.0),4)
3 3.5
```

4.7.2 Newton's method

Listing 4.58: in file: nlib.py

```
1 def optimize_newton(f, x, ap=1e-6, rp=1e-4, ns=20):
2     x = float(x) # make sure it is not int
3     (f, Df) = (D(f), DD(f))
4     for k in xrange(ns):
5         (fx, Dfx) = (f(x), Df(x))
6         if Dfx==0: return x
7         if norm(Dfx) < ap:
8             raise ArithmeticError('unstable solution')
9         (x_old, x) = (x, x-fx/Dfx)
10        if norm(x-x_old)<max(ap,norm(x)*rp): return x
11    raise ArithmeticError('no convergence')
```

Listing 4.59: in file: nlib.py

```
1 >>> def f(x): return (x-2)*(x-5)
2 >>> print round(optimize_newton(f,3.0),3)
3 3.5
```

4.7.3 Secant method

As in the Newton case, the secant method can also be used to find extrema, by replacing f with f' :

Listing 4.60: in file: nlib.py

```

1 def optimize_secant(f, x, ap=1e-6, rp=1e-4, ns=100):
2     x = float(x) # make sure it is not int
3     (f, Df) = (D(f), DD(f))
4     (fx, Dfx) = (f(x), Df(x))
5     for k in xrange(ns):
6         if fx==0: return x
7         if norm(Dfx) < ap:
8             raise ArithmeticError('unstable solution')
9         (x_old, fx_old, x) = (x, fx, x-fx/Dfx)
10        if norm(x-x_old)<max(ap,norm(x)*rp): return x
11        fx = f(x)
12        Dfx = (fx - fx_old)/(x-x_old)
13    raise ArithmeticError('no convergence')

```

Listing 4.61: in file: nlib.py

```

1 >>> def f(x): return (x-2)*(x-5)
2 >>> print round(optimize_secant(f,3.0),3)
3 3.5

```

4.7.4 Golden section search

If the function we want to optimize is continuous but not differentiable, then the previous algorithms do not work. In this case, there is one algorithm that comes to our rescue, the golden section [41] search. It is similar to the bisection method, with one caveat; in the bisection method, at each point, we need to know if a function changes sign in between two points, therefore two points are all we need. If instead we are looking for a max or min, we need to know if the function is concave or convex in between those two points. This requires one extra point in between the two. So while the bisection method only needs one point in between $[a, b]$, the golden search needs two points, x_1 and x_2 , in between $[a, b]$, and from them it can determine whether the extreme is in $[a, x_2]$ or in $[x_1, b]$. This is also represented pictorially in fig. 4.7.4. The two points are chosen in an optimal way so that at the next iteration, one of the two points can be recycled by leaving the ratio between $x_1 - a$ and $b - x_2$ fixed and equal to 1:

Listing 4.62: in file: nlib.py

```

1 def optimize_golden_search(f, a, b, ap=1e-6, rp=1e-4, ns=100):

```

```

2  a,b=float(a),float(b)
3  tau = (sqrt(5.0)-1.0)/2.0
4  x1, x2 = a+(1.0-tau)*(b-a), a+tau*(b-a)
5  fa, f1, f2, fb = f(a), f(x1), f(x2), f(b)
6  for k in xrange(ns):
7      if f1 > f2:
8          a, fa, x1, f1 = x1, f1, x2, f2
9          x2 = a+tau*(b-a)
10         f2 = f(x2)
11     else:
12         b, fb, x2, f2 = x2, f2, x1, f1
13         x1 = a+(1.0-tau)*(b-a)
14         f1 = f(x1)
15     if k>2 and norm(b-a)<max(ap,norm(b)*rp): return b
16     raise ArithmeticError('no convergence')

```

Here is an example:

Listing 4.63: in file: nlib.py

```

1 >>> def f(x): return (x-2)*(x-5)
2 >>> print round(optimize_golden_search(f,2.0,5.0),3)
3 3.5

```




Figure 4.7: Pictorial representation of the golden search method. If the function is concave ($f'''(x) > 0$), then knowledge of the function in 4 points (a, x_1, x_2, b) permits us to determine whether a minimum is between $[a, x_2]$ or between $[x_1, b]$.

4.8 Functions of many variables

To be able to work with functions of many variables, we need to introduce the concept of the partial derivative:

$$\frac{\partial f(\mathbf{x})}{\partial x_i} = \lim_{h \rightarrow 0} \frac{f(\mathbf{x} + \mathbf{h}_i) - f(\mathbf{x} - \mathbf{h}_i)}{2h} \quad (4.122)$$

where \mathbf{h}_i is a vector with components all equal to zero but $h_i = h > 0$.

We can implement it as follows:

Listing 4.64: in file: nlib.py

```

1 def partial(f,i,h=1e-4):
2     def df(x,f=f,i=i,h=h):
3         x = list(x) # make copy of x
4         x[i] += h
5         f_plus = f(x)
6         x[i] -= 2*h
7         f_minus = f(x)
8         if isinstance(f_plus,(list,tuple)):
9             return [(f_plus[i]-f_minus[i])/(2*h) for i in xrange(len(f_plus))]
10        else:
11            return (f_plus-f_minus)/(2*h)
12    return df

```

Similarly to $D(f)$, we have implemented it in such a way that `partial(f,i)` returns a function that can be evaluated at any point x . Also notice that the function f may return a scalar, a matrix, a list, or a tuple. The `if` condition allows the function to deal with the difference between two lists or tuples.

Here is an example:

Listing 4.65: in file: nlib.py

```

1 >>> def f(x): return 2.0*x[0]+3.0*x[1]+5.0*x[1]*x[2]
2 >>> df0 = partial(f,0)
3 >>> df1 = partial(f,1)
4 >>> df2 = partial(f,2)
5 >>> x = (1,1,1)
6 >>> print round(df0(x),4), round(df1(x),4), round(df2(x),4)
7 2.0 8.0 5.0

```


4.8.1 Jacobian, gradient, and Hessian

A generic function $f(x_0, x_1, x_2, \dots)$ of multiple variables $\mathbf{x} = (x_0, x_1, x_2, \dots)$ can be expanded in Taylor series to the second order as

$$\begin{aligned} f(x_0, x_1, x_2, \dots) &= f(\bar{x}_0, \bar{x}_1, \bar{x}_2, \dots) + \sum_i \frac{\partial f(\bar{\mathbf{x}})}{\partial x_i} (x_i - \bar{x}_i) \\ &+ \sum_{ij} \frac{1}{2} \frac{\partial^2 f}{\partial x_i \partial x_j} (\bar{\mathbf{x}}) (x_i - \bar{x}_i) (x_j - \bar{x}_j) + \dots \end{aligned} \quad (4.123)$$

We can rewrite the above expression in terms of the vector \mathbf{x} as follows:

$$f(\mathbf{x}) = f(\bar{\mathbf{x}}) + \nabla f(\bar{\mathbf{x}})(\mathbf{x} - \bar{\mathbf{x}}) + \frac{1}{2}(\mathbf{x} - \bar{\mathbf{x}})^t H_f(\bar{\mathbf{x}})(\mathbf{x} - \bar{\mathbf{x}}) + \dots \quad (4.124)$$

where we introduce the gradient vector

$$\nabla_f(x) \equiv \begin{pmatrix} \partial f(x)/\partial x_0 \\ \partial f(x)/\partial x_1 \\ \partial f(x)/\partial x_2 \\ \dots \end{pmatrix} \quad (4.125)$$

and the Hessian matrix

$$H_f(x) \equiv \begin{pmatrix} \partial^2 f(x)/\partial x_0 \partial x_0 & \partial^2 f(x)/\partial x_0 \partial x_1 & \partial^2 f(x)/\partial x_0 \partial x_2 & \dots \\ \partial^2 f(x)/\partial x_1 \partial x_0 & \partial^2 f(x)/\partial x_1 \partial x_1 & \partial^2 f(x)/\partial x_1 \partial x_2 & \dots \\ \partial^2 f(x)/\partial x_2 \partial x_0 & \partial^2 f(x)/\partial x_2 \partial x_1 & \partial^2 f(x)/\partial x_2 \partial x_2 & \dots \\ \dots & \dots & \dots & \dots \end{pmatrix} \quad (4.126)$$

Given the definition of partial, we can compute the gradient and the Hessian using the two functions

Listing 4.66: in file: `nlib.py`

```

1 def gradient(f, x, h=1e-4):
2     return Matrix(len(x),1,fill=lambda r,c: partial(f,r,h)(x))
3
4 def hessian(f, x, h=1e-4):
5     return Matrix(len(x),len(x),fill=lambda r,c: partial(partial(f,r,h),c,h)(x))

```

Here is an example:

Listing 4.67: in file: nlib.py

```

1 >>> def f(x): return 2.0*x[0]+3.0*x[1]+5.0*x[1]*x[2]
2 >>> print gradient(f, x=(1,1,1))
3 [[1.999999...], [7.999999...], [4.999999...]]
4 >>> print hessian(f, x=(1,1,1))
5 [[0.0, 0.0, 0.0], [0.0, 0.0, 5.000000...], [0.0, 5.000000..., 0.0]]

```

When dealing with functions returning multiple values like

$$f(\mathbf{x}) = (f_0(\mathbf{x}), f_1(\mathbf{x}), f_2(\mathbf{x}), \dots) \quad (4.127)$$

we need to Taylor expand each component:

$$f(\mathbf{x}) = \begin{pmatrix} f_0(\mathbf{x}) \\ f_1(\mathbf{x}) \\ f_2(\mathbf{x}) \\ \dots \end{pmatrix} = \begin{pmatrix} f_0(\bar{\mathbf{x}}) + \nabla_{f_0}(\mathbf{x} - \bar{\mathbf{x}}) + \dots \\ f_1(\bar{\mathbf{x}}) + \nabla_{f_1}(\mathbf{x} - \bar{\mathbf{x}}) + \dots \\ f_2(\bar{\mathbf{x}}) + \nabla_{f_2}(\mathbf{x} - \bar{\mathbf{x}}) + \dots \\ \dots \end{pmatrix} \quad (4.128)$$

which we can rewrite as

$$f(\mathbf{x}) = f(\bar{\mathbf{x}}) + J_f(\bar{\mathbf{x}})(\mathbf{x} - \bar{\mathbf{x}}) + \dots \quad (4.129)$$

where J_f is called Jacobian and is defined as

$$J_f \equiv \begin{pmatrix} \partial f_0(x)/\partial x_0 & \partial f_0(x)/\partial x_1 & \partial f_0(x)/\partial x_2 & \dots \\ \partial f_1(x)/\partial x_0 & \partial f_1(x)/\partial x_1 & \partial f_1(x)/\partial x_2 & \dots \\ \partial f_2(x)/\partial x_0 & \partial f_2(x)/\partial x_1 & \partial f_2(x)/\partial x_2 & \dots \\ \dots & \dots & \dots & \dots \end{pmatrix} \quad (4.130)$$

which we can implement as follows:

Listing 4.68: in file: nlib.py

```

1 def jacobian(f, x, h=1e-4):
2     partials = [partial(f,c,h)(x) for c in xrange(len(x))]
3     return Matrix(len(partial[0]),len(x),fill=Lambda r,c: partials[c][r])

```

Here is an example:

Listing 4.69: in file: nlib.py

```

1 >>> def f(x): return (2.0*x[0]+3.0*x[1]+5.0*x[1]*x[2], 2.0*x[0])
2 >>> print jacobian(f, x=(1,1,1))
3 [[1.9999999..., 7.999999..., 4.999999...], [1.999999..., 0.0, 0.0]]

```

4.8.2 Newton's method (solver)

We can now solve eq. 4.129 iteratively as we did for the one-dimensional Newton solver with only one change—the first derivative of f is replaced by the Jacobian:

Listing 4.70: in file: nlib.py

```

1 def solve_newton_multi(f, x, ap=1e-6, rp=1e-4, ns=20):
2     """
3     Computes the root of a multidimensional function f near point x.
4
5     Parameters
6     f is a function that takes a list and returns a scalar
7     x is a list
8
9     Returns x, solution of f(x)=0, as a list
10    """
11    n = len(x)
12    x = Matrix(len(x))
13    for k in xrange(ns):
14        fx = Matrix(f(x.flatten()))
15        J = jacobian(f,x.flatten())
16        if norm(J) < ap:
17            raise ArithmeticError('unstable solution')
18        (x_old, x) = (x, x-(1.0/J)*fx)
19        if k>2 and norm(x-x_old)<max(ap,norm(x)*rp): return x.flatten()
20    raise ArithmeticError('no convergence')

```

Here is an example:

Listing 4.71: in file: nlib.py

```

1 >>> def f(x): return [x[0]+x[1], x[0]+x[1]**2-2]
2 >>> print solve_newton_multi(f, x=(0,0))
3 [1.0..., -1.0...]

```

4.8.3 Newton's method (optimize)

As for the one-dimensional case, we can approximate $f(\mathbf{x})$ with its Taylor expansion at the first order,

$$f(\mathbf{x}) = f(\bar{\mathbf{x}}) + \nabla_f(\bar{\mathbf{x}})(\mathbf{x} - \bar{\mathbf{x}}) + \frac{1}{2}(\mathbf{x} - \bar{\mathbf{x}})^t H_f(\bar{\mathbf{x}})(\mathbf{x} - \bar{\mathbf{x}}) \quad (4.131)$$

set its derivative to zero, and solve it, thus obtaining

$$\mathbf{x} = \bar{\mathbf{x}} - H_f^{-1} \nabla_f \quad (4.132)$$

which constitutes the core of the multidimensional Newton optimizer:

Listing 4.72: in file: nlib.py

```

1 def optimize_newton_multi(f, x, ap=1e-6, rp=1e-4, ns=20):
2     """
3     Finds the extreme of multidimensional function f near point x.
4
5     Parameters
6     f is a function that takes a list and returns a scalar
7     x is a list
8
9     Returns x, which maximizes or minimizes f(x)=0, as a list
10    """
11    x = Matrix(list(x))
12    for k in xrange(ns):
13        (grad,H) = (gradient(f,x.flatten()), hessian(f,x.flatten()))
14        if norm(H) < ap:
15            raise ArithmeticError('unstable solution')
16        (x_old, x) = (x, x-(1.0/H)*grad)
17        if k>2 and norm(x-x_old)<max(ap,norm(x)*rp): return x.flatten()
18    raise ArithmeticError('no convergence')

```

Listing 4.73: in file: nlib.py

```

1 >>> def f(x): return (x[0]-2)**2+(x[1]-3)**2
2 >>> print optimize_newton_multi(f, x=(0,0))
3 [2.0, 3.0]

```

4.8.4 Improved Newton's method (optimize)

We can further improve the Newton multidimensional optimizer by using the following technique. At each step, if the next guess does not

reduce the value of f , we revert to the previous point, and we perform a one-dimensional Newton optimization along the direction of the gradient. This method greatly increases the stability of the multidimensional Newton optimizer:

Listing 4.74: in file: nlib.py

```

1 def optimize_newton_multi_improved(f, x, ap=1e-6, rp=1e-4, ns=20, h=10.0):
2     """
3     Finds the extreme of multidimensional function f near point x.
4
5     Parameters
6     f is a function that takes a list and returns a scalar
7     x is a list
8
9     Returns x, which maximizes or minimizes f(x)=0, as a list
10    """
11    x = Matrix(list(x))
12    fx = f(x.flatten())
13    for k in xrange(ns):
14        (grad,H) = (gradient(f,x.flatten()), hessian(f,x.flatten()))
15        if norm(H) < ap:
16            raise ArithmeticError('unstable solution')
17        (fx_old, x_old, x) = (fx, x, x-(1.0/H)*grad)
18        fx = f(x.flatten())
19        while fx>fx_old: # revert to steepest descent
20            (fx, x) = (fx_old, x_old)
21            norm_grad = norm(grad)
22            (x_old, x) = (x, x - grad/norm_grad*h)
23            (fx_old, fx) = (fx, f(x.flatten()))
24            h = h/2
25        h = norm(x-x_old)*2
26        if k>2 and h/2<max(ap,norm(x)*rp): return x.flatten()
27    raise ArithmeticError('no convergence')

```

4.9 Nonlinear fitting

Finally, we have all the ingredients to implement a very generic fitting function that will work linear and nonlinear least squares.

Here we consider a generic experiment or simulated experiment that generates points of the form $(x_i, y_i \pm \delta y_i)$. Our goal is to minimize the χ^2

defined as

$$\chi^2(\mathbf{a}, \mathbf{b}) = \sum_i \left| \frac{y_i - f(x_i, \mathbf{a}, \mathbf{b})}{\delta y_i} \right|^2 \quad (4.133)$$

where the function f is known but depends on unknown parameters $\mathbf{a} = (a_0, a_1, \dots)$ and $\mathbf{b} = (b_0, b_1, \dots)$. In terms of these parameters, the function f can be written as follows:

$$f(x, \mathbf{a}, \mathbf{b}) = \sum_j a_j f_j(x, \mathbf{b}) \quad (4.134)$$

Here is an example:

$$f(x, \mathbf{a}, \mathbf{b}) = a_0 e^{-b_0 x} + a_1 e^{-b_1 x} + a_2 e^{-b_2 x} + \dots \quad (4.135)$$

The goal of our algorithm is to efficiently determine the parameters \mathbf{a} and \mathbf{b} that minimize the χ^2 .

We proceed by defining the following two quantities:

$$\mathbf{z} = \begin{pmatrix} y_0 / \delta y_0 \\ y_1 / \delta y_1 \\ y_2 / \delta y_2 \\ \dots \end{pmatrix} \quad (4.136)$$

and

$$A(\mathbf{b}) = \begin{pmatrix} f_0(x_0, \mathbf{b}) / \delta y_0 & f_1(x_0, \mathbf{b}) / \delta y_0 & f_2(x_0, \mathbf{b}) / \delta y_0 & \dots \\ f_0(x_1, \mathbf{b}) / \delta y_1 & f_1(x_1, \mathbf{b}) / \delta y_1 & f_2(x_1, \mathbf{b}) / \delta y_1 & \dots \\ f_0(x_2, \mathbf{b}) / \delta y_2 & f_1(x_2, \mathbf{b}) / \delta y_2 & f_2(x_2, \mathbf{b}) / \delta y_2 & \dots \\ \dots & \dots & \dots & \dots \end{pmatrix} \quad (4.137)$$

In terms of A and \mathbf{z} , the χ^2 can be rewritten as

$$\chi^2(\mathbf{a}, \mathbf{b}) = |A(\mathbf{b})\mathbf{a} - \mathbf{z}|^2 \quad (4.138)$$

We can minimize this function in \mathbf{a} using the linear least squares algorithm, exactly:

$$\mathbf{a}(\mathbf{b}) = (A(\mathbf{b})A(\mathbf{b})^t)^{-1}(A(\mathbf{b})^t\mathbf{z}) \quad (4.139)$$

We define a function that returns the minimum χ^2 for a fixed input \mathbf{b} :

$$g(\mathbf{b}) = \min_{\mathbf{a}} \chi^2(\mathbf{a}, \mathbf{b}) = \chi^2(\mathbf{a}(\mathbf{b}), \mathbf{b}) = |A(\mathbf{b})\mathbf{a}(\mathbf{b}) - \mathbf{z}|^2 \quad (4.140)$$

Therefore we have reduced the original problem to a simple problem by reducing the number of unknown parameters from $N_a + N_b$ to N_b .

The following code takes as input the data as a list of $(x_i, y_i, \delta y_i)$, a list of functions (or a single function), and a guess for the \mathbf{b} values. If the `fs` argument is not a list but a single function, then there is no \mathbf{a} to compute, and the function proceeds by minimizing the χ^2 using the improved Newton optimizer (the one-dimensional or the improved multidimensional one, as appropriate). If the argument \mathbf{b} is missing, then the fitting parameters are all linear, and the algorithm reverts to regular linear least squares. Otherwise, run the more complex algorithm described earlier:

Listing 4.75: in file: `nlib.py`

```

1 def fit(data, fs, b=None, ap=1e-6, rp=1e-4, ns=200, constraint=None):
2     if not isinstance(fs, (list, tuple)):
3         def g(b, data=data, f=fs, constraint=constraint):
4             chi2 = sum(((y-f(b,x))/dy)**2 for (x,y,dy) in data)
5             if constraint: chi2+=constraint(b)
6             return chi2
7         if isinstance(b, (list, tuple)):
8             b = optimize_newton_multi_improved(g,b,ap,rp,ns)
9         else:
10            b = optimize_newton(g,b,ap,rp,ns)
11        return b, g(b,data,constraint=None)
12    elif not b:
13        a, chi2, ff = fit_least_squares(data, fs)
14        return a, chi2
15    else:
16        na = len(fs)
17        def core(b,data=data,fs=fs):
18            A = Matrix([[fs[k](b,x)/dy for k in xrange(na)] \
19                        for (x,y,dy) in data])

```

```

20     z = Matrix([[y/dy] for (x,y,dy) in data])
21     a = (1/(A.T*A))*(A.T*z)
22     chi2 = norm(A*a-z)**2
23     return a.flatten(), chi2
24 def g(b,data=data,fs=fs,constraint=constraint):
25     a, chi2 = core(b, data, fs)
26     if constraint:
27         chi += constraint(b)
28     return chi2
29 b = optimize_newton_multi_improved(g,b,ap,rp,ns)
30 a, chi2 = core(b,data,fs)
31 return a+b,chi2

```

Here is an example:

```

1 >>> data = [(i, i+2.0*i**2+300.0/(i+10), 2.0) for i in xrange(1,10)]
2 >>> fs = [(lambda b,x: x), (lambda b,x: x*x), (lambda b,x: 1.0/(x+b[0]))]
3 >>> ab, chi2 = fit(data,fs,[5])
4 >>> print ab, chi2
5 [0.999..., 2.000..., 300.000..., 10.000...] ...

```

In the preceding implementation, we added a somewhat mysterious argument `constraint`. This is a function of `b`, and its output gets added to the value of χ^2 , which we are minimizing. By choosing the appropriate function, we can set constraints about the expected values `b`. These constraints represent a priori knowledge about the parameters, that is, knowledge that does not come from the data being fitted.

For example, if we know that b_i must be close to some \bar{b}_i with some uncertainty δb_i , then we can use

```

1 def constraint(b, bar_b, delta_b):
2     return sum(((b[i]-bar_b[i])/delta_b[i])**2 for i in xrange(len(b)))

```

and pass the preceding function as a constraint. From a practical effect, this stabilizes our fit. From a theoretical point of view, the \bar{b}_i are the *priors* of Bayesian statistics.

4.10 Integration

Consider the integral of $f(x)$ for x in domain $[a, b]$, which we normally represent as

$$I = \int_a^b f(x) dx \quad (4.141)$$

and which measures the area under the curve $y = f(x)$ delimited on the left by $x = a$ and on the right by $x = b$.




Figure 4.8: Visual representation of the concept of an integral as the area under a curve.

As we did in the previous subsection, we can approximate the possible values taken by x as discrete values $x \equiv hi$, where $h = (b - a)/n$. At those values, the function f evaluates to $f_i \equiv f(hi)$. Thus the integral can be approximated as a sum of trapezoids:

$$I_n \simeq \sum_{i=0}^{i \leq n} \frac{h}{2} (f_i + f_{i+1}) \quad (4.142)$$

If a function is discontinuous only in a finite number of points in the domain $[a, b]$, then the following limit exists:

$$\lim_{n \rightarrow \infty} I_n \rightarrow I \quad (4.143)$$

We can implement the naive integration as a function of N as follows:




Figure 4.9: Visual representation of the trapezoid method for numerical integration.

Listing 4.76: in file: nlib.py

```

1 def integrate_naive(f, a, b, n=20):
2     """
3     Integrates function, f, from a to b using the trapezoidal rule
4     >>> from math import sin
5     >>> integrate(sin, 0, 2)
6     1.416118...
7     """
8     a,b= float(a),float(b)
9     h = (b-a)/n
10    return h/2*(f(a)+f(b))+h*sum(f(a+h*i) for i in xrange(1,n))

```

And here we implement the limit by doubling the number of points until convergence is achieved:

Listing 4.77: in file: nlib.py

```

1 def integrate(f, a, b, ap=1e-4, rp=1e-4, ns=20):
2     """
3     Integrates function, f, from a to b using the trapezoidal rule
4     converges to precision
5     """
6     I = integrate_naive(f,a,b,1)
7     for k in xrange(1,ns):
8         I_old, I = I, integrate_naive(f,a,b,2**k)
9         if k>2 and norm(I-I_old)<max(ap,norm(I)*rp): return I
10    raise ArithmeticError('no convergence')

```

We can test the convergence as follows:

Listing 4.78: in file: nlib.py

```

1 >>> from math import sin, cos
2 >>> print integrate_naive(sin,0,3,n=2)
3 1.6020...

```

```

4 >>> print integrate_naive(sin,0,3,n=4)
5 1.8958...
6 >>> print integrate_naive(sin,0,3,n=8)
7 1.9666...
8 >>> print integrate(sin,0,3)
9 1.9899...
10 >>> print 1.0-cos(3)
11 1.9899...

```

4.10.1 Quadrature

In the previous integration, we divided the domain $[a, b]$ into subdomains, and we computed the area under the curve f in each subdomain by approximating it with a trapezoid; for example, we approximated the function in between x_i and x_{i+1} with a straight line. We can do better by approximating the function with a polynomial of arbitrary degree n and then compute the area in the subdomain by explicitly integrating the polynomial.

This is the basic idea of quadrature. For a subdomain delimited by $(0, 1)$, we can impose

$$\int_0^1 1 dx = h = \sum_i c_i (i/n)^0 \quad (4.144)$$

$$\int_0^1 x dx = h^2/2 = \sum_i c_i (i/n)^1 \quad (4.145)$$

$$\dots = \dots = \dots \quad (4.146)$$

$$\int_0^1 x^{n-1} dx = h^n/n = \sum_i c_i (i/n)^2 \quad (4.147)$$

where c_i are coefficients to be determined:

Listing 4.79: in file: nlib.py

```

1 class QuadratureIntegrator:
2     """
3     Calculates the integral of the function f from points a to b
4     using n Vandermonde weights and numerical quadrature.
5     """
6     def __init__(self, order=4):
7         h = 1.0/(order-1)

```

```

8      A = Matrix(order, order, fill = lambda r,c: (c*h)**r)
9      s = Matrix(order, 1, fill = lambda r,c: 1.0/(r+1))
10     w = (1/A)*s
11     self.w = w
12     def integrate(self,f,a,b):
13         w = self.w
14         order = len(w.rows)
15         h = float(b-a)/(order-1)
16         return (b-a)*sum(w[i,0]*f(a+i*h) for i in xrange(order))
17
18 def integrate_quadrature_naive(f,a,b,n=20,order=4):
19     a,b = float(a),float(b)
20     h = float(b-a)/n
21     q = QuadratureIntegrator(order=order)
22     return sum(q.integrate(f,a+i*h,a+i*h+h) for i in xrange(n))

```

Here is an example of usage:

Listing 4.80: in file: nlib.py

```

1 >>> from math import sin
2 >>> print integrate_quadrature_naive(sin,0,3,n=2,order=2)
3 1.60208248595
4 >>> print integrate_quadrature_naive(sin,0,3,n=2,order=3)
5 1.99373945223
6 >>> print integrate_quadrature_naive(sin,0,3,n=2,order=4)
7 1.99164529955

```

4.11 Fourier transforms

A function with a domain over a finite interval $[a, b]$ can be approximated with a vector. For example, consider a function $f(x)$ with domain $[0, T]$. We can sample the function at points $x_k = a + (b - a)k/N$ and represent the discretized function with a vector

$$\mathbf{u}_f \equiv \{cf(x_0), cf(x_1), cf(x_2), \dots, cf(x_N)\} \quad (4.148)$$

where c is an arbitrary constant that we choose to be $c = \sqrt{(b - a)/N}$. This choice simplifies our later algebra. Summarizing, we define

$$u_{fk} \equiv \sqrt{\frac{b - a}{N}} f(x_k) \quad (4.149)$$

Given any two functions, we can define their scalar product as the limit of $N \rightarrow \infty$ of the scalar product between their corresponding vectors:

$$f \cdot g \equiv \lim_{N \rightarrow \infty} \mathbf{u}_f \cdot \mathbf{u}_g = \lim_{N \rightarrow \infty} \frac{b - a}{N} \sum_k f(x_k) g(x_k) \quad (4.150)$$

Using the definition of integral, it can be proven that, in the limit $N \rightarrow \infty$, this is equivalent to

$$f \cdot g = \int_a^b f(x) g(x) dx \quad (4.151)$$

This is because we have chosen c such that c^2 is the width of a rectangle in the Riemann integration.

From now on, we will omit the f subscript in \mathbf{u} and simply use different letters for vectors representing different sampled functions (\mathbf{u} , \mathbf{v} , \mathbf{b} , etc.).

Because we are interested in numerical algorithms, we will keep N finite and work with the sum instead of the integral.

Given a fixed N , we can always find N vectors $\mathbf{b}_0, \mathbf{b}_1, \dots, \mathbf{b}_{N-1}$ that are lin-

early independent, normalized, and orthogonal, that is,

$$\mathbf{b}_i \cdot \mathbf{b}_j = \sum_k b_{ik} b_{jk} = \delta_{ij} \quad (4.152)$$

Here b_{jk} is the k component of vector \mathbf{b}_j and δ_{ij} is the Kronecker delta defined as 0 when $i \neq j$ and 1 when $i = j$.

Any set of vectors $\{\mathbf{b}_j\}$ meeting the preceding condition is called an orthonormal basis. Any other vector \mathbf{u} can be represented by its projections over the basis vectors:

$$u_i = \sum_j v_j b_{ji} \quad (4.153)$$

where v_j is the projection of u along \mathbf{b}_j , which can be computed as

$$v_j = \sum_i u_i b_{ji} \quad (4.154)$$

In fact, by direct substitution, we obtain

$$v_j = \sum_k u_k b_{jk} \quad (4.155)$$

$$= \sum_k \left(\sum_i v_i b_{ik} \right) b_{jk} \quad (4.156)$$

$$= \sum_i v_i \left(\sum_k b_{ik} b_{jk} \right) \quad (4.157)$$

$$= \sum_i v_i \delta_{ij} \quad (4.158)$$

$$= v_j \quad (4.159)$$

In other words, once we have a basis of vectors, the vector \mathbf{u} can be represented in terms of the vector \mathbf{v} of v_j coefficients and, conversely, \mathbf{v} can be computed from \mathbf{u} ; \mathbf{u} and \mathbf{v} contain the same information.

The transformation from \mathbf{u} to \mathbf{v} , and vice versa, is a linear transformation. We call T^+ the transformation from \mathbf{u} to \mathbf{v} and T^- its inverse:

$$\mathbf{v} = T^+(\mathbf{u}) \quad \mathbf{u} = T^-(\mathbf{v}) \quad (4.160)$$

From the definition, and without attempting any optimization, we can implement these operators as follows:

```

1 def transform(u,b):
2     return [sum(u[k]*bi[k] for k in xrange(len(u))) for bi in b]
3
4 def antitransform(v,b):
5     return [sum(v[i]*bi[k] for i,bi in enumerate(b)) for k in xrange(len(v))]

```

Here is an example of usage:

```

1 >>> def make_basis(N):
2 >>>     return [[1 if i==j else 0 for i in xrange(N)] for j in xrange(N)]
3 >>> b = make_basis(4)
4 >>> print b
5 [[1, 0, 0, 0], [0, 1, 0, 0], [0, 0, 1, 0], [0, 0, 0, 1]]
6 >>> u = [1.0, 2.0, 3.0, 4.0]
7 >>> v = transform(u,b)
8 >>> print antitransform(v,b)
9 [1.0, 2.0, 3.0, 4.0]

```

Of course, this example is trivial because of the choice of basis which makes \mathbf{v} the same as \mathbf{u} . Yet our argument works for any basis \mathbf{b}_i . In particular, we can make the following choice:

$$b_{ji} = \frac{1}{\sqrt{2\pi}} e^{2\pi Iij/N} \quad (4.161)$$

where I is the imaginary unit. With this choice, the T^+ and T^- functions become

$$v_j = N^{-\frac{1}{2}} \sum_i u_i e^{2\pi Iij/N} \quad (4.162)$$

$$u_i = N^{-\frac{1}{2}} \sum_j v_j e^{-2\pi Iij/N} \quad (4.163)$$

and they take the names of Fourier transform and anti-transform [42], respectively; we can implement them as follows:

```

1 from cmath import exp as cexp
2
3 def fourier(u, sign=1):
4     N, D = len(u), xrange(len(u))
5     coeff, omega = 1.0/sqrt(N), 2.0*pi*sign*(1j)/N
6     return [sum(coeff*u[i]*cexp(omega*i*j) for i in D) for j in D]
7
8 def anti_fourier(v):
9     return fourier(v, sign=-1)

```

Here $1j$ is the Python notation for I and *cexp* is the exponential function for complex numbers.

Notice how the transformation works even when u is a vector of complex numbers.

Something special happens when u is real:

$$\operatorname{Re}(v_j) = +\operatorname{Re}(v_{N-j-1}) \quad (4.164)$$

$$\operatorname{Im}(v_j) = -\operatorname{Im}(v_{N-j-1}) \quad (4.165)$$

We can speed up the code even more using recursion and by observing that if N is a power of 2, then

$$v_j = N^{-\frac{1}{2}} \sum_i u_{2i} e^{2\pi I(2i)j/N} + N^{-\frac{1}{2}} \sum_i u_{2i+1} e^{2\pi I(2i+1)j/N} \quad (4.166)$$

$$= 2^{-\frac{1}{2}} (v_j^{\text{even}} + e^{2\pi j/N} v_j^{\text{even}}) \quad (4.167)$$

where v_j^{even} is the Fourier transform of the even terms and v_j^{odd} is the Fourier transform of the odd terms.

The preceding recursive expression can be implemented using dynamic programming, thus obtaining

```

1 from cmath import exp as cexp
2
3 def fast_fourier(u, sign=1):
4     N, sqrtN, D = len(u), sqrt(len(u)), xrange(len(u))
5     v = [ui/sqrtN for ui in u]
6     k = N/2
7     while k:
8         omega = cexp(2.0*pi*1j*k/N)

```



```

9     for i in D:
10         j = i ^ k
11         if i < k:
12             ik, jk = int(i/k), int(j/k)
13             v[i], v[j] = v[i]+(omega**ik)*v[j], v[i]+(omega**jk)*v[j]
14         k/=2
15     return v
16
17 def fast_anti_fourier(v):
18     return fast_fourier(v, sign=-1)

```

This implementation of the Fourier transform is equivalent to the previous one in the sense that it produces the same result (up to numerical issues), but it is faster as it runs in $\Theta(N \log_2 N)$ versus $\Theta(N^2)$ of the naive implementation. Here $i \wedge j$ is a binary operator, specifically a XOR. For each binary digit of i , it returns a flipped bit if the corresponding bit in j is 1. For example:

```

1 i : 10010010101110
2 j : 00010001000010
3 i^j: 100000111001110

```

4.12 Differential equations

In this section, we deal specifically with differential equations of the following form:

$$a_0 f(x) + a_1 f'(x) + a_2 f''(x) + \cdots = g(x) \quad (4.168)$$

where $f(x)$ is an unknown function to be determined; f' , f'' , and so on, are its derivatives; a_i are known input coefficients; and $g(x)$ is a known input function:

$$f''(x) - 4f'(x) + f(x) = \sin(x) \quad (4.169)$$

In this case, $a_2(x) = 1$, $a_1(x) = -4$, $a_0(x) = 1$, and $g(x) = \sin(x)$.

This can be solved using Fourier transforms by observing that if the Fourier transform of $f(x)$ is $\tilde{f}(y)$, then the Fourier transform of $f'(x)$ is $iy\tilde{f}(y)$.

Hence, if we Fourier transform both the left and right side of

$$\sum_k a_k f^{(k)}(x) = g(x) \quad (4.170)$$

we obtain

$$(\sum_k a_k (iy)^k) \tilde{f}(y) = \tilde{g}(y) \quad (4.171)$$

therefore $f(x)$ is the anti-Fourier transform of

$$\tilde{f}(y) = \frac{\tilde{g}(y)}{(\sum_k a_k (iy)^k)} \quad (4.172)$$

In one equation, the solution of eq. 4.168 is

$$f(x) = T^-(T^+(g)/(\sum_k a_k (iy)^k)) \quad (4.173)$$

This is fine and useful when the Fourier transformations are easy to compute.

A more practical numerical solution is the following. We define

$$y_i(x) \equiv f^{(i)}(x) \quad (4.174)$$

and we rewrite the differential equation as

$$y'_0 = y_1 \quad (4.175)$$

$$y'_1 = y_2 \quad (4.176)$$

$$y'_2 = y_3 \quad (4.177)$$

$$\dots = \dots \quad (4.178)$$

$$y'_{N-1} = y_N = (g(x) - \sum_{k < N} a_k y_k(x)) / a_N(x) \quad (4.179)$$

or equivalently

$$\mathbf{y}' = F(\mathbf{y}) \quad (4.180)$$

where

$$F(\mathbf{y}) = \mathbf{y} + \begin{pmatrix} y_1 \\ y_2 \\ y_3 \\ \dots \\ (g(x) - \sum_{k < N} a_k(x)y_k(x))/a_N(x) \end{pmatrix} \quad (4.181)$$

The naive solution is due to Euler:

$$\mathbf{y}(x+h) = \mathbf{y}(x) + hF(\mathbf{y}, x) \quad (4.182)$$

The solution is found by iterating the latest equation. Here h is an arbitrary discretization step. Euler's method works even if the a_k coefficients depend on x .

Although the Euler integrator works in theory, its systematic error adds up and does not disappear in the limit $h \rightarrow 0$. More accurate integrators are the Runge–Kutta and the Adam–Bashforth. In the fourth-order Runge–Kutta, the *classical Runge–Kutta method*, we also solve the differential equation by iteration, except that eq. 4.182 is replaced with

$$\mathbf{y}(x+h) = \mathbf{y}(x) + h/6(\mathbf{k}_1 + 2\mathbf{k}_2 + 2\mathbf{k}_3 + \mathbf{k}_4) \quad (4.183)$$

where

$$\mathbf{k}_1 = F(\mathbf{y}, x) \quad (4.184)$$

$$\mathbf{k}_2 = F(\mathbf{y} + hk_1/2, x + h/2) \quad (4.185)$$

$$\mathbf{k}_3 = F(\mathbf{y} + hk_2/2, x + h/2) \quad (4.186)$$

$$\mathbf{k}_4 = F(\mathbf{y} + hk_3, x + h) \quad (4.187)$$

